

COMPUTER ARCHITECTURE

Chapter 4: Microarchitecture



Introduction

- CPU performance factors
 - Instruction count
 - Determined by ISA and compiler
 - CPI and Cycle time
 - Determined by CPU hardware
- We will examine two MIPS implementations
 - A simplified version: CPI = 1
 - A more realistic pipelined version: CPI \approx 1
- Simple subset, shows most aspects
 - Memory reference: `lw`, `sw`
 - Arithmetic/logical: `add`, `sub`, `and`, `or`, `slt`
 - Control transfer: `beq`, `j`



Instruction execution

1. PC → **instruction memory (cache)**, fetch instruction
2. Register numbers → **registers file**, read registers
3. Depending on instruction class
 - 3.1. Use ALU to calculate
 - Arithmetic result → **done**
 - Memory address for load/store → **3.2**
 - Branch target address → **3.3**
 - 3.2. Access data memory for load/store
 - 3.3. $\text{PC} \leftarrow \text{target address or PC} + 4$



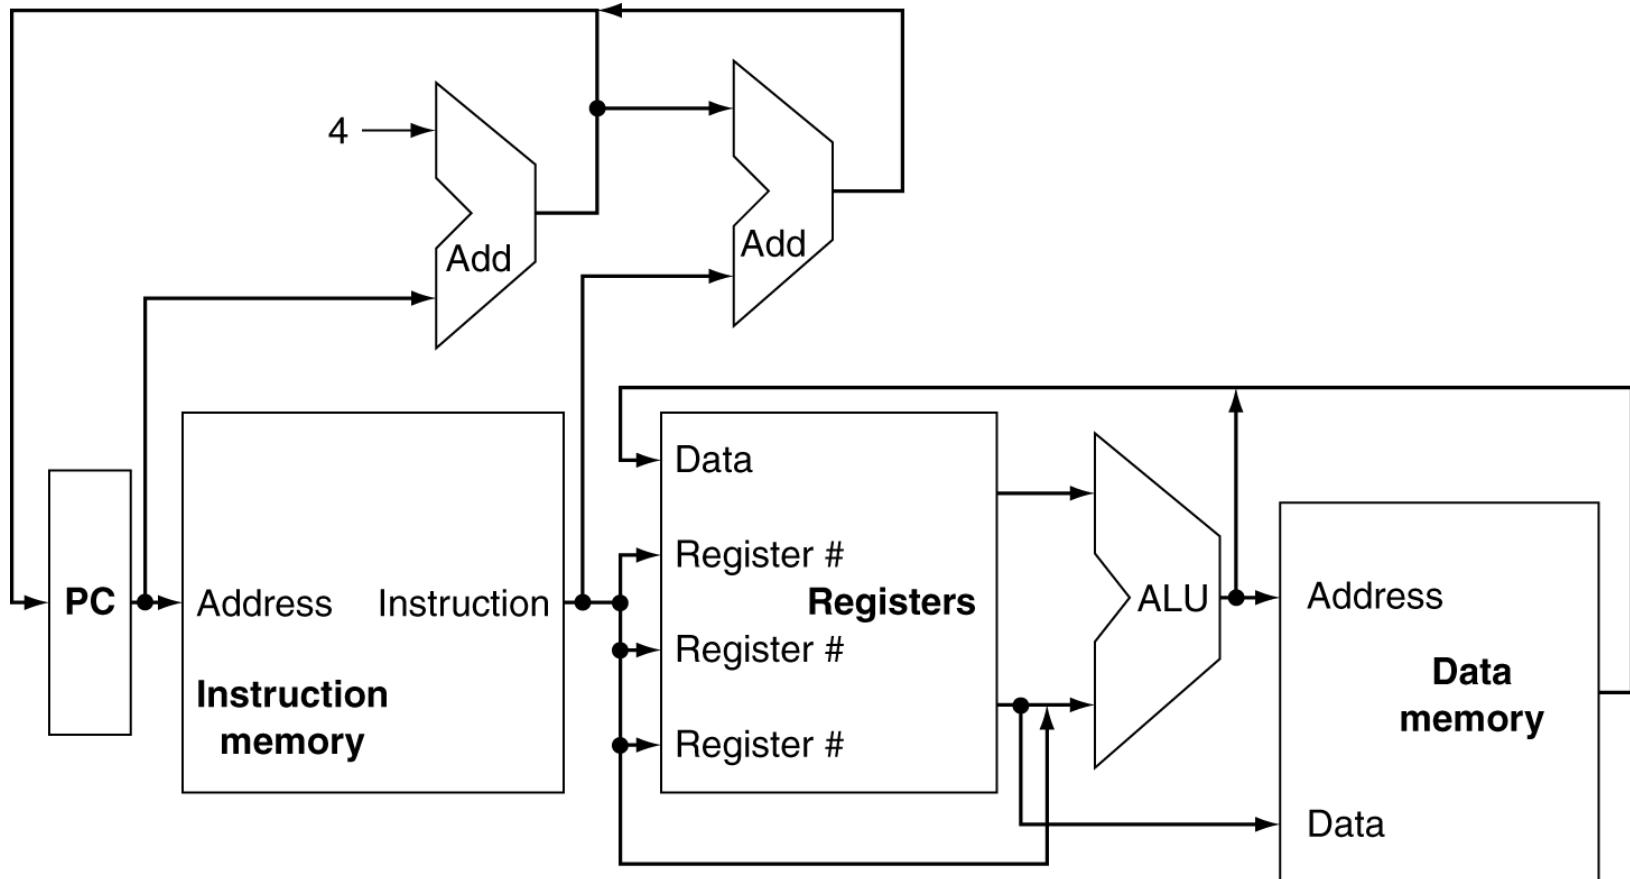
Execution stages

1. **Instruction fetch (IF)**: PC → instruction address
2. **Instruction decode (ID)**: register operands → register file
3. **Execute (EXE)**:
 - Load/store: compute a memory address
 - Arithmetic/logical: compute an arithmetic/logical result
4. **Memory access (MEM)**:
 - Load: read data memory
 - Store: write data memory
5. **Write back (WB)**:
 - Store a result of register file

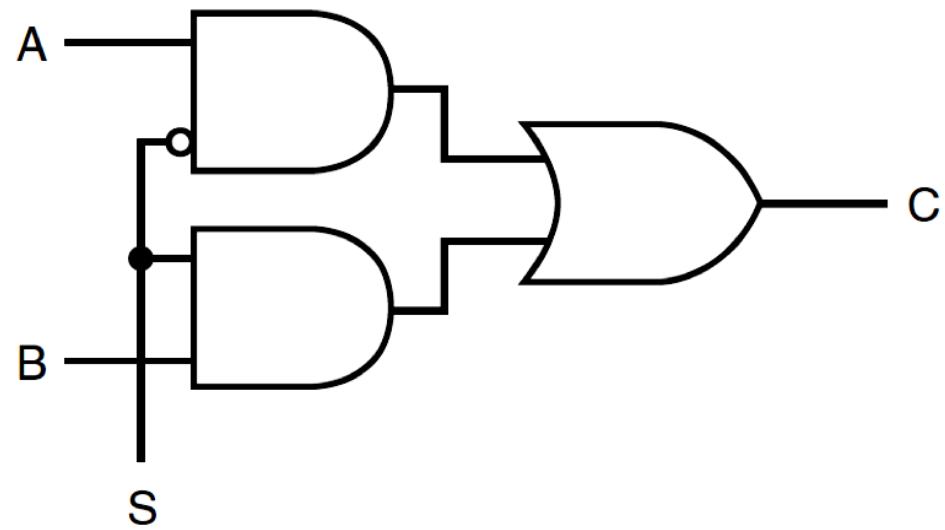
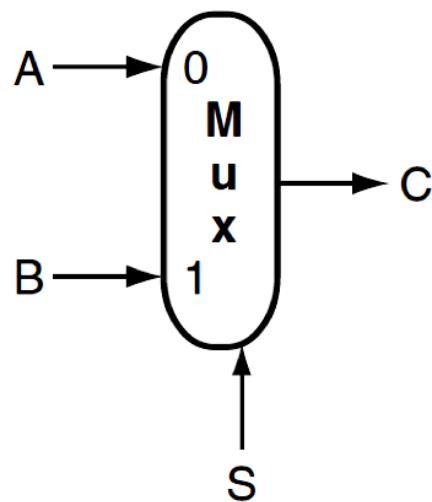
Datapath - controller

- **Datapath:** contains information that is operated on by a functional unit
 - Instruction memory: contain instructions (**IF**)
 - Registers file: 32 32-bit registers (**ID** & **WB**)
 - ALU: calculate arithmetic and logical operations (**EXE**)
 - Data memory: contain data (**MEM**)
- **Control signals:** used for multiplexer selection or for directing the operation of a functional unit
 - Control unit
 - Multiplexer

Datapath overview

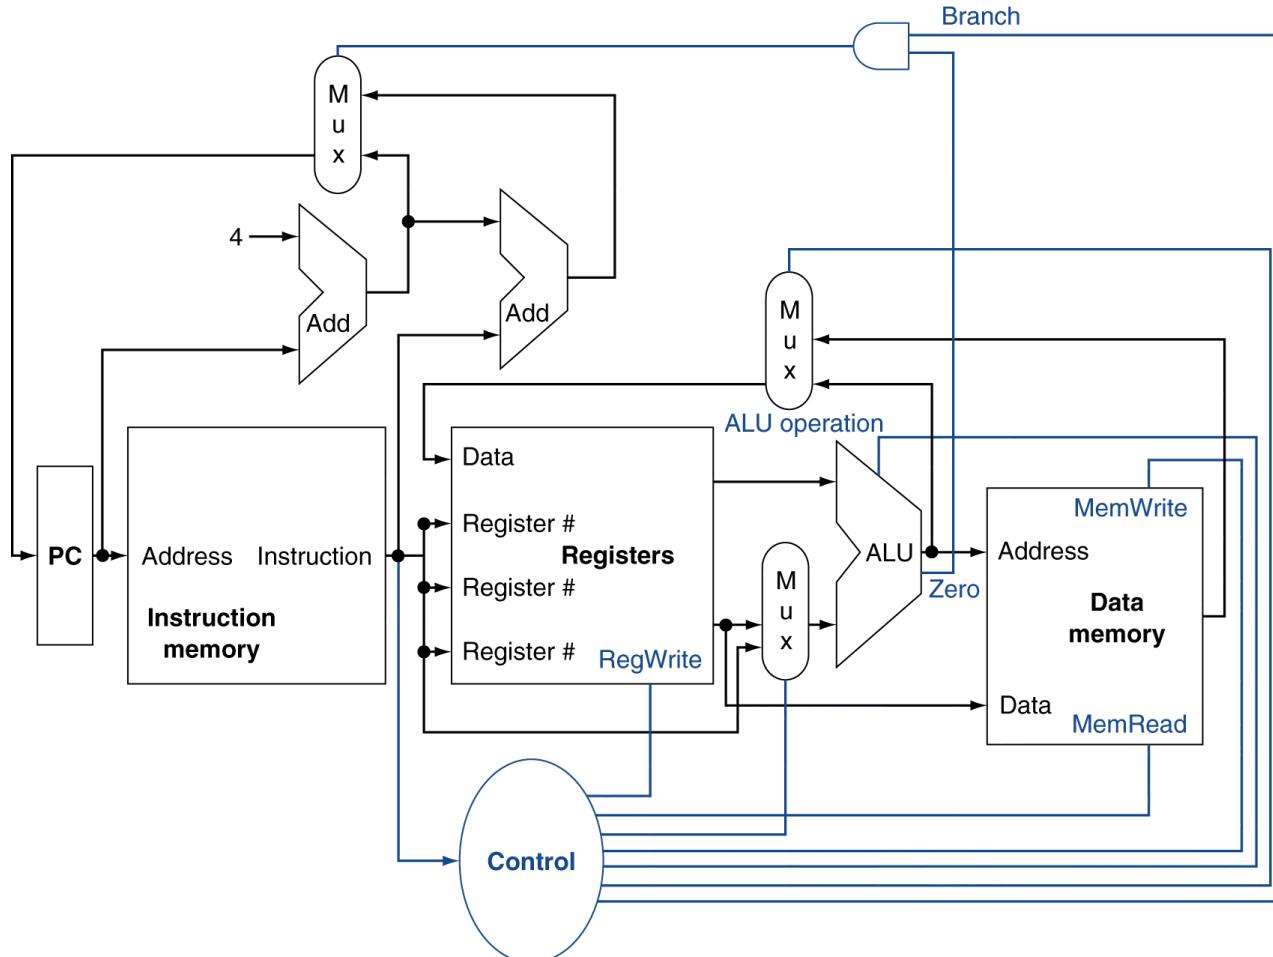


Multiplexer - MUX



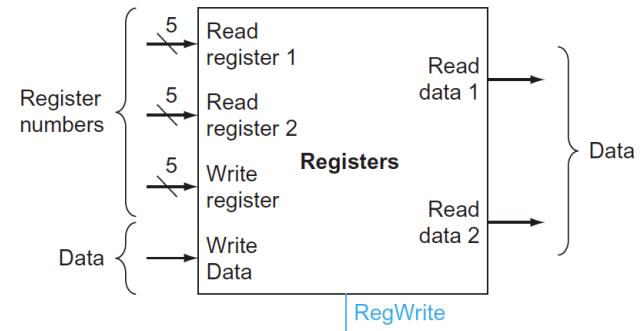
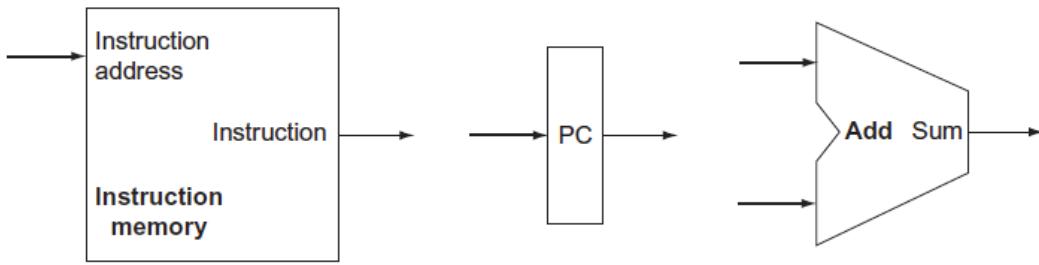
$$\begin{aligned} S = 0 &\rightarrow C = A \\ S = 1 &\rightarrow C = B \end{aligned}$$

Controller overview

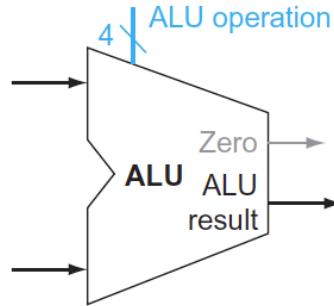


Building datapath

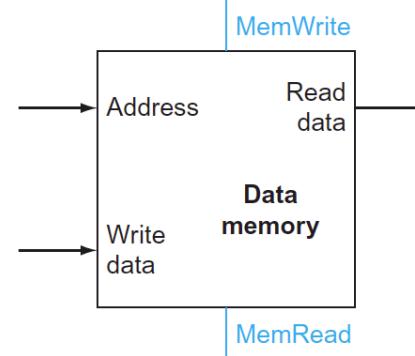
- Hardware components:



IF

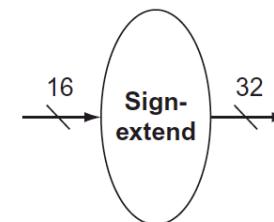


EXE



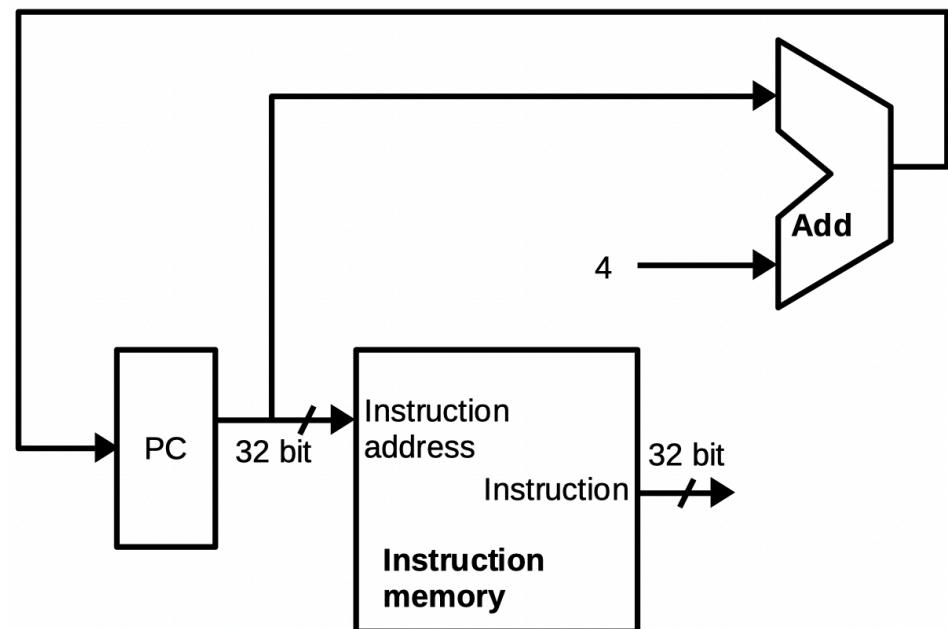
MEM

ID & WB



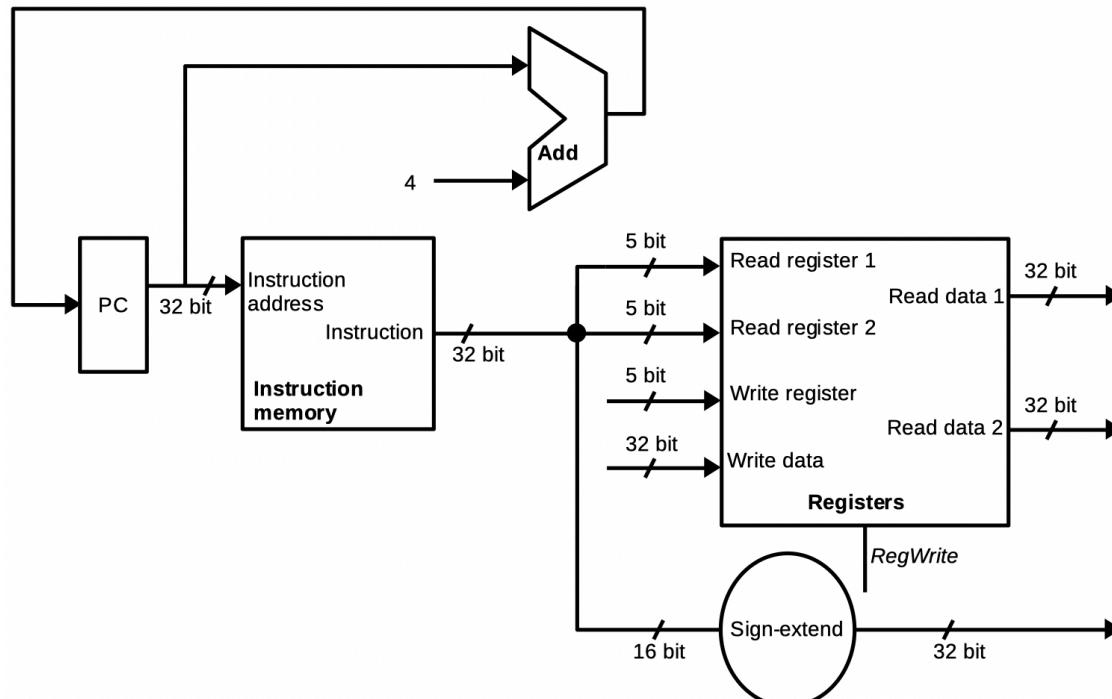
Instruction fetch

- **Main operations:** fetching the next instruction from Instruction memory
 - PC → instruction address
 - Instruction memory → instruction (32 bits)
 - $\text{PC} \leftarrow \text{PC} + 4$ (using the Add component)
- **Results:**
 - 32 bits machine instruction (Instruction)
 - Address of the next instruction in PC



Instruction decode

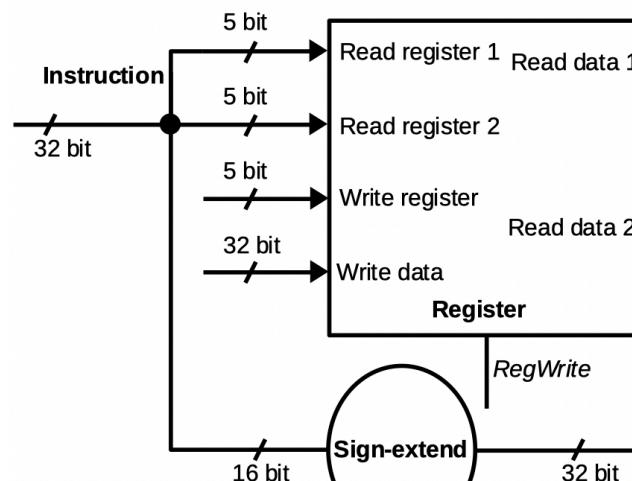
- **Main operations:** Extract machine instructions
 - R-format, branch, and store: $rs, rt \rightarrow$ Registers \rightarrow values
 - I-format: $rs \rightarrow$ Registers \rightarrow values & immediate \rightarrow sign-extend
- **Results:**
 - Values (32-bit) for the next stage



Instruction execution

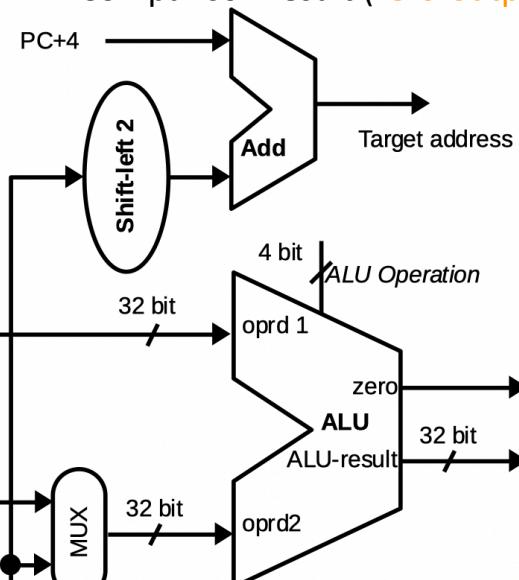
- Main operation

- Calculate the arithmetic operations/ address of memory/register comparison
 - R-format and **bne&beq**: both operands collected from **Registers**
 - I-format (except **bne&beq**): one operands collected from Registers and one from **sign-extend**



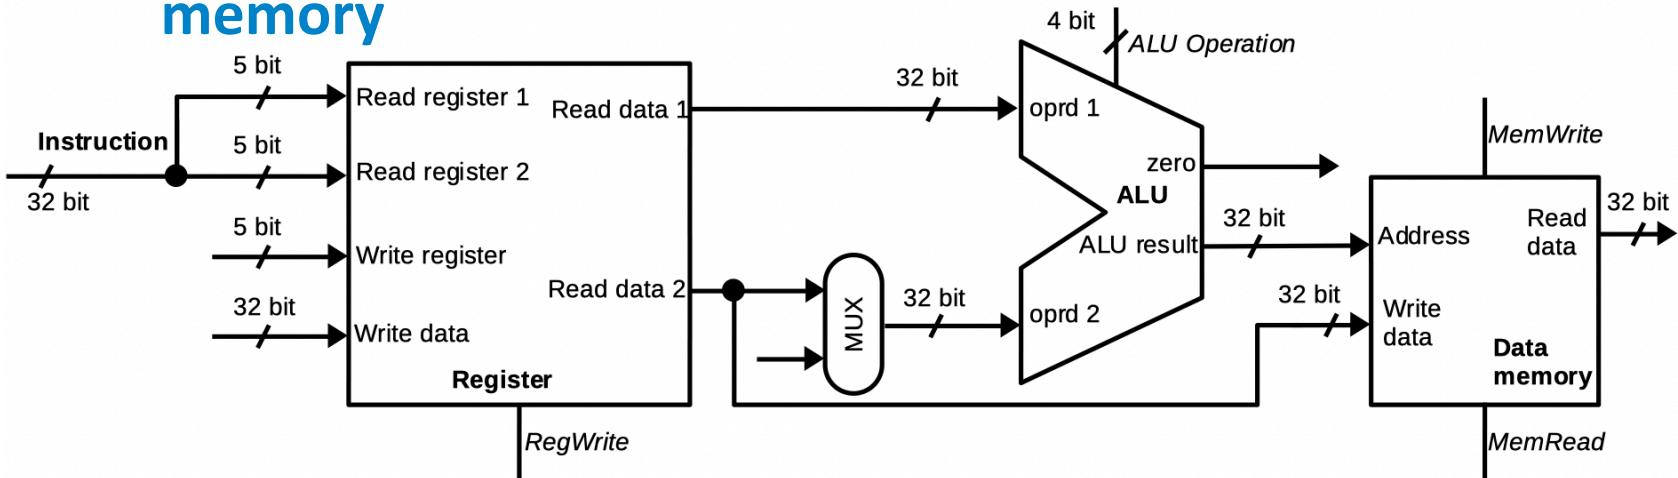
- Results:

- Arithmetic result/memory address (**ALU-result**)
- Comparison result (**zero-output**)



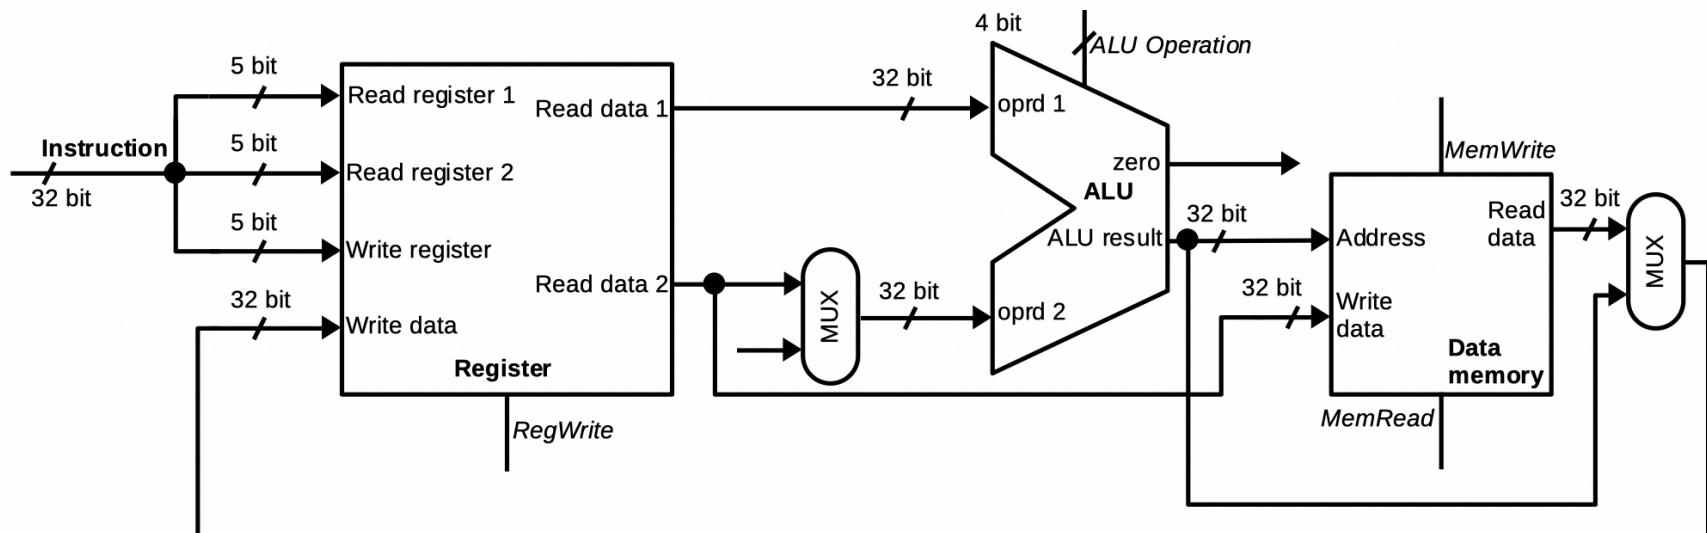
Memory access

- Main operations
 - Load: get data from **Data memory**
 - Store: write data from **Registers** to **Data memory**
- Results:
 - Load: values of **Data memory** (**Read data**)
 - Store: N/A

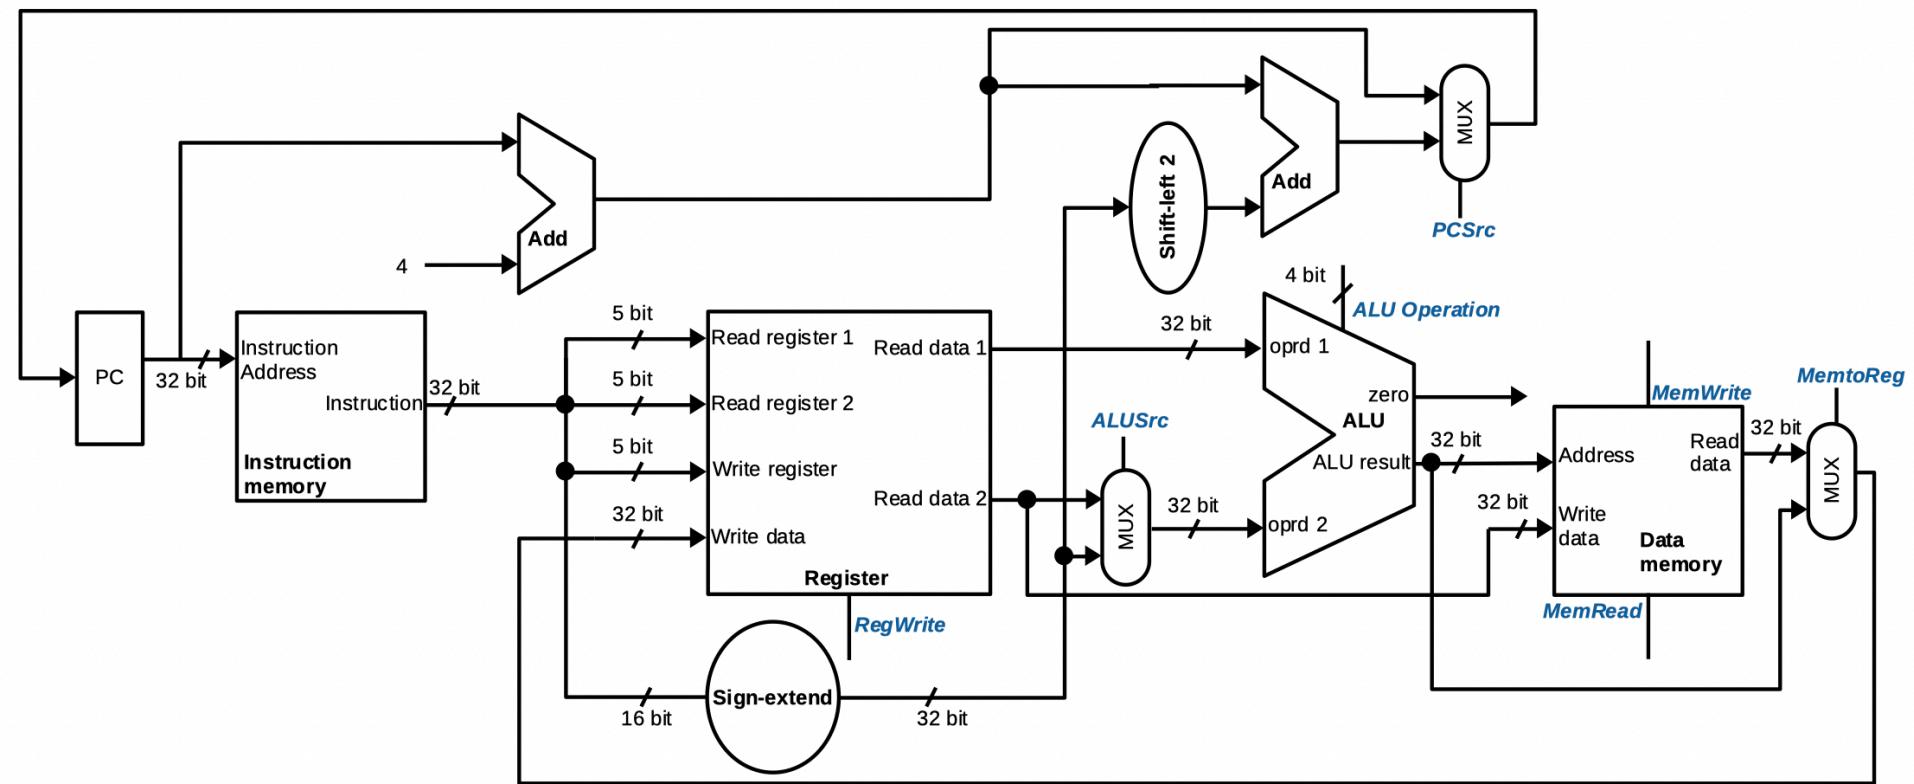


Write back

- Main operations:
 - Write values back to **Registers** (arithmetic/load)
- Result:
 - N/A



Full datapath



Example

- **Question:** Assume that the processor is executing

add \$s0, \$s1, \$s2

- Identify values of functional units' inputs/outputs
- **Answer:**

- The machine code is: 000000_10001_10010_10000_00000_100000₂
- Instruction memory:
 - **Instruction address** = PC (word address where we store the above instruction)
 - **Instruction** = machine code
- Registers:
 - **Read register 1** = 10001₂ => **Read data 1** = content (\$s1)
 - **Read register 2** = 10010₂ => **Read data 2** = content (\$s2)
 - **Write register** = 10000₂ & **Write data** = content (\$s1) + content (\$s2)
- ...sign-extend: input = 10000_00000_100000

Building controller

- Extracting bits from 32-bit instructions
 - Read register 1, Read register 2, and Write registers
 - Sign-extend
 - **Control** block
- Building a **Control** block:
 - Handling multiplexers
 - Handling control signals of functional units
 - *RegWrite*
 - *MemRead*
 - *MemWrite*
 - ...



Extracting bits

NOT update any Register: sw, beq

Update some Register: addi, lw, slli....

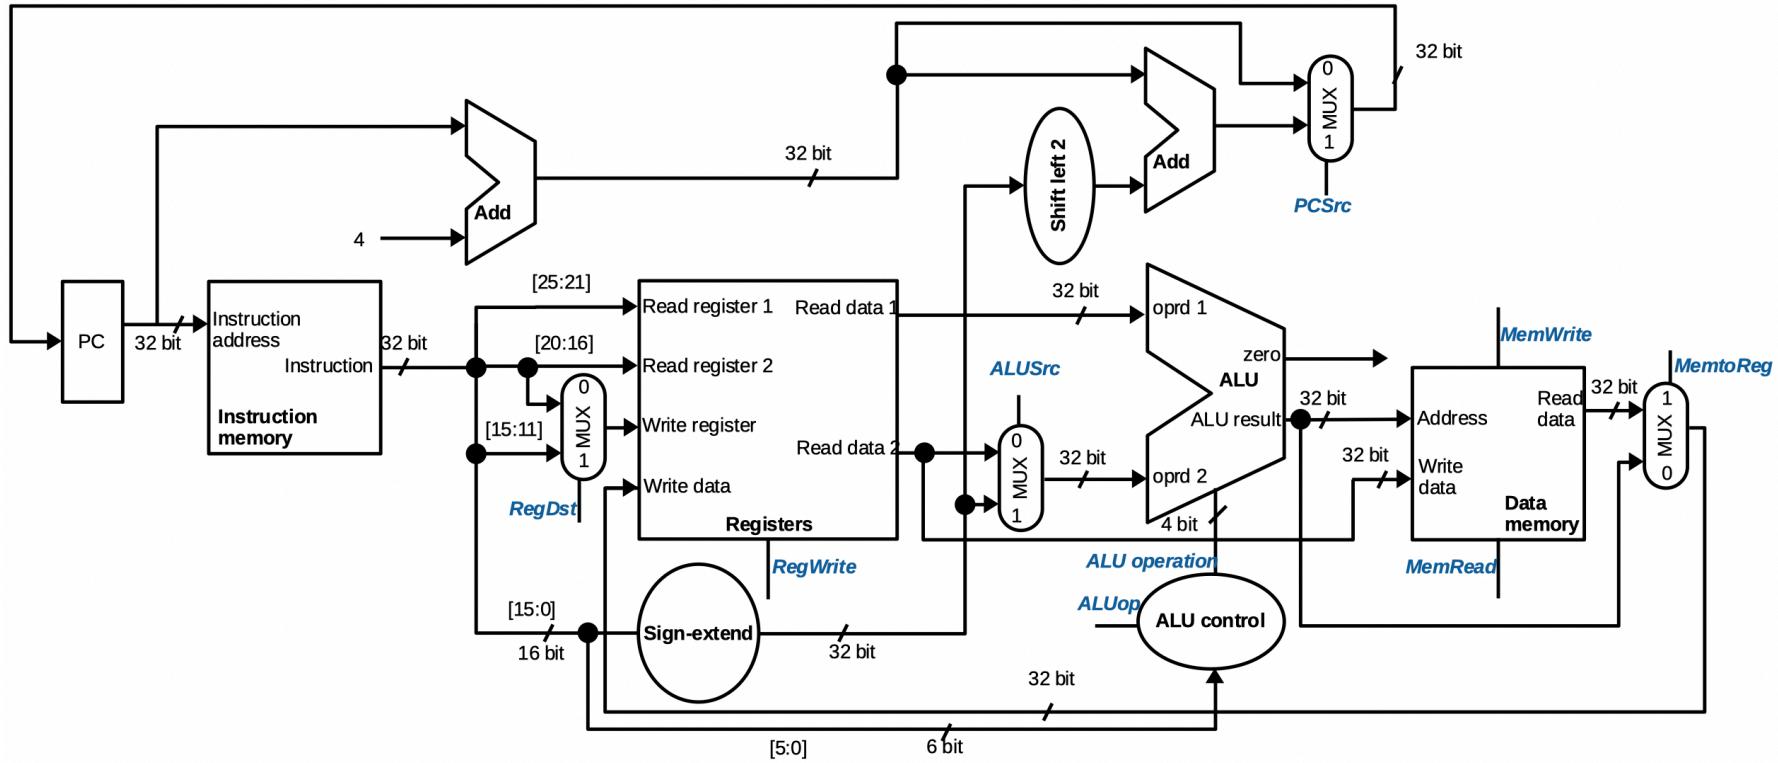
- **Registers:**

- Read register 1 \Leftarrow rs (instruction[25:21])
- Read register 2 \Leftarrow rt (instruction[20:16])
- Write register \Leftarrow rt/rd \rightarrow need a multiplexer

- **Sign-extend** \Leftarrow address (instruction[15:0])

R-type	0	rs	rt	rd	shamt	funct
	31:26	25:21	20:16	15:11	10:6	5:0
Load/ Store	35 or 43	rs	rt	address		
	31:26	25:21	20:16		15:0	
Branch	4	rs	rt	address		
	31:26	25:21	20:16		15:0	

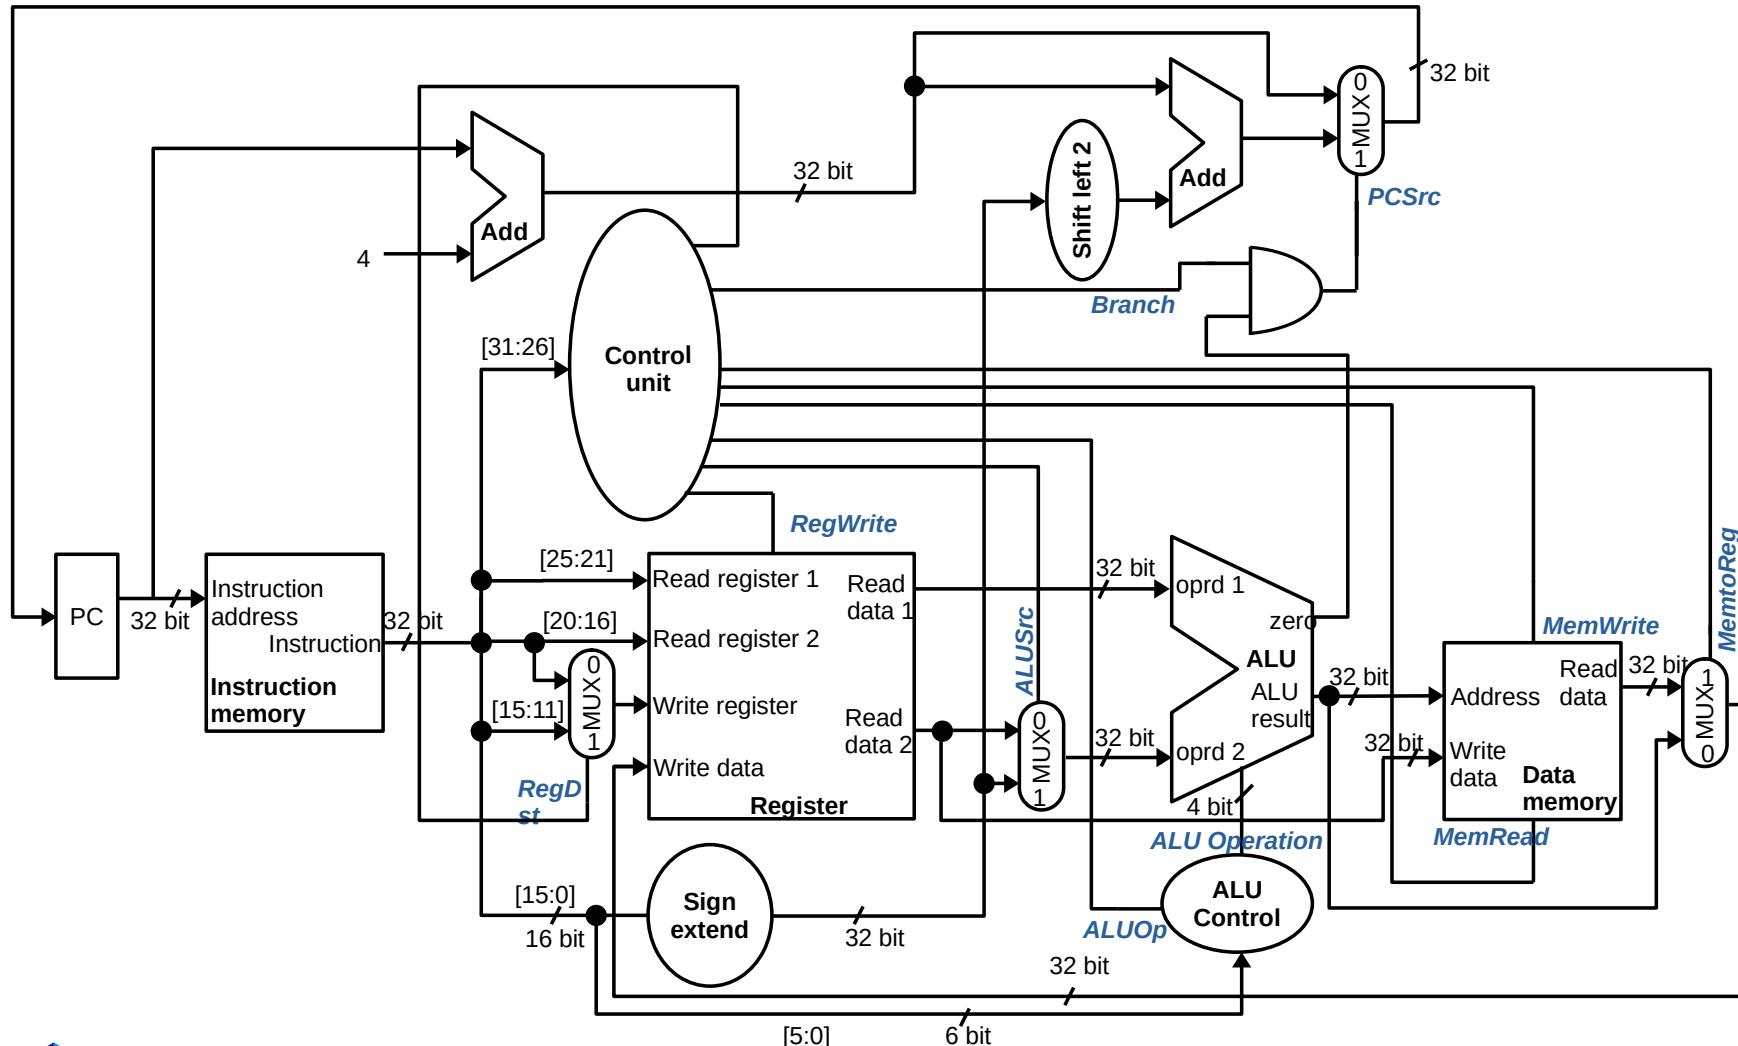
Datapath with bit-selection



Control block

- Main function: handling multiplexers and functional units' control signals
 - ALU: two levels of decoding
 - Level 1: 6 bit **opcode** → 2 bit **ALUop**
 - Level 2: 2 bit **ALUop** (+ 6 bit function field) → 4 bit **ALU operation**
 - Muxes and control signals of functional units (except ALU): 6 bit **opcode**
 - Predefined **ALU operation** values
- | Operator | ALU Opeation |
|----------|--------------|
| and | 0000 |
| or | 0001 |
| add | 0010 |
| sub | 0110 |
| slt | 0111 |
- **ALUop**:
 - opcode → add operator → **ALUop** = 00
 - opcode → sub operator → **ALUop** = 01
 - opcode → unknown → **ALUop** = 10

Full microarchitecture



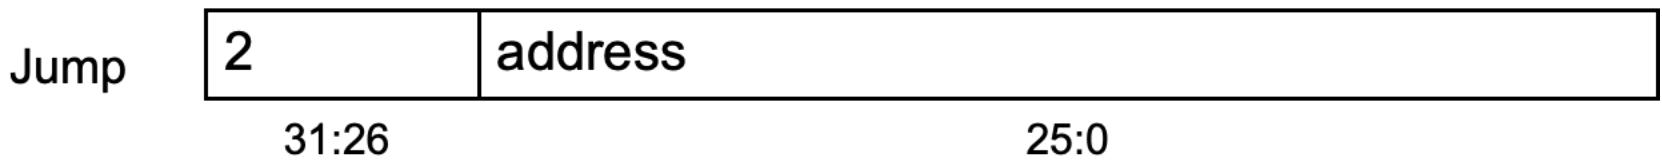
Exercise

- **Question:** assume that registers store values of two times their numbers, for instance \$s1 stores values of $17 \times 2 = 34$, please identify:
 1. Which functional units contribute to the processing of following instructions
 2. Values of inputs/outputs of functional units when processing following instructions
 3. Values of control signals when processing following instructions
 - add \$t0, \$t1, \$t2
 - addi \$s0, \$s1, 100
 - lw \$s0, 100(\$s2) # memory word at address 136 stores values of 2025
 - sw \$s0, 100(\$s2)
 - beq \$s1, \$s0, L1

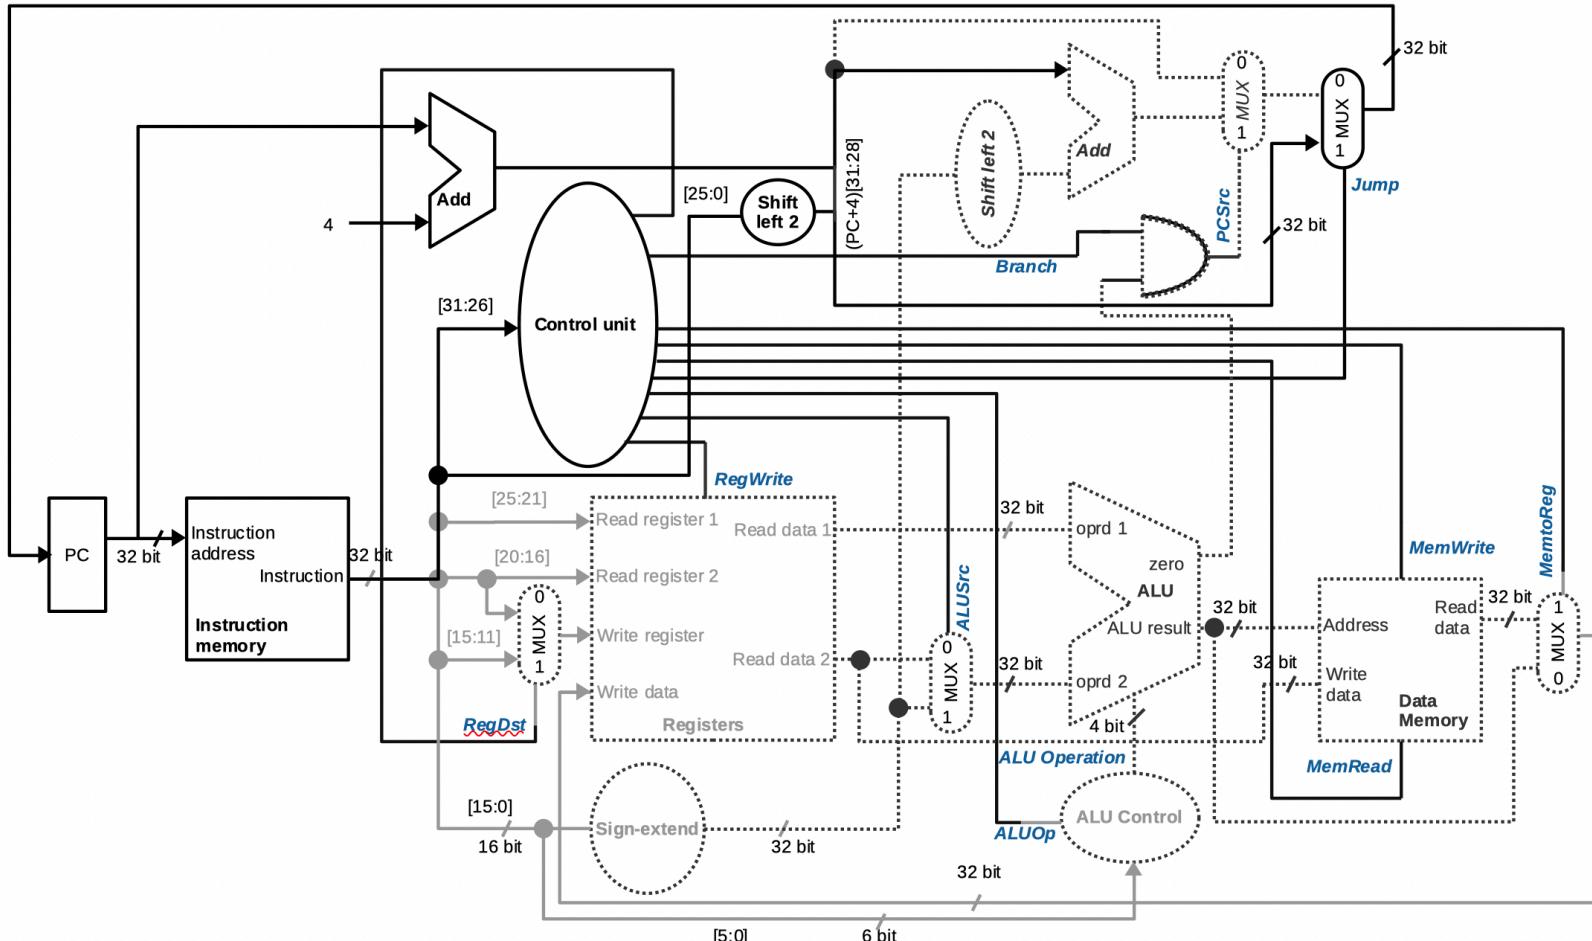


Unconditional jump instructions

- Chapter 2: update **PC** with concatenation of
 - Top 4 bits of old **PC**
 - 26-bit jump address
 - 00
- One more way to update **PC** => Need a multiplexer & an extra control signal decoded from opcode



Datapath with Jumps added



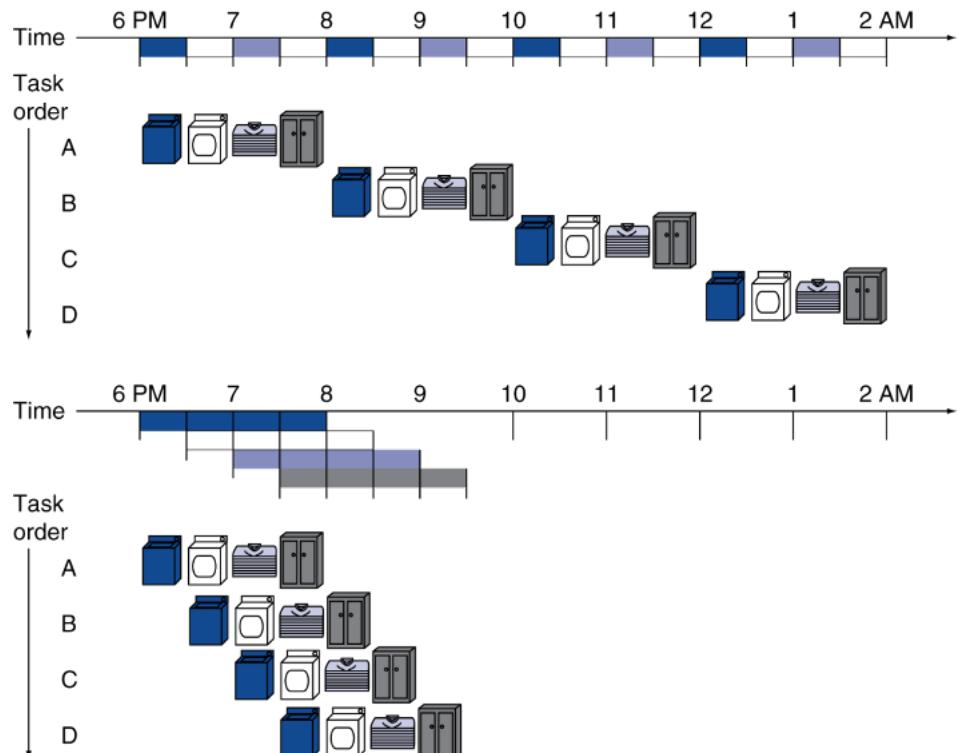
Performance issue

- **Simplified version** ($CPI = 1$): every instruction executed in **only one cycle**
 - **Longest delay** determines clock period
 - What is the longest instruction?
 - Instruction memory → register file → ALU → data memory → register file
- Not feasible to vary period for different instructions
- We will improve performance by pipelining



Pipelining analogy

- Pipeline laundry:
 - Overlapping execution
 - Improving performance (time for entire group)
- Four loads
 - Speed-up = 2.3×
 - Not impressive
- Non-stop ($\# \text{loads} \rightarrow \infty$)
 - Speed-up?
 - Number of stages



MIPS pipeline

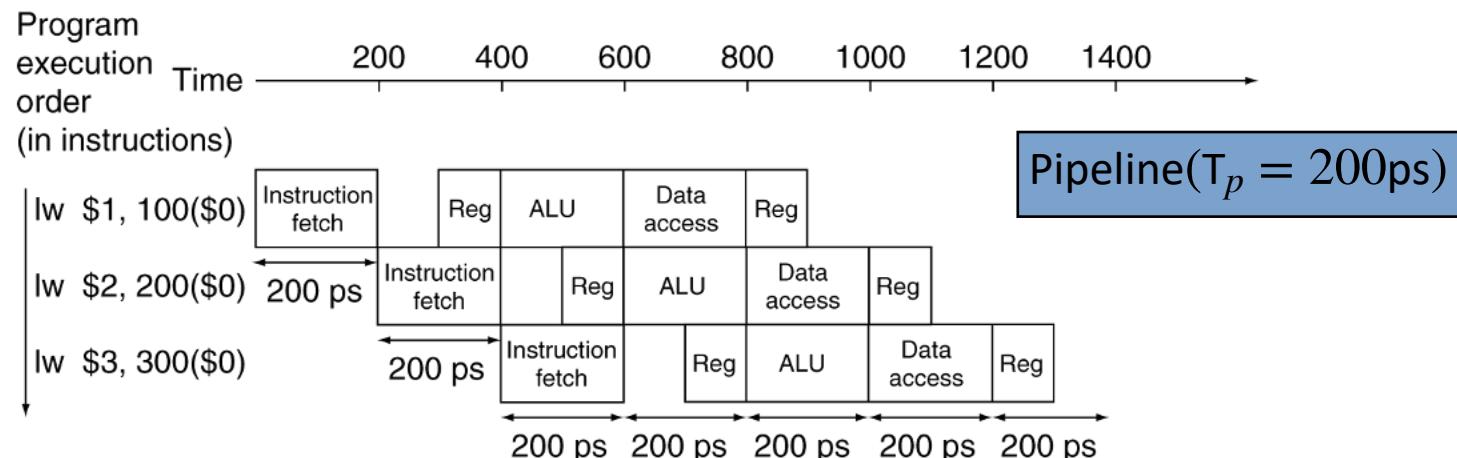
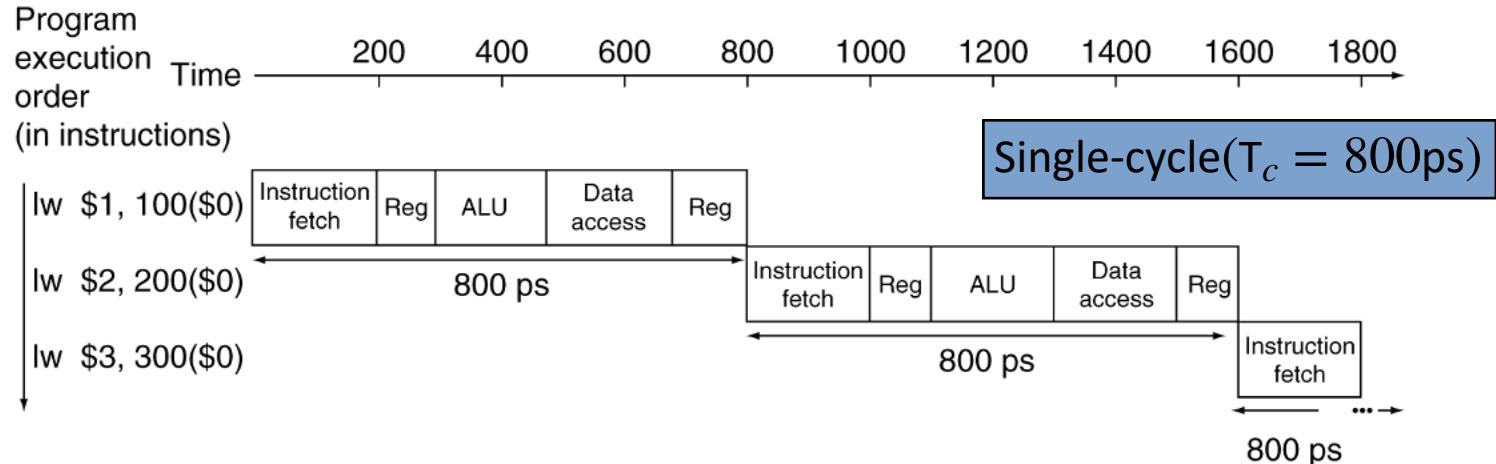
- **Five stages**, one step per stage
 1. Instruction fetch (**IF**): PC → instruction address
 2. Instruction decode (**ID**): register operands → register file
 3. Execute (**EXE**):
 - Load/store: compute a memory address
 - Arithmetic/logical: compute an arithmetic/logical result
 4. Memory access (**MEM**):
 - Load: read data memory
 - Store: write data memory
 5. Write back (**WB**):
 - Store a result of register file

Pipeline performance

- Assume time for stages is
 - 100ps for register read or write (**ID** & **WB**)
 - 200ps for other stages (**IF**, **EXE**, & **MEM**)
- Compare **pipelined datapath** with **single-cycle datapath**

Instruction	Instr fetch	Register read	ALU op	Memory access	Register write	Total time
lw	200ps	100 ps	200ps	200ps	100 ps	800ps
sw	200ps	100 ps	200ps	200ps		700ps
R-format	200ps	100 ps	200ps		100 ps	600ps
beq	200ps	100 ps	200ps			500ps

Pipeline performance

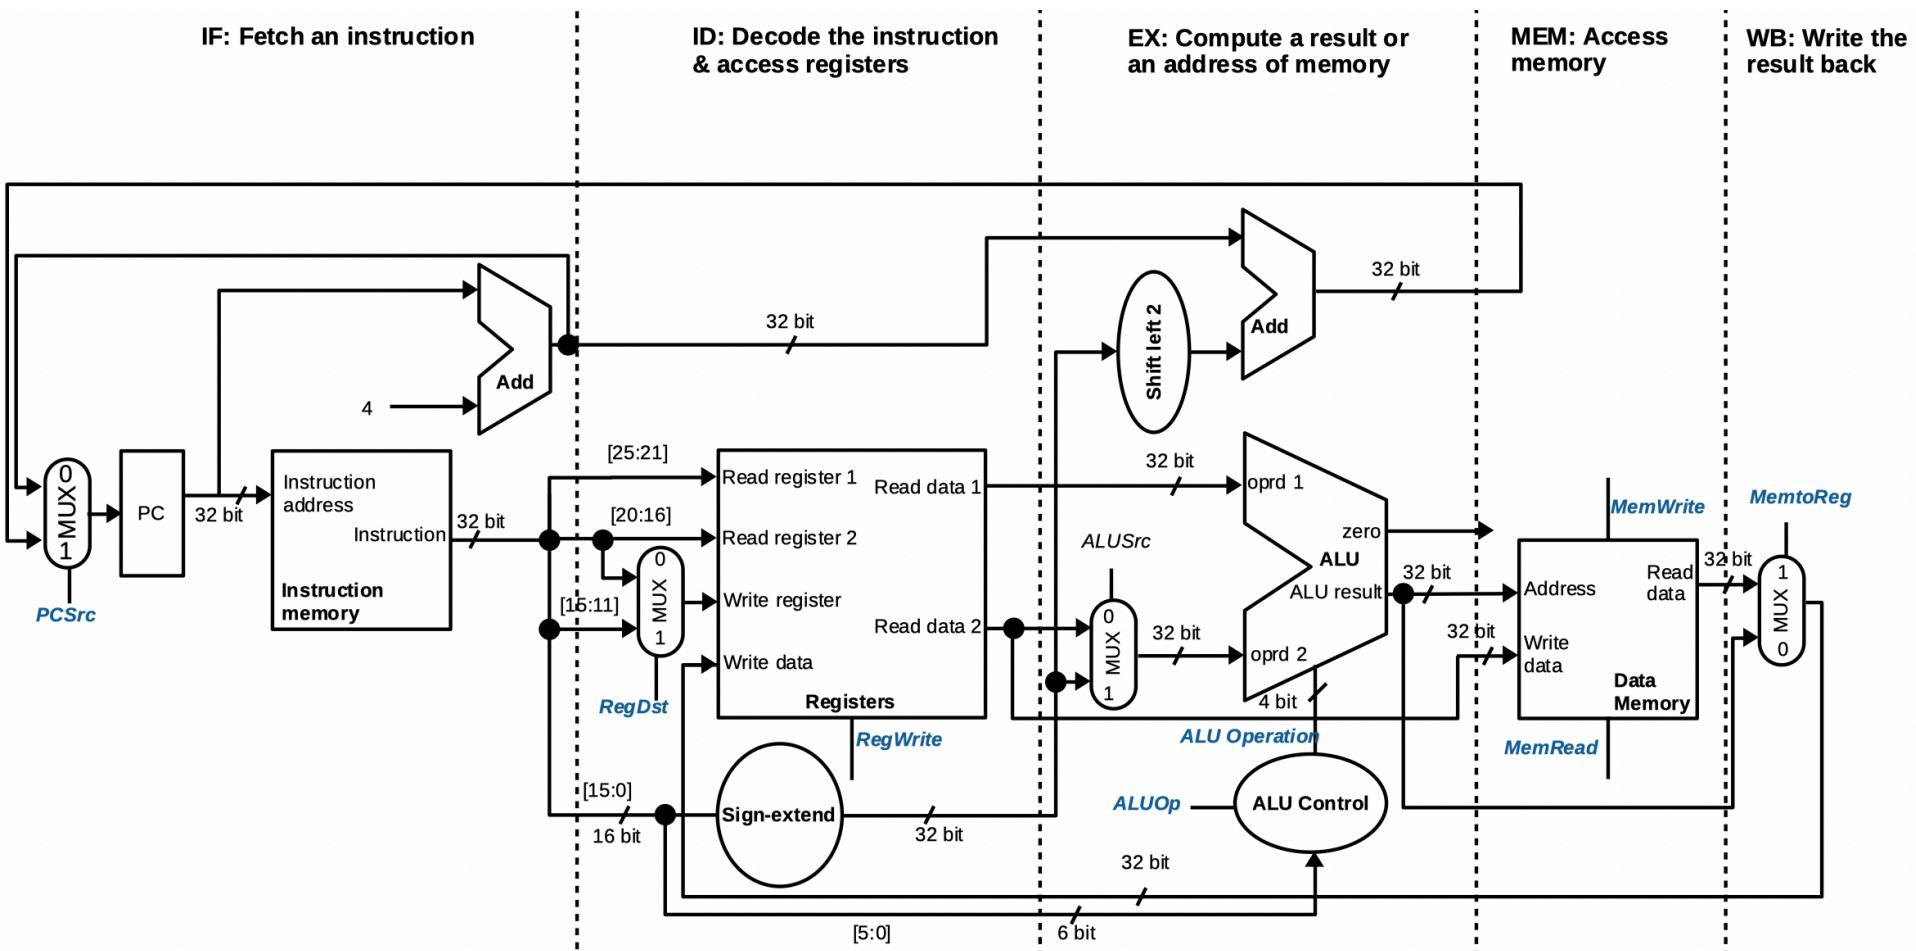


Pipeline speed-up

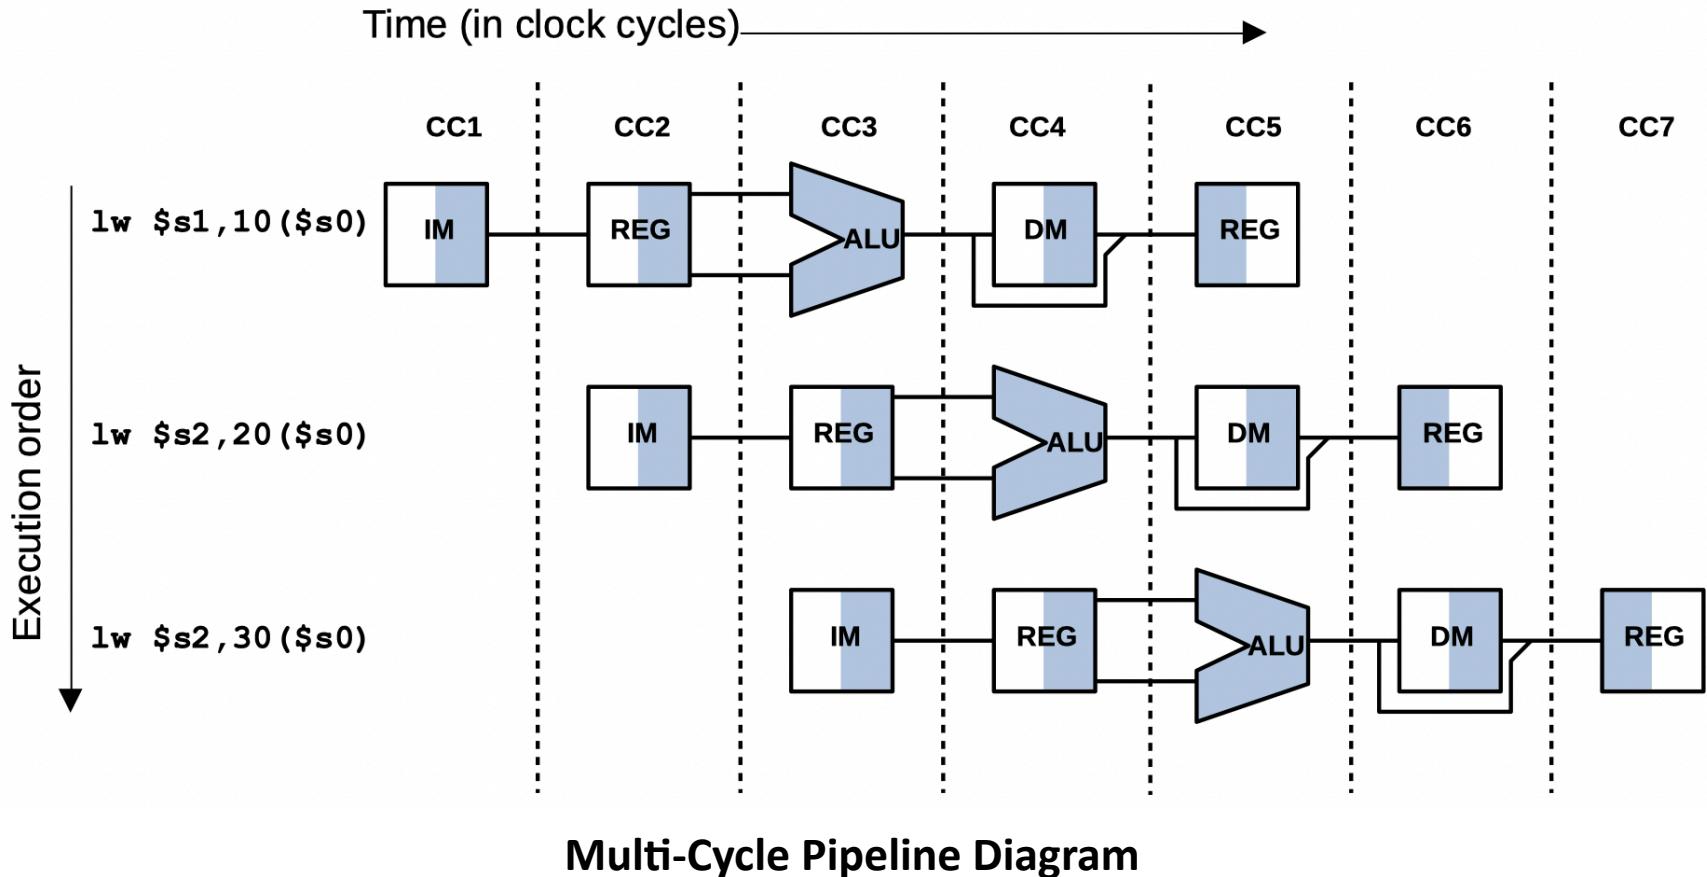
$$\text{speed-up} = \frac{\text{time of single-cycle}}{\text{time of pipelined}} = \frac{\text{time bw instruction}_{\text{single-cycle}}}{\text{time bw instruction}_{\text{pipelined}}}$$

- If all stages are **balanced**:
 - $\text{speed_up} = \text{number of pipe stages}$
- If **not balanced**, speedup is less
- Source of speedup
 - **Throughput** increased
 - **Latency** (time for each instruction) does not decrease
 - Sometimes **increased**

Pipeline datapath

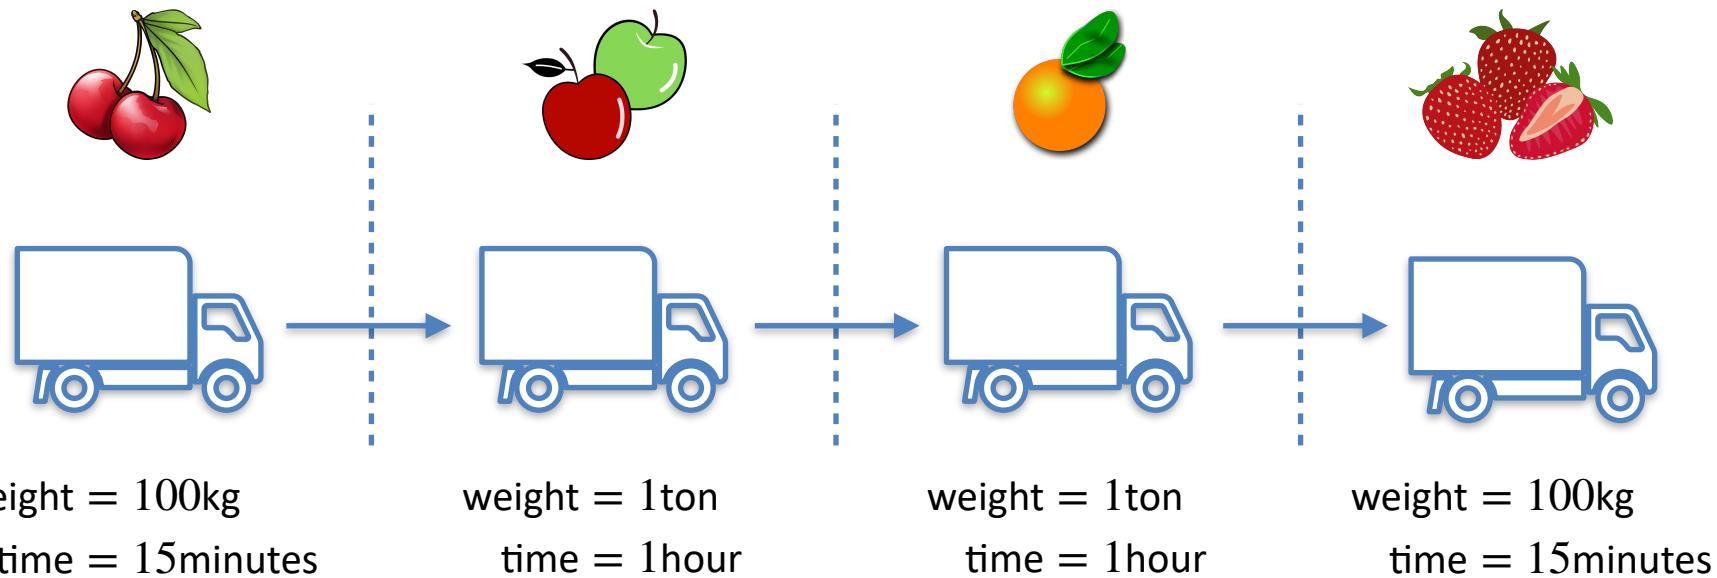


Instructions execution



How can we keep data when stages are not balanced?

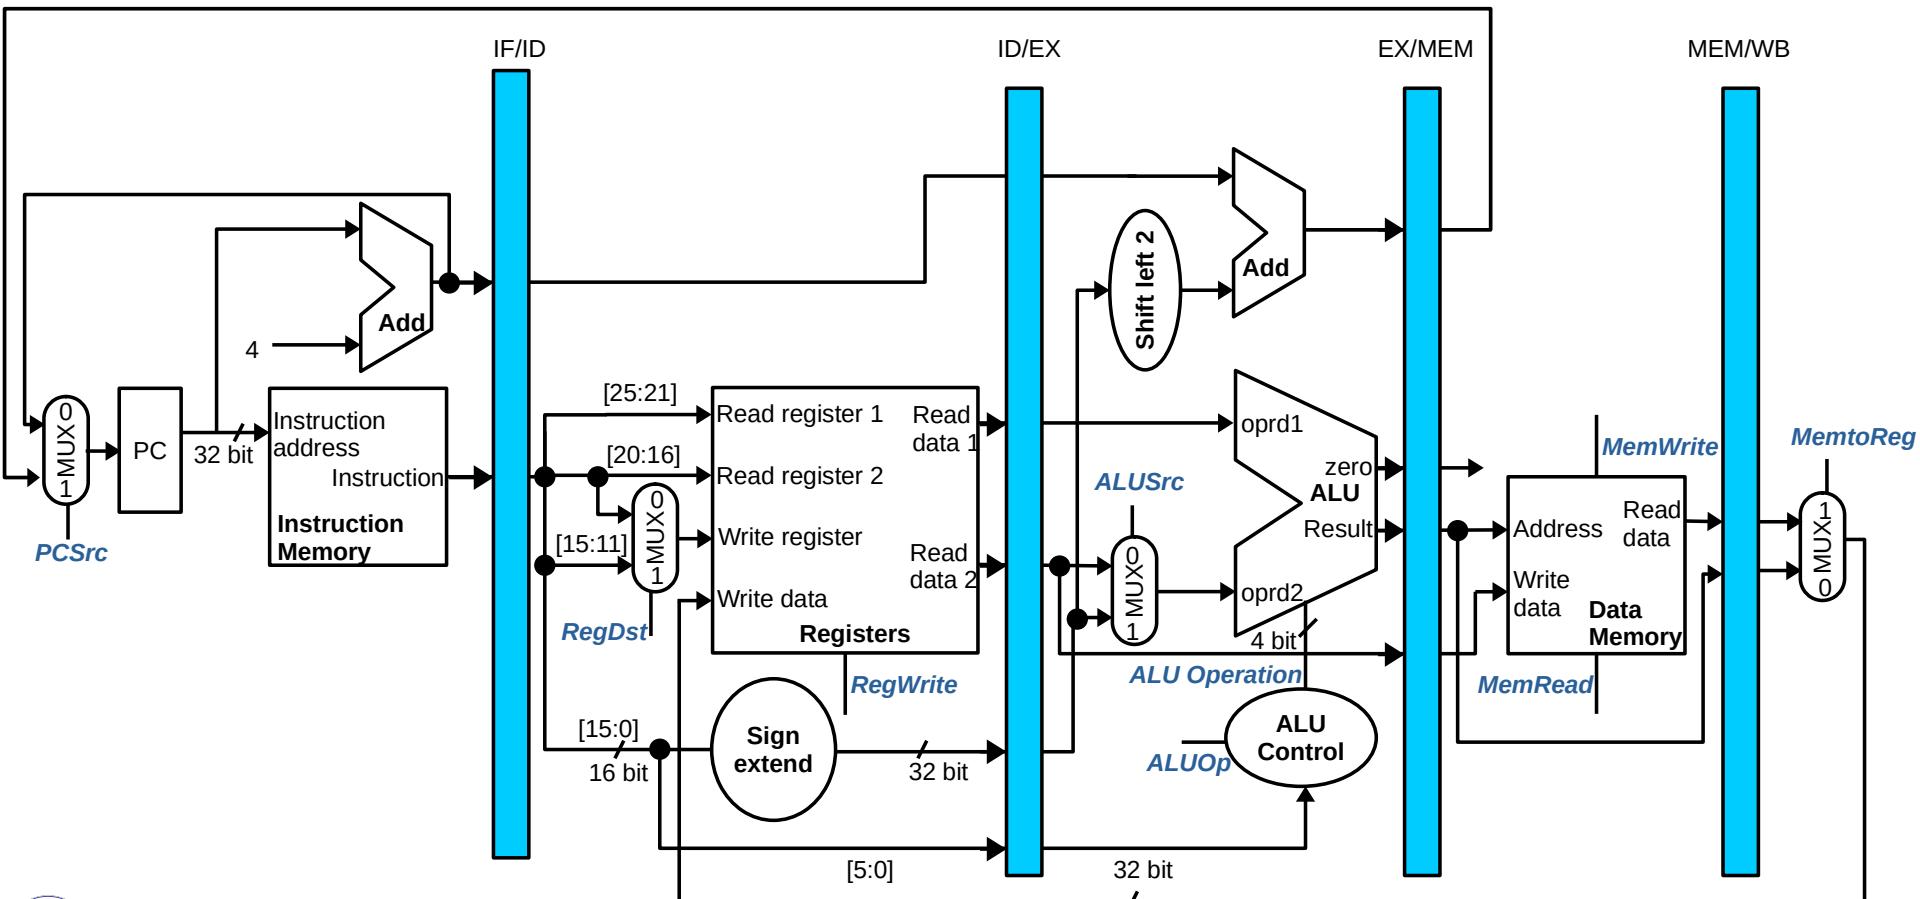
Wholesale market example



- Trucks move forward when completed
 - **Accident** at the Apple store
 - **Barriers** can help: open every hour (**cycle**)

Pipeline registers

Barriers in wholesale markets \Leftrightarrow Registers in digital circuits



Example

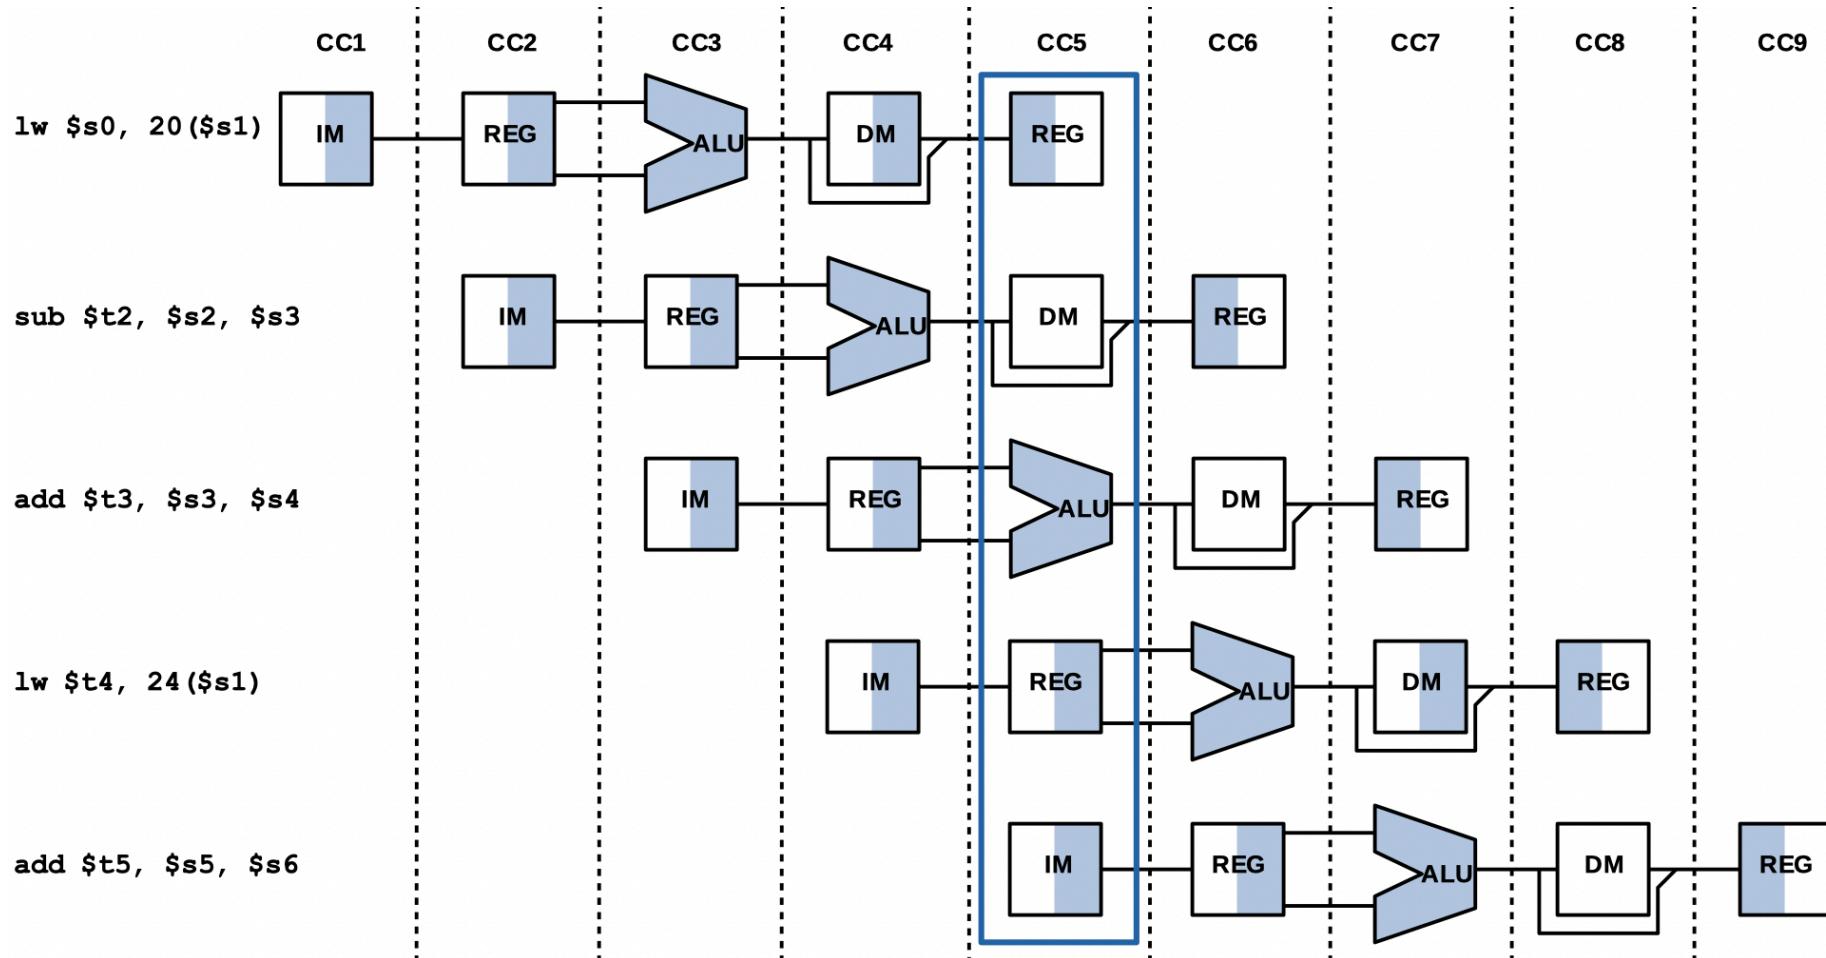
- **Question:** Given the following MIPS sequence:

1. lw \$s0, 20(\$s1)
2. sub \$t2, \$s2, \$s3
3. add \$t3, \$s3, \$s4
4. lw \$t4, 24(\$s1)
5. add \$t5, \$s5, \$s6

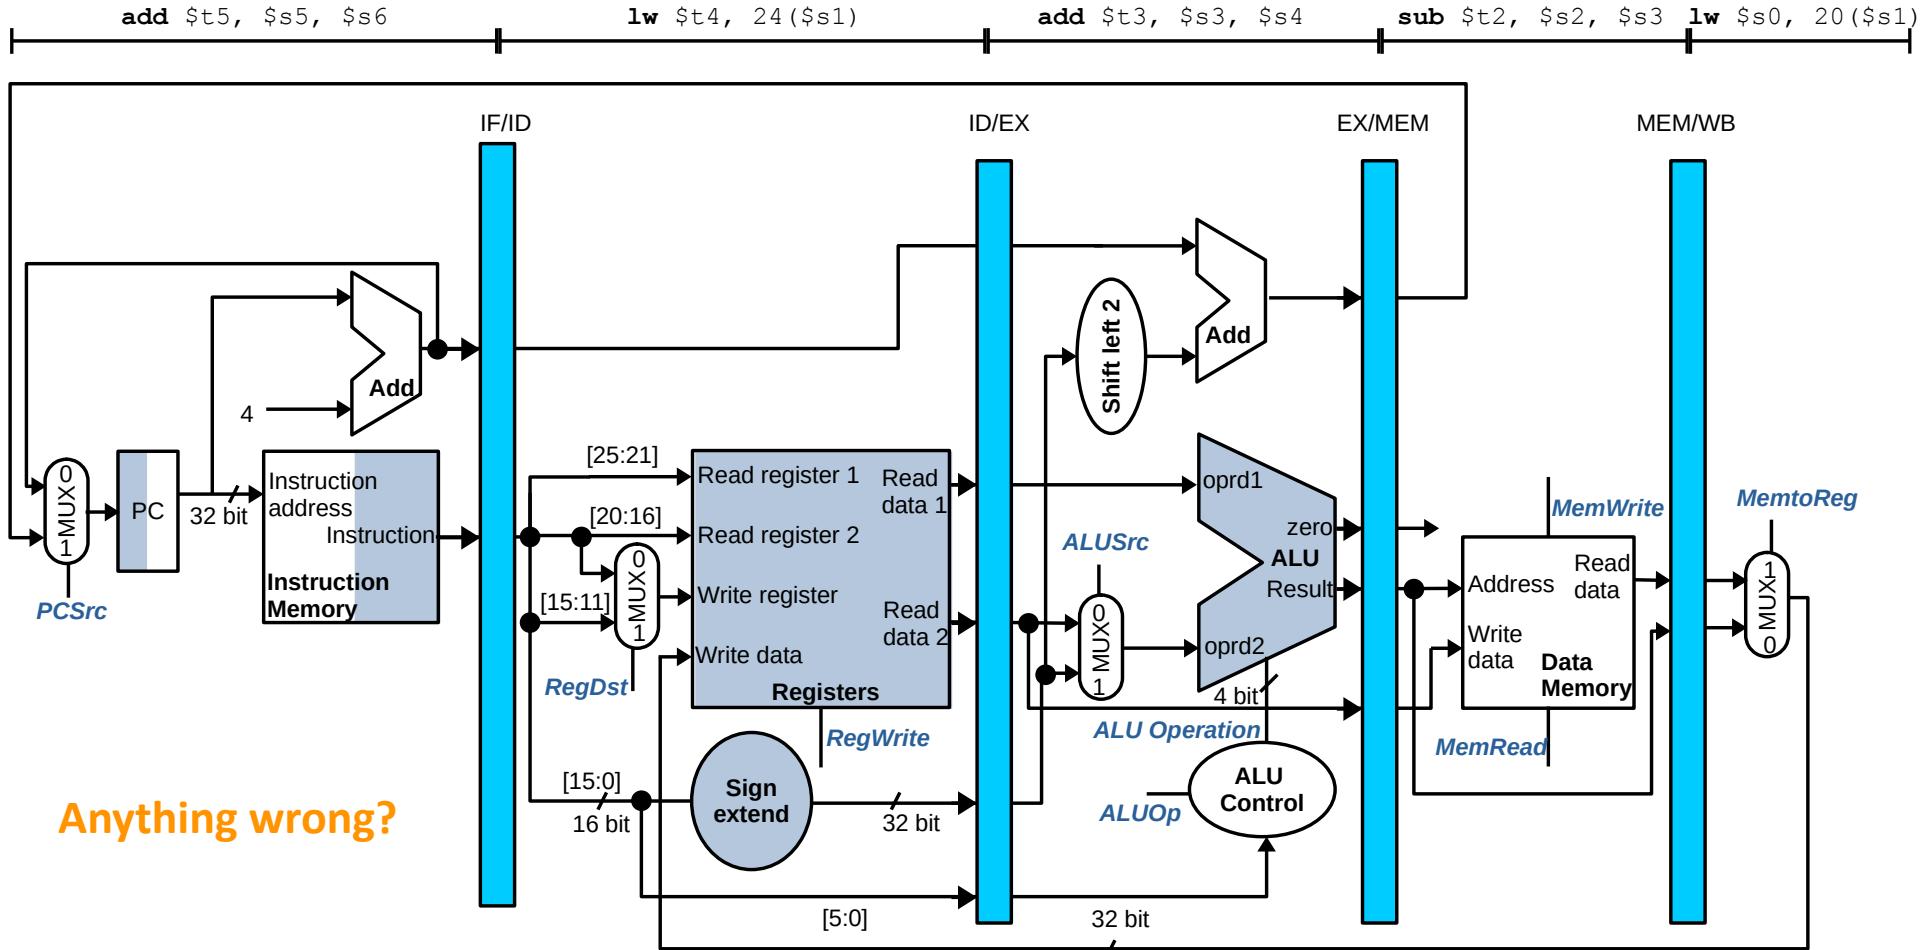
Assume that the sequence is executed by a 5-stage pipelined MIPS processor

- a) Draw a multi-cycle pipeline diagram for the sequence
- b) Analyze the 5th cycle with the datapath diagram in the previous slide

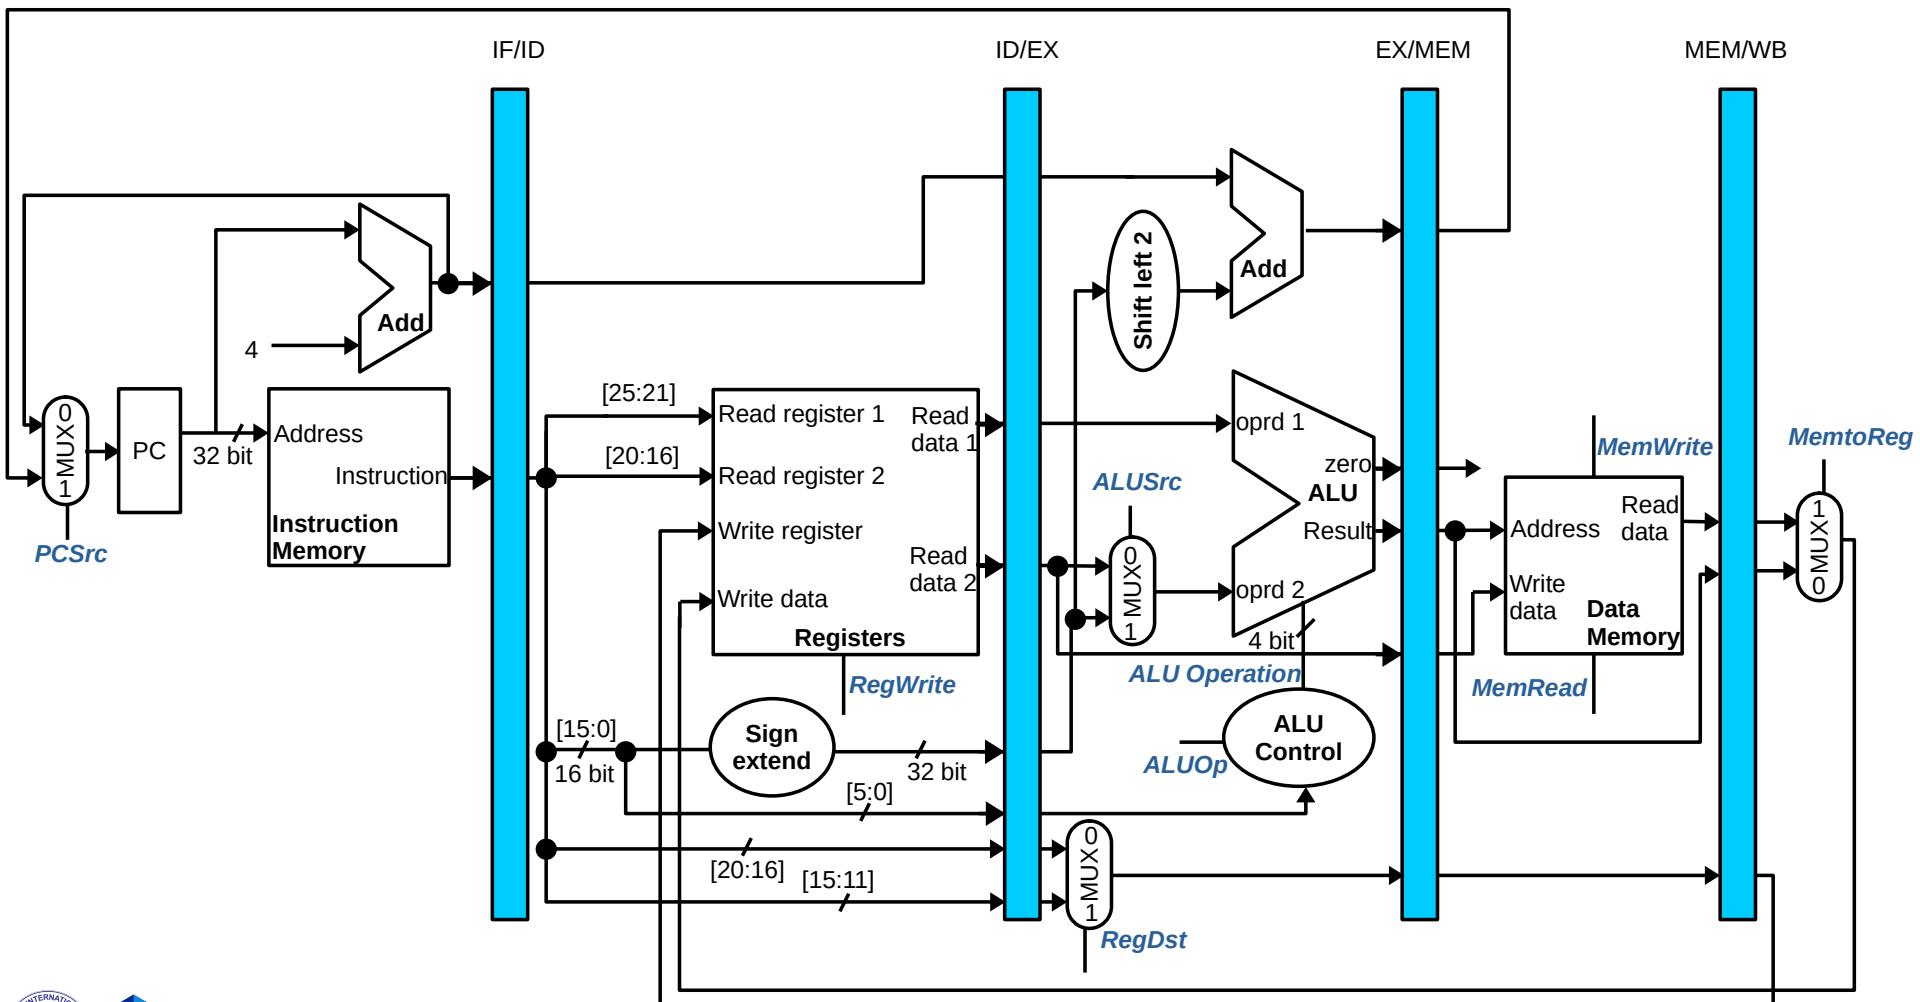
Multi-cycle pipeline diagram



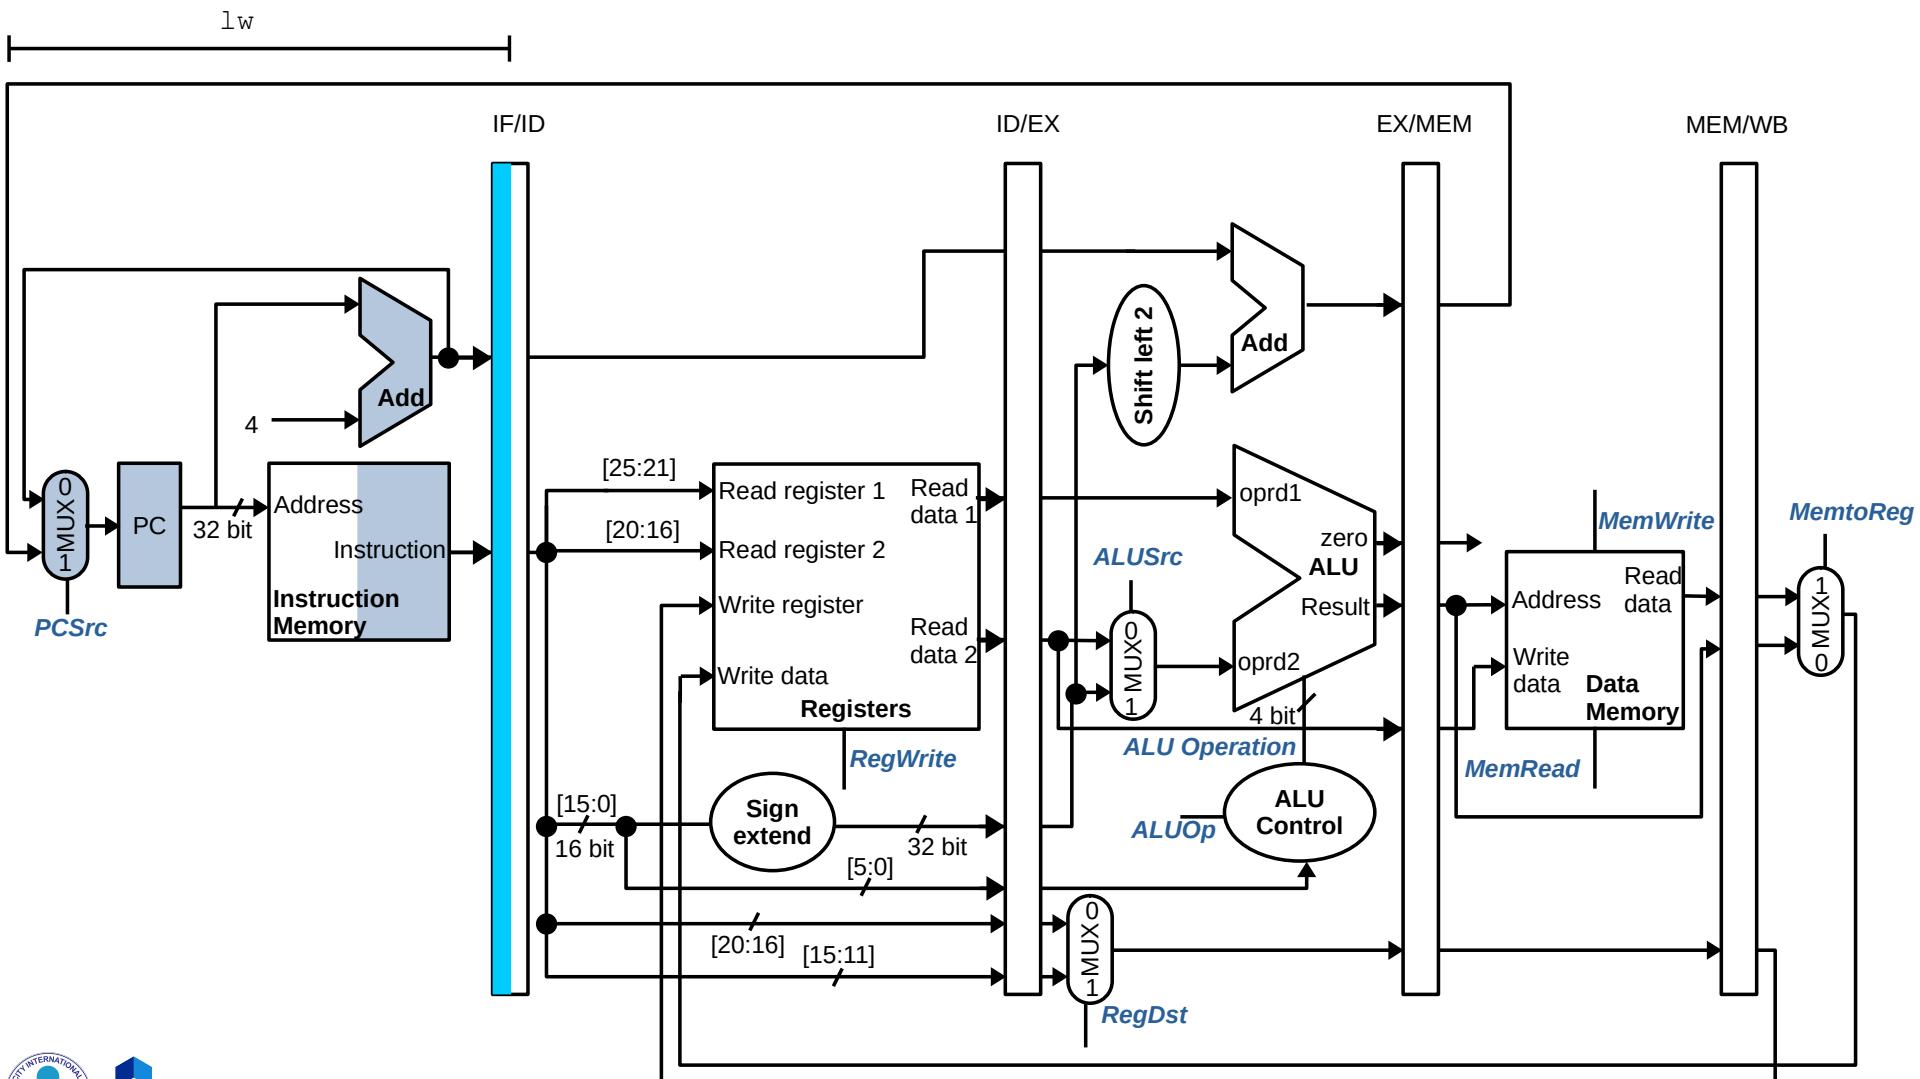
Single-cycle pipeline diagram



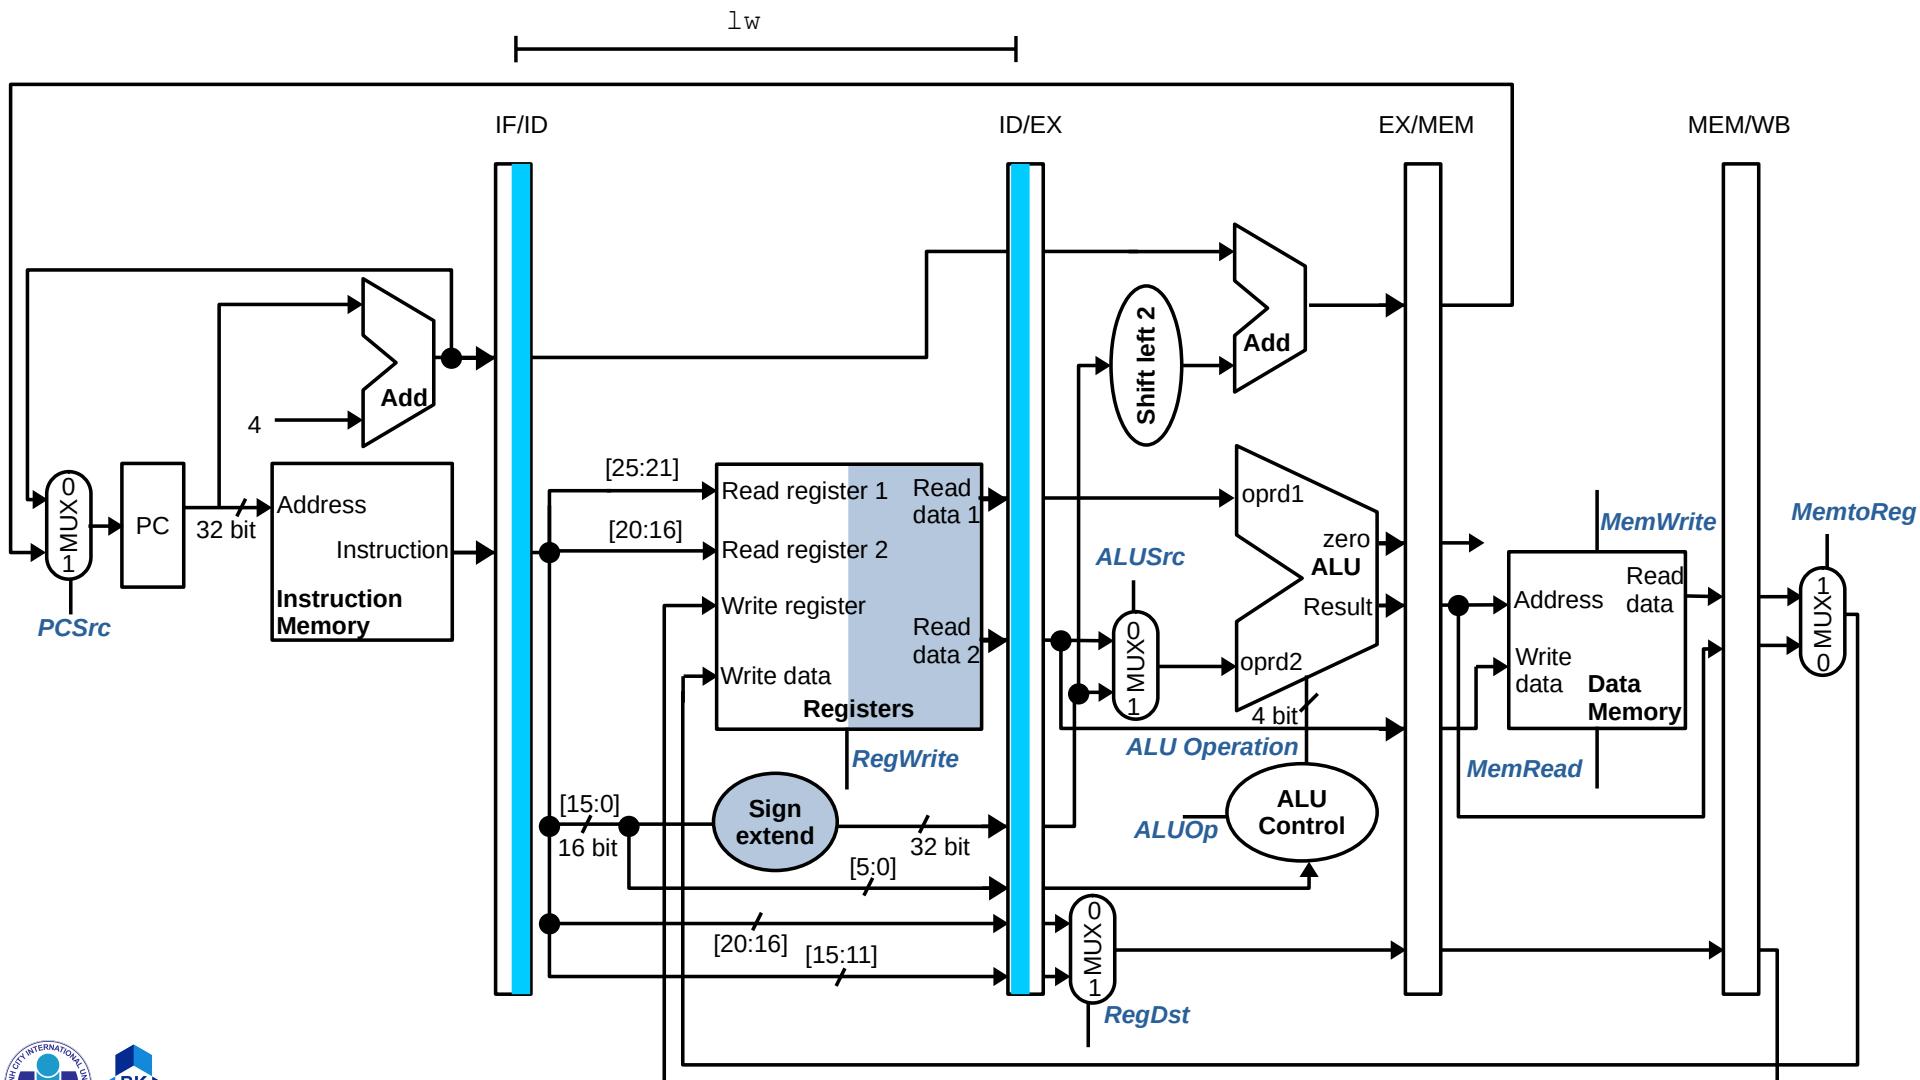
Corrected pipeline datapath



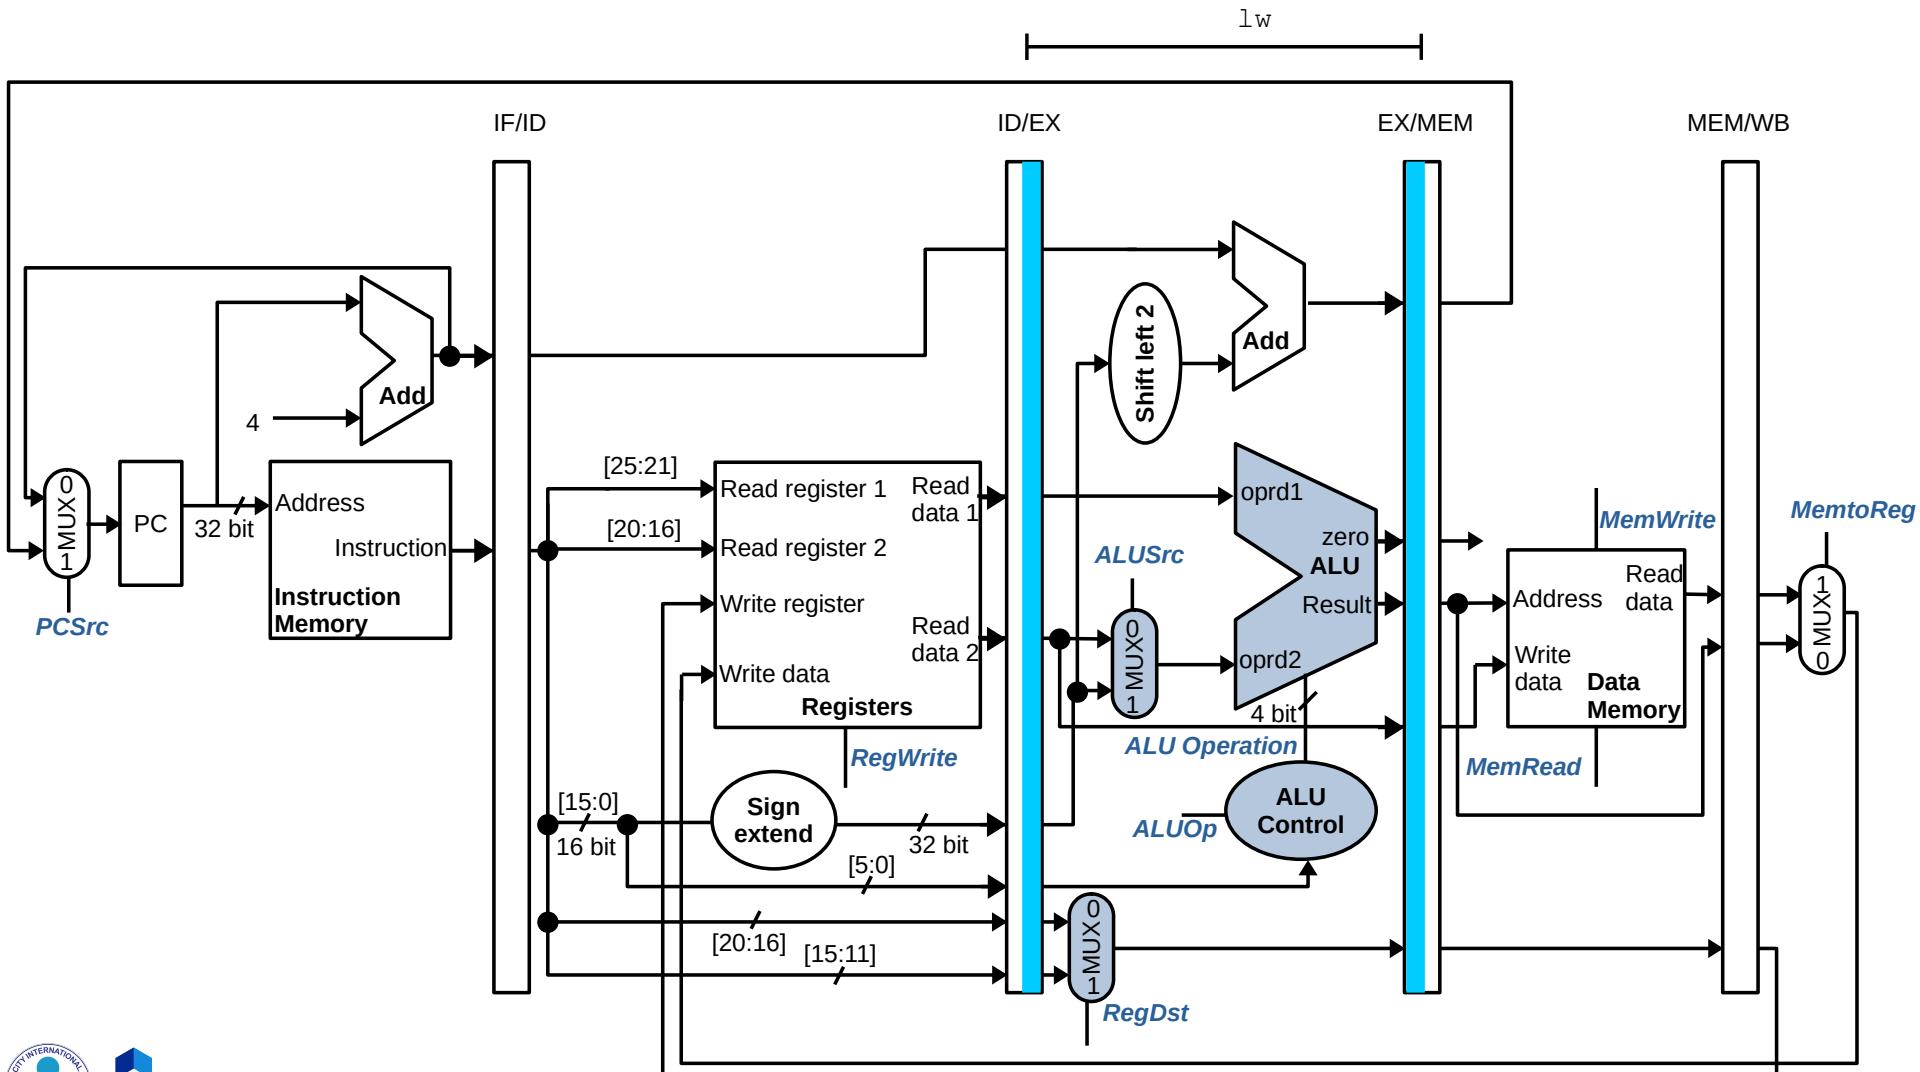
Example of lw - IF



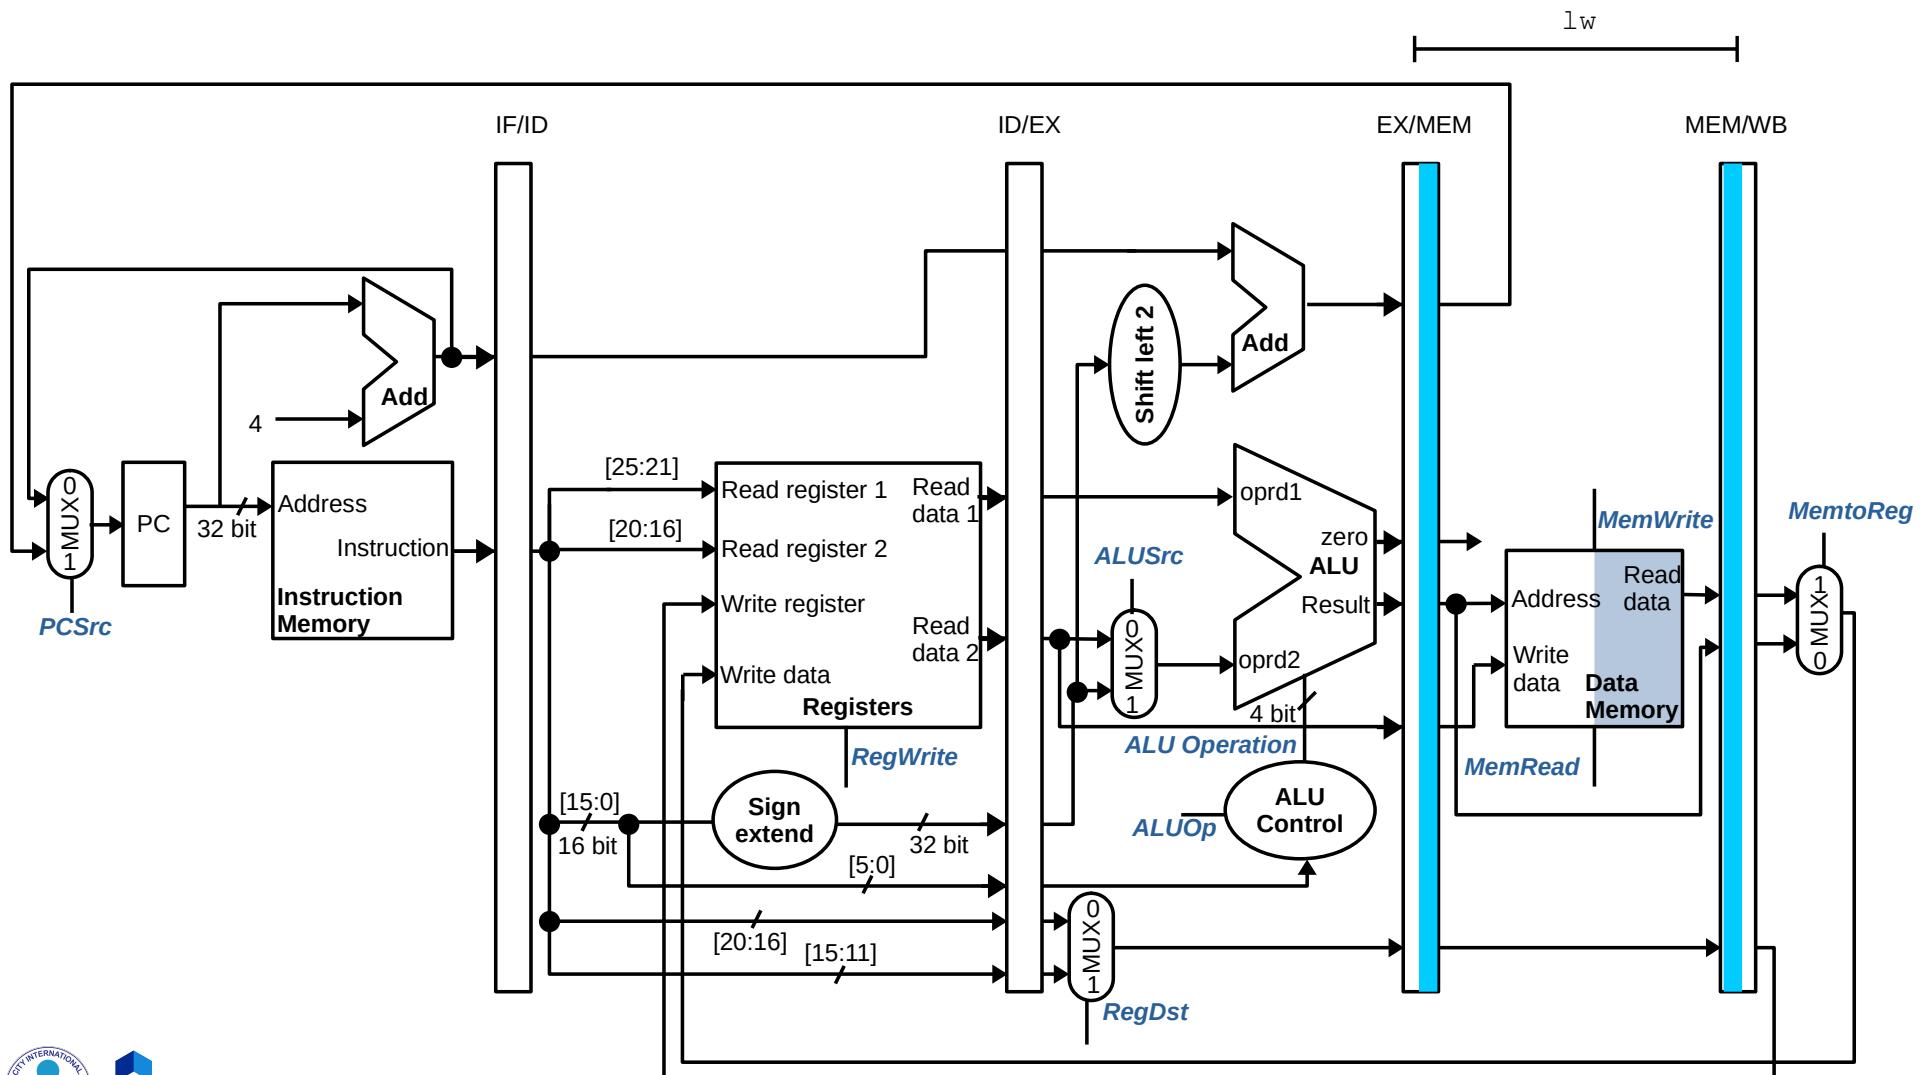
Example of lw - ID



Example of lw - EXE

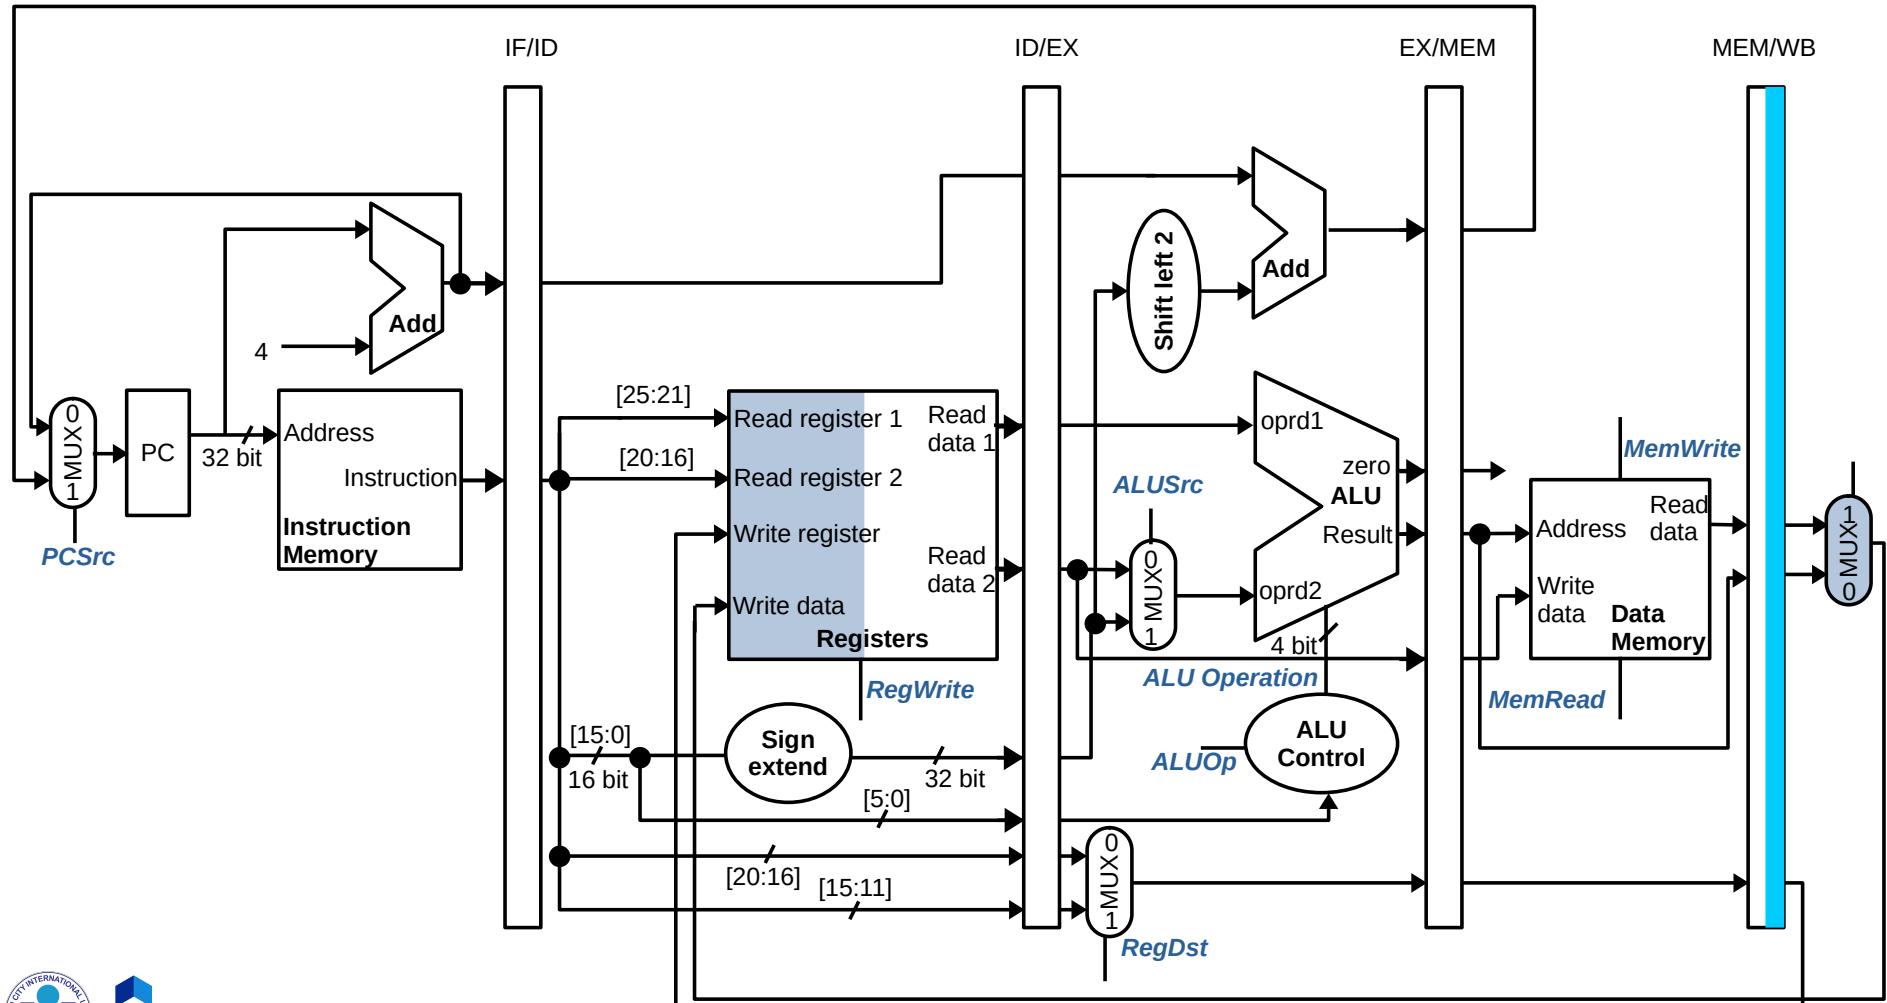


Example of lw - MEM



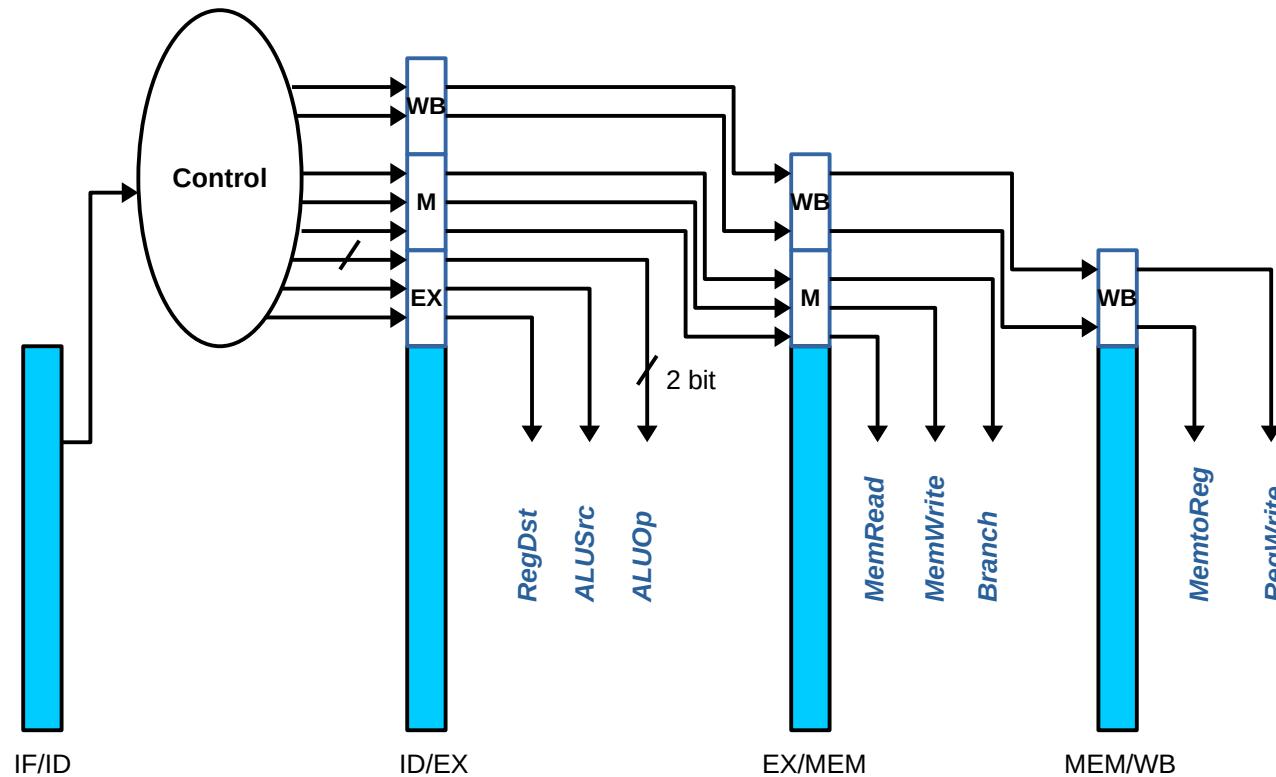
Example of lw - WB

lw



Control signals

- Control signals travel across pipeline registers



Exercise

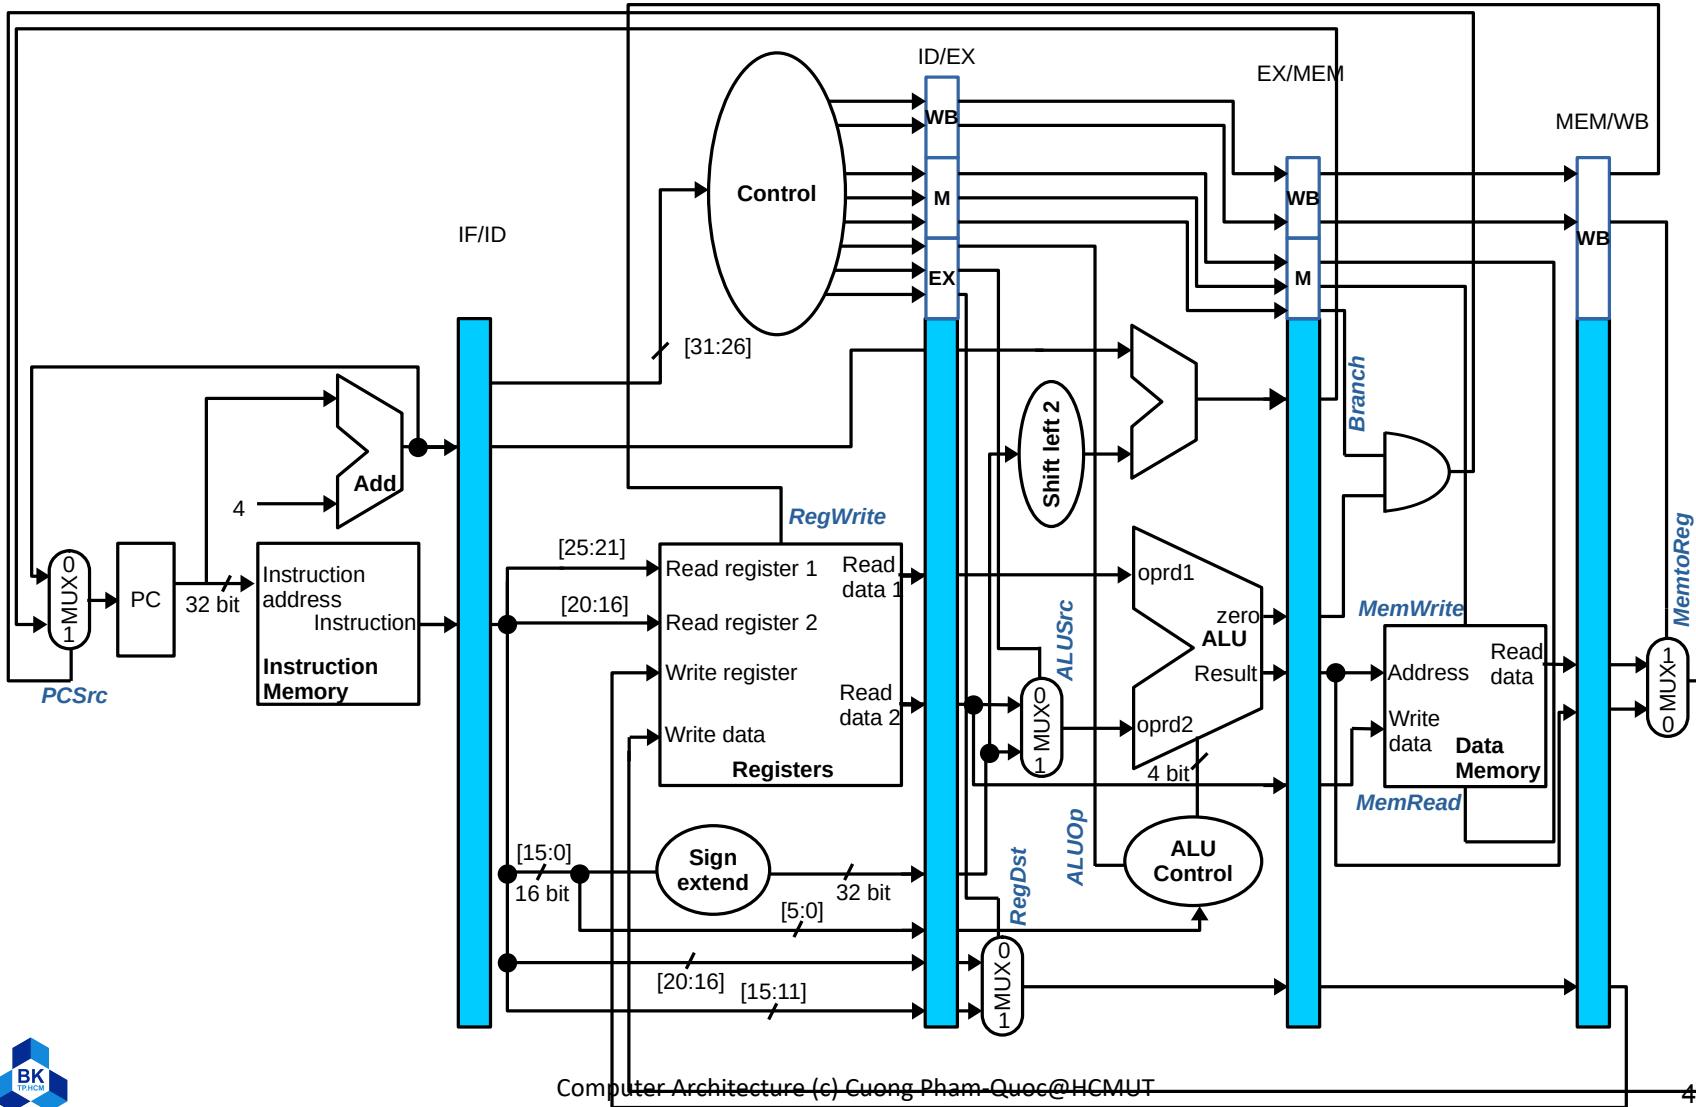
- **Question:** Given the following MIPS sequence:

```
lw    $s0, 20($s1)
sub  $t2, $s2, $s3
add  $t3, $s3, $s4
lw    $t4, 24($s1)
add  $t5, $s5, $s6
```

Assume that the sequence is executed by a 5-stage pipelined MIPS processor

- a) Identify values of control signals in cycle 5 **at the functional units**
- b) Identify values of control signals in cycle 5 **at the Control block**

Full pipeline micro-architecture



Hazards

- Situations that **prevent** starting the next instruction in the next cycle
- Structure hazard
 - A required **resource** is busy
- Data hazard
 - Need to wait for previous instruction to complete its **data** read/write
- Control hazard
 - Deciding on **control action** depends on previous instruction

Structure hazards

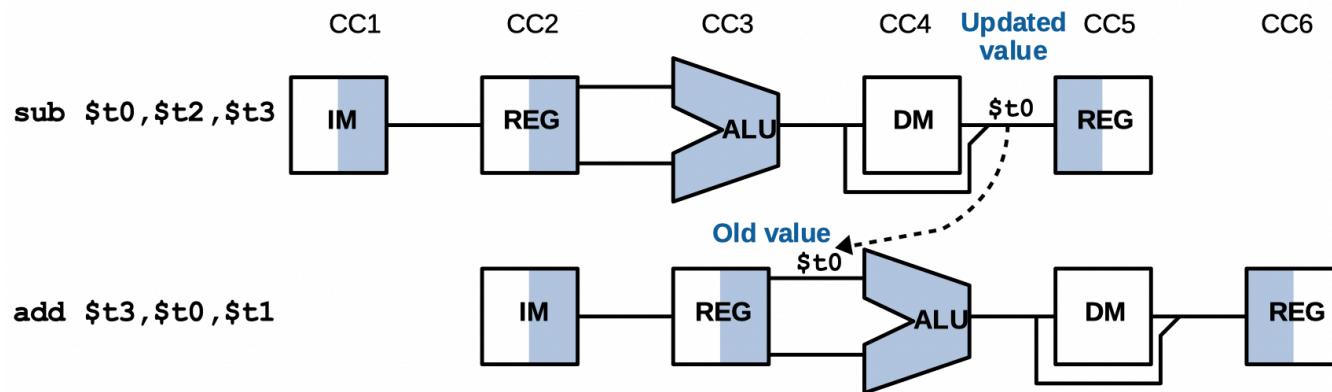
- **Conflict** for use of a resource
 - e.g.,: the laundry process, B forgot to bring clothes from the washing machine to dryer
- Should be eliminated entirely
 - Stall the pipe for that cycle
 - May repeat many times
- MIPS processors already **solved** all structure hazards
 - Separated instructions and data memory (caches)
 - Read and write registers use **different ports**

Data hazards

- An instruction depends on completion of data access by a previous instruction

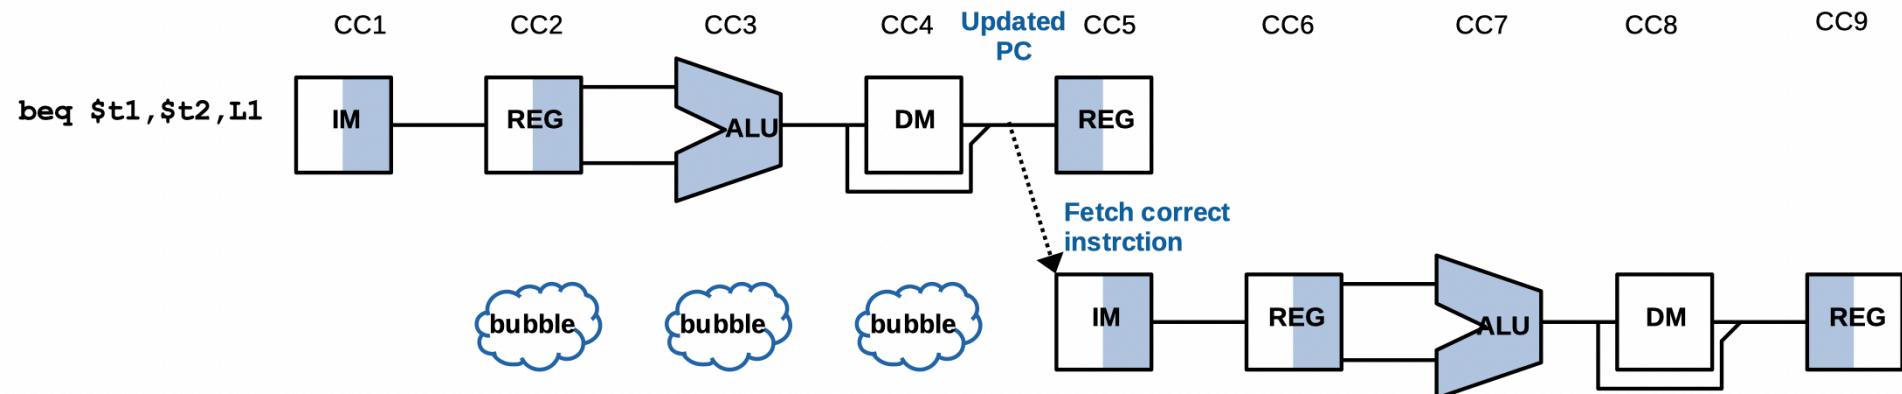
```
sub $t0, $t2, $t3  
add $t3, $t0, $t1
```

- No any issues with the **simplified version**
- Problem** with pipeline



Control hazards

- Branch determines flow of control
 - Fetching next instruction depends on branch outcome
- Pipeline updates PC in the MEM stage
 - Still working on ID stage of branch



Data hazards solutions

1. Code rescheduling

- Done by **compiler** (software level)
- *Sometimes cannot* find solutions

2. Delay or stalls insertion

- Done by **hardware**
- **Always** can find solutions
- *Increase* execution time & need *extra hardware* resources

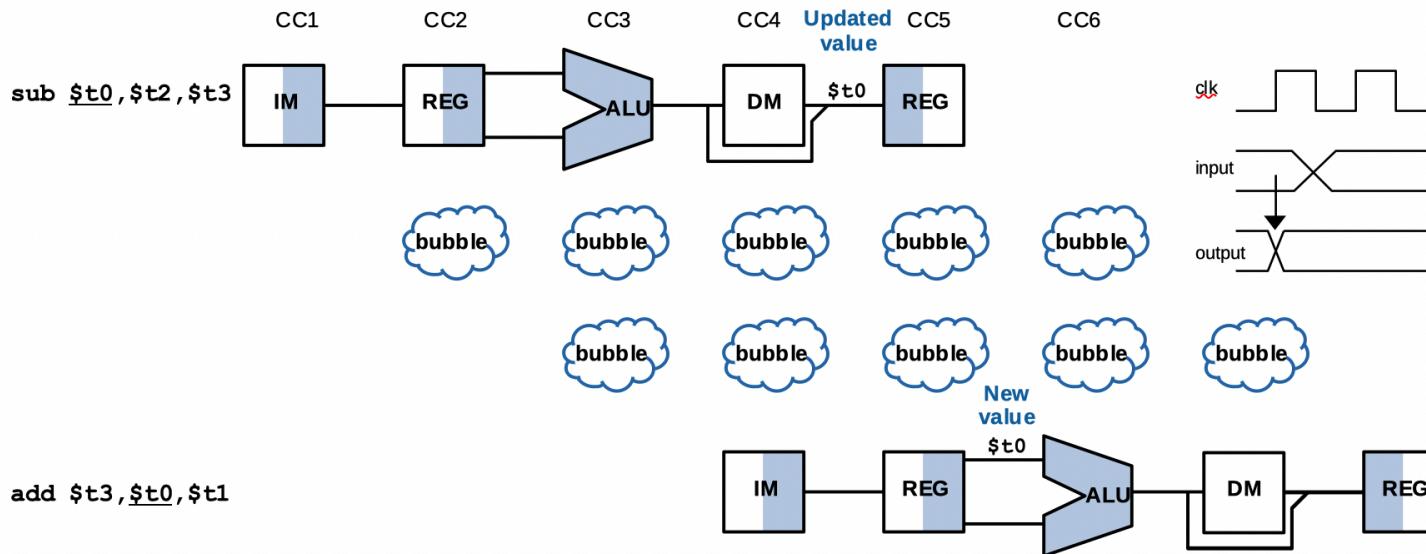
3. Forwarding

- Done by **hardware**
- *Cannot* find solutions in a *special case*
- Requires *extra hardware* resources



Code re-scheduling

- When haven't data hazards happened?
 - Read after or in the same cycle with Write (**RAW**)



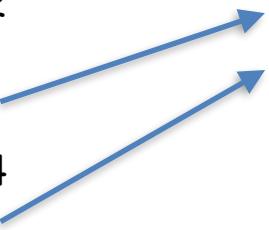
- Find and swap **data-independent** instructions

Example

- **Question:** given the following sequence of MIPS instructions

```
1: lw    $t1, 0($t0)  
2: lw    $t2, 4($t0)  
3: add  $t3, $t1, $t2  
4: sw    $t5, 12($t0)  
5: lw    $t4, 8($t0)  
6: add  $t5, $t1, $t4  
7: addi $t6, $t0, 4
```

```
1: lw    $t1, 0($t0)  
2: lw    $t2, 4($t0)  
5: lw    $t4, 8($t0)  
7: addi $t6, $t0, 4  
3: add  $t3, $t1, $t2  
4: sw    $t5, 12($t0)  
6: add  $t5, $t1, $t4
```



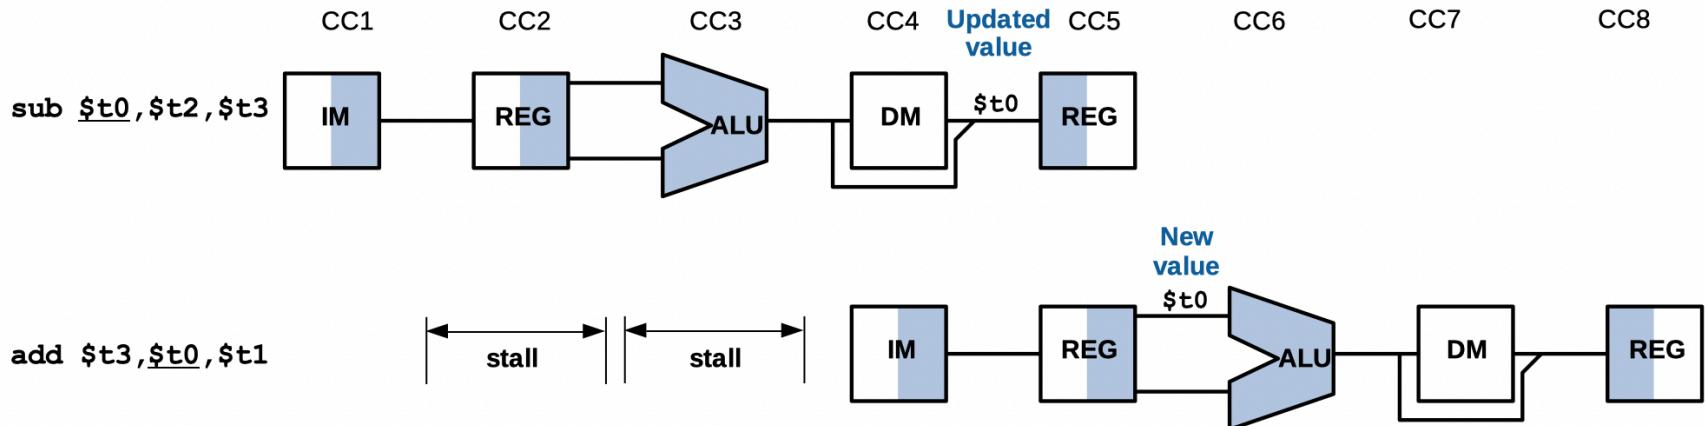
- Identify data hazards and solve by code re-scheduling

- **Answer:**

- The third and the sixth instructions (add \$t3, \$t1, \$t2 and add \$t5, \$t1, \$t4) have data hazards
- Move two data-independent instructions 5 and 7 to right after the second instruction

Stalls insertion (Delay)

- When haven't data hazards happened?
 - Read after or in the same cycle with Write (RAW)



- Delay instructions that use data:
 - ID stage of the instruction using data \equiv WB stage of the instruction producing data

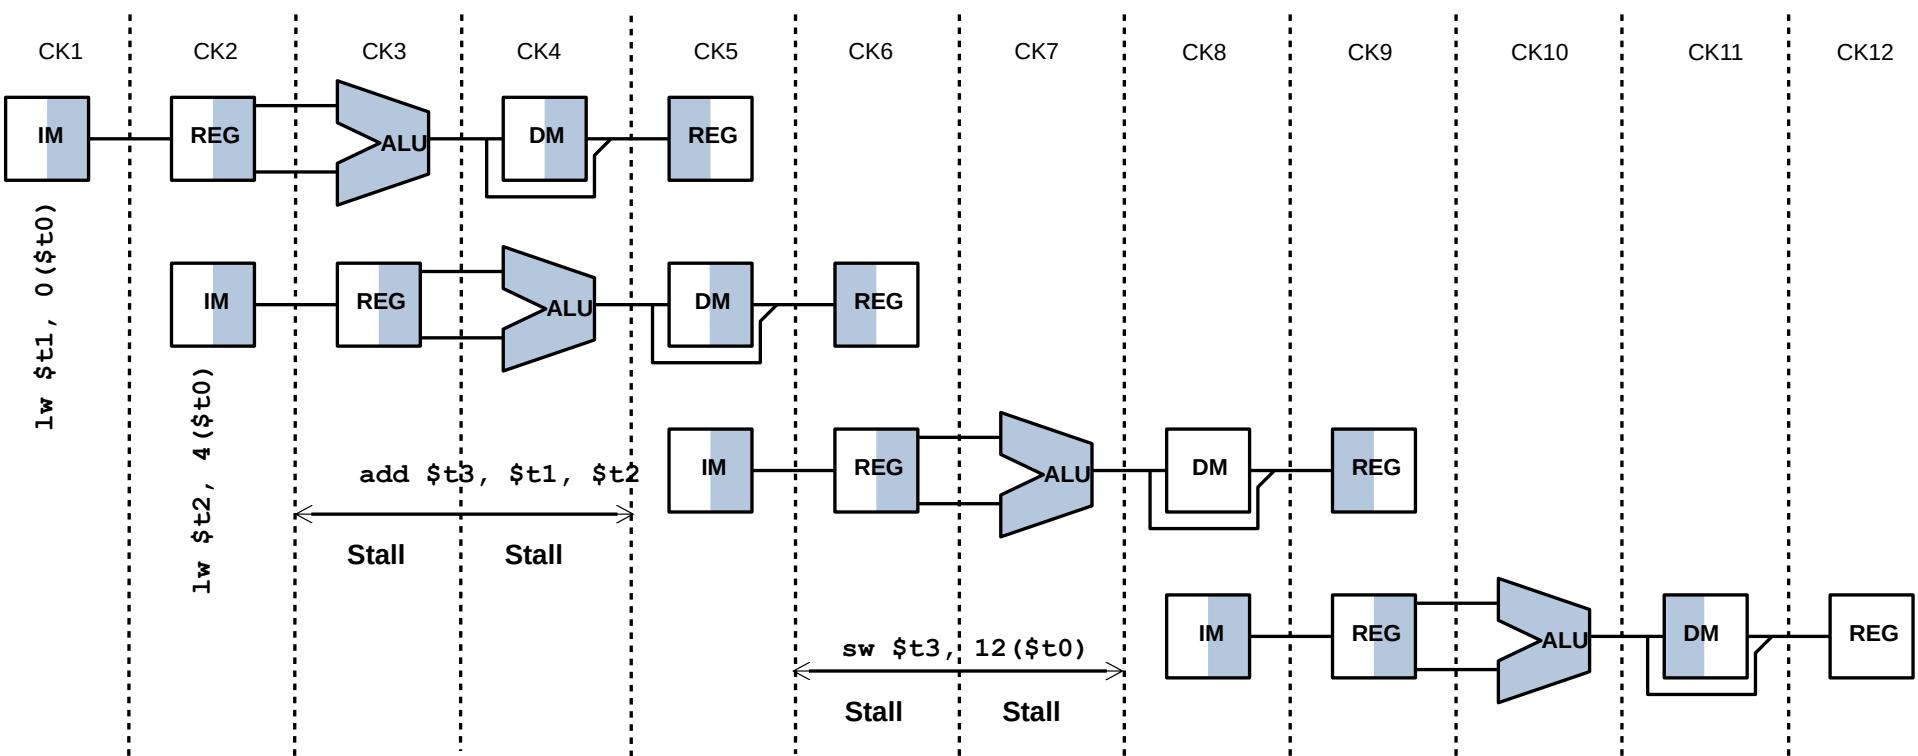
Example

- **Questions:** given the following MIPS sequence of instructions

1: lw \$t1, 0(\$t0)
2: lw \$t2, 4(\$t0)
3: add \$t3, \$t1, \$t2
4: sw \$t3, 12(\$t0)

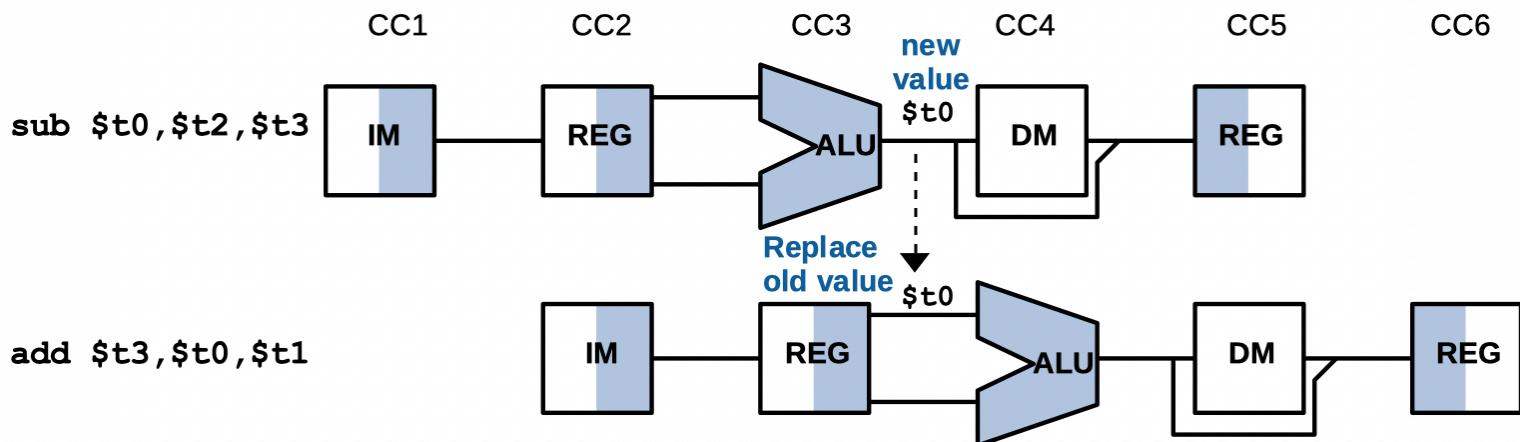
- Identify data hazards and solve them by the stalls insertion method;
how many cycles needed for the sequence?
- **Answer:**
 - Data hazards (2) - (3) & (3) - (4)
 - Insert two stalls for the 3rd instruction & two stalls for the 4th instruction
 - 12 cycles needed

Example (cont.)



Forwarding

- When can an instruction use data at **soonest**?
 - **Right after** data is produced (**EXE** stage & **MEM** stage)
- When is data **processed**?
 - **Apply** operators: **EXE** state



- **Use** data right after created

Example

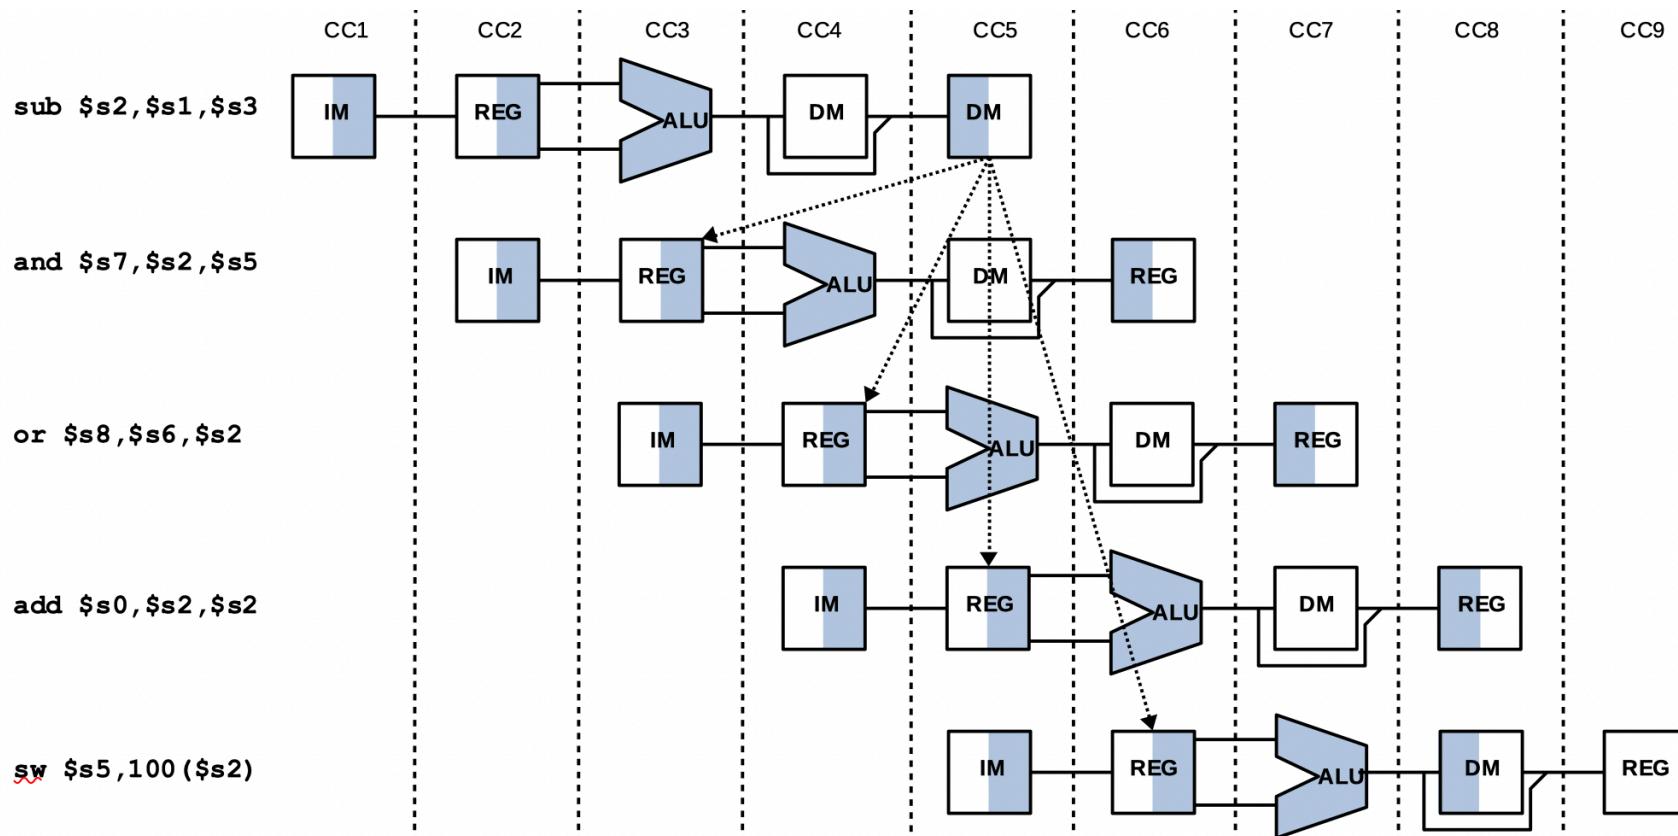
- **Question:** given the following MIPS sequence

```
1: sub $s2, $s1, $s3  
2: and $s7, $s2, $s5  
3: or $s8, $s6, $s2  
4: add $s0, $s2, $s2  
5: sw $s5, 100($s2)
```

- Analyze data dependencies & identify data hazards
- **Answer:**
 - **\$s2 produced** by the 1^{st} instruction **used** by all other instructions
 - **Data hazards:** 2^{nd} & 3^{rd} instructions

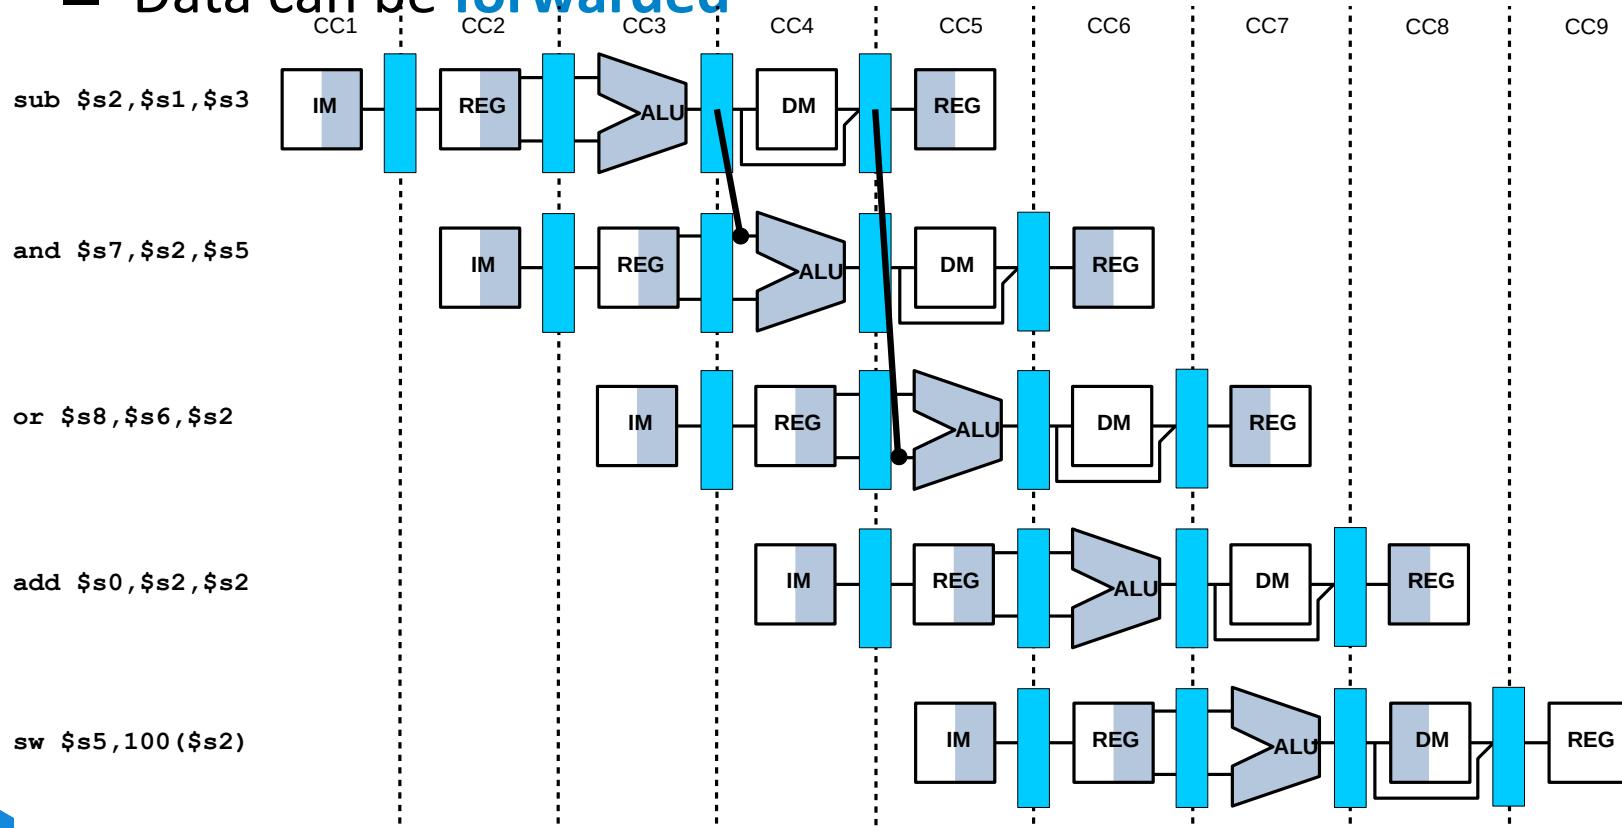
Example (cont.)

- Answer (cont.):



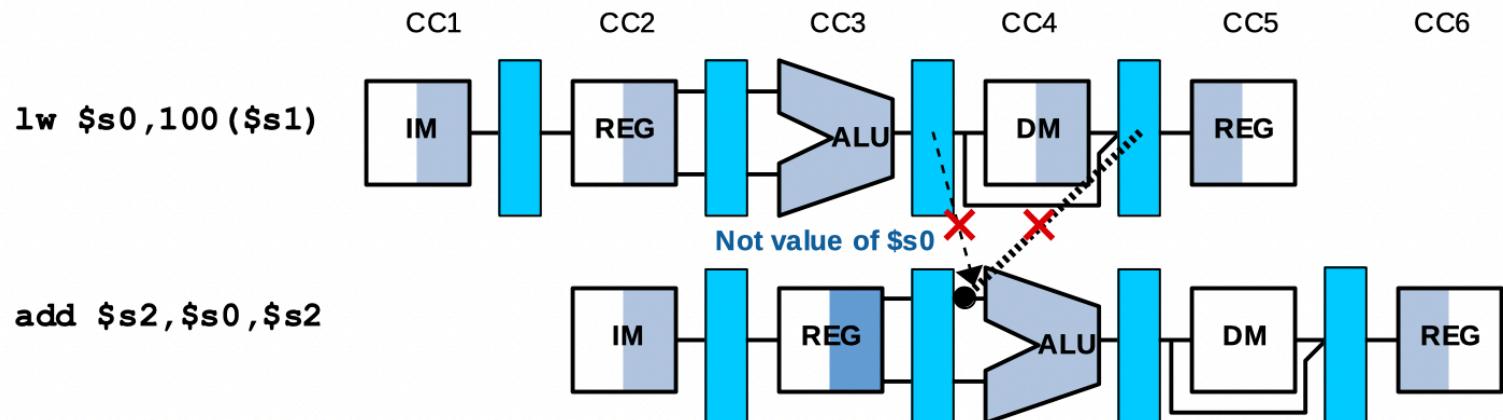
Forwarding

- Consider the previous sequence of MIPS instructions
 - Data can be **forwarded**



Load used data hazards

- When instruction producing data is a **load**, can data be forwarded?
 - Cannot forward** to the next instruction since data is produced later



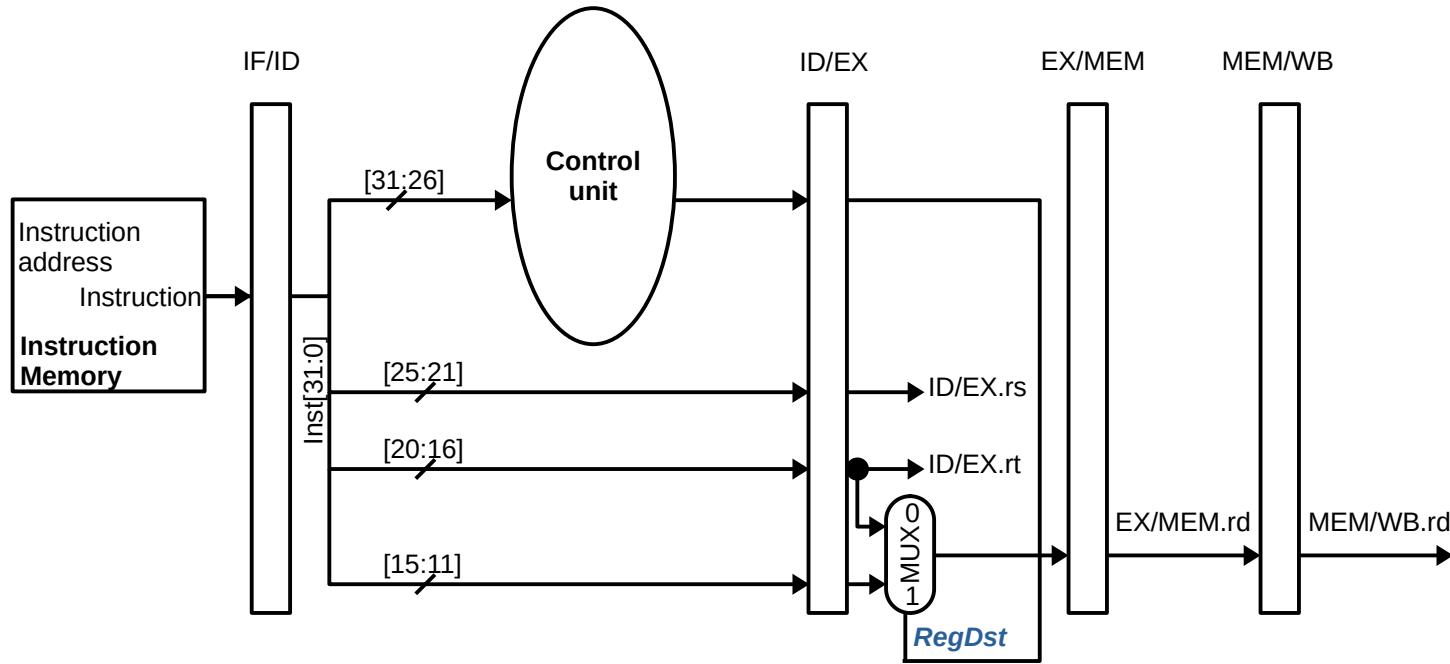
- Delay:**
 - 1 stall with forwarding or 2 stalls without forwarding**

Detecting data hazards

- **EXE** hazard **versus** **MEM** hazard
 - **EXE** hazard: forward from the **EX/MEM** register
 - **Destination register** in the MEM stage \equiv one of the source registers of the instruction in EXE
 - **MEM** hazard: forward from the **MEM/WB** register
 - **Destination register** in the WB stage \equiv one of the source registers of the instruction in EXE
- **Passing** register numbers along the pipe
 - ID/EX.rs & ID/EX.rt: first & second sources
 - EX/MEM.rd; MEM/WB.rd: destinations



EX hazard conditions



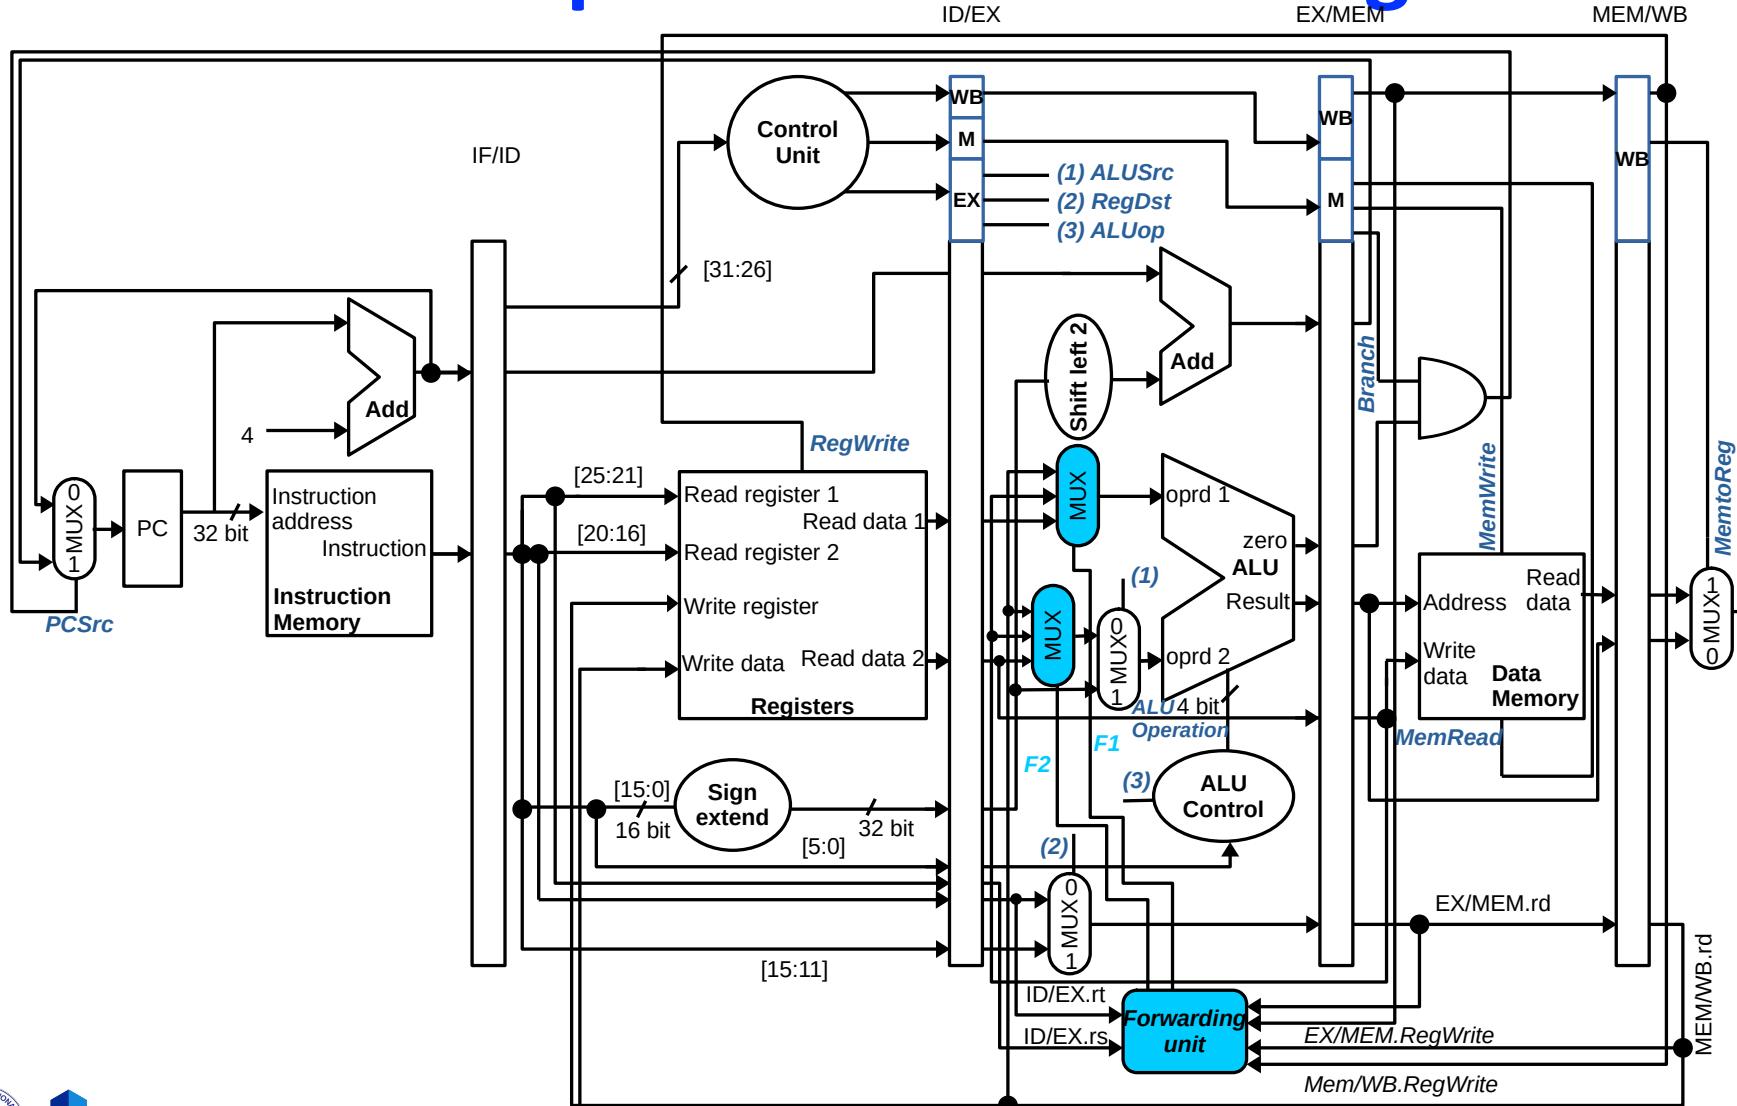
1a. $ID/EX.rs = EX/MEM.rd$

1b. $ID/EX.rt = EX/MEM.rd$ (not happened when an I-format instruction is in the EXE stage)

MEM hazard conditions

- Consider the following MIPS sequence of instructions
 - 1: add \$s0, \$s0, \$s1 1. IF ID EX ME WB
 - 2: add \$s0, \$s0, \$s2 2. IF ID EX ME WB
 - 3: add \$s0, \$s0, \$s3 3. IF ID EX
- Both **EX** and **MEM hazards** seem occur
 - EX hazard is **correct**
- MEM hazard conditions
 - 2a. (ID/EX.rs = MEM/WB.rd) & (ID/EX.rs != EX/MEM.rd)
 - 2b. (ID/EX.rt = MEM/WB.rd) & (ID/EX.rt != EX/MEM.rd)
(not happened when an **I-format** instruction is in the **EXE** stage)

Datapath with forwarding



Forwarding control values

Mux values (binary)	Source	Explanation
F1 = 00	ID/EX	The first ALU operand comes from the registers file
F1 = 10	EX/MEM	The first ALU operand is forwarded from the prior ALU result
F1 = 01	MEM/WB	The first ALU operand is forwarded from data memory of an earlier ALU result
F2 = 00	ID/EX	The second ALU operand comes from the registers file
F2 = 10	EX/MEM	The second ALU operand is forwarded from the prior ALU result
F2 = 01	MEM/WB	The second ALU operand is forwarded from data memory of an earlier ALU result

Forwarding conditions

- EX hazard

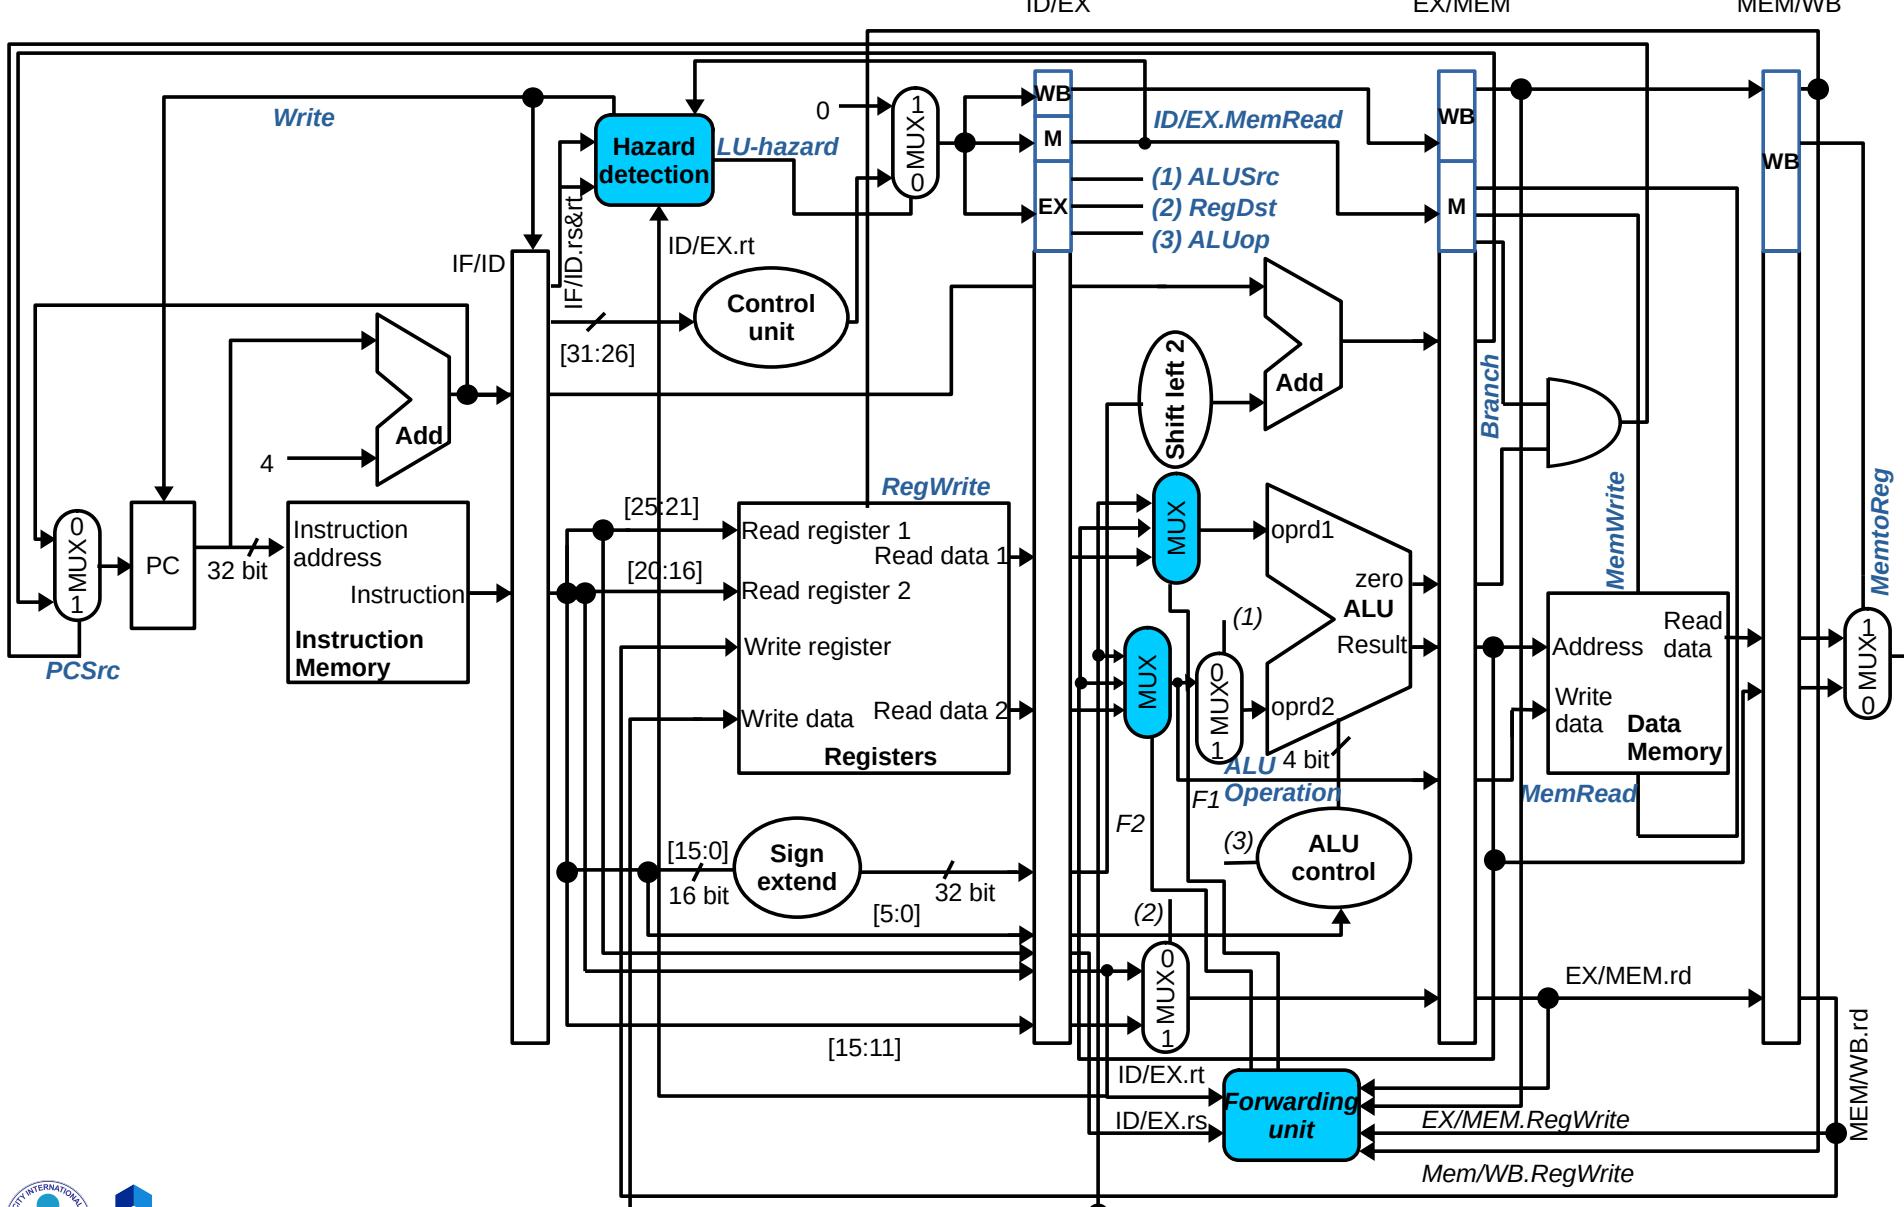
- 1a: if (EX/MEM.RegWrite and (EX/MEM.rd ≠ 0) and (EX/MEM.rd = ID/EX.rs)) **F1 = 10**
- 1b: if (EX/MEM.RegWrite and (EX/MEM.rd ≠ 0) and (EX/MEM.rd = ID/EX.rt)) **F2 = 10**

- MEM hazard

- 2a: if (MEM/WB.RegWrite and (MEM/WB.rd ≠ 0) and not (EX/MEM.RegWrite and (EX/MEM.rd ≠ 0) and (EX/MEM.rd = ID/EX.rs)) and (MEM/WB.rd = ID/EX.rs)) **F1 = 01**
- 2b: if (MEM/WB.RegWrite and (MEM/WB.rd ≠ 0) and not (EX/MEM.RegWrite and (EX/MEM.rd ≠ 0) and (EX/MEM.rd = ID/EX.rt)) and (MEM/WB.rd = ID/EX.rt)) **F2 = 01**

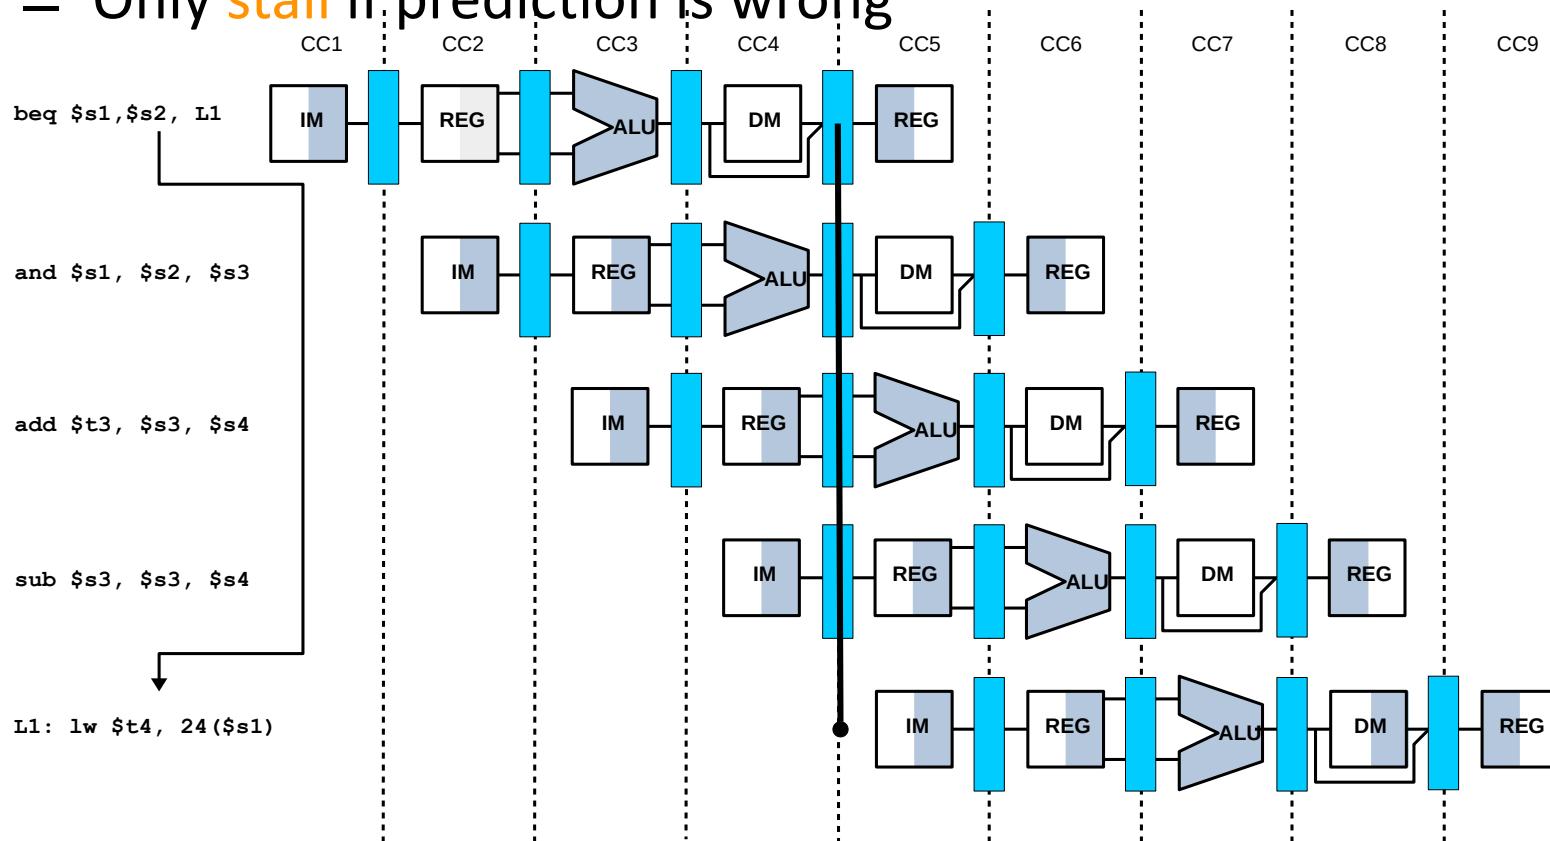
Detecting Load-use data hazards

- Check when using data instruction is decoded in **ID** stage
 - Producing data instruction (**load**) is in the **EXE** stage
 - The **sooner** the **better** due to a stall inserted
- Load-use hazard when
 - **ID/EX.MemRead** and ((**ID/EX.rt** = **IF/ID.rs**) **or** (**ID/EX.rt** = **IF/ID.rt**))
- **How to stall the pipeline?**
 - **Force** control signals in **ID/EX** register to 0 \Rightarrow EXE, MEM, and WB do nothing
 - **Prevent** updating PC and **ID/EX** register



Branch hazards solutions

- **Predict** outcome of branch
 - Only **stall** if prediction is wrong



Branch prediction

- **Static** branch prediction
 - Based on **typical branch behavior**
 - Example: loop and if-statement branches
 - Predict backward branches taken
 - Predict forward branches not taken
- **Dynamic** branch prediction
 - Hardware measures **actual branch behavior**
 - e.g., record recent history of each branch
 - Assume future behavior will continue the trend
 - When wrong, stall while re-fetching, and **update history**

Concluding remarks

- ISA influences design of **datapath** and **control**
- Datapath and control influence design of ISA
- **Pipelining** improves instruction throughput using parallelism
 - More instructions completed per second
 - Latency for each instruction not reduced
- **Hazards**: structural, data, control



The end

