

*Operating  
Systems:  
Internals  
and  
Design  
Principles*

# Chapter 4

# Threads

Seventh Edition  
By William Stallings

# Operating Systems: Internals and Design Principles

*The basic idea is that the several components in any complex system will perform particular subfunctions that contribute to the overall function.*

—THE SCIENCES OF THE ARTIFICIAL,

*Herbert Simon*



# Processes and Threads

\*Processes have two characteristics:

## Resource Ownership

Process includes a virtual address space to hold the process image

- the OS performs a protection function to prevent unwanted interference between processes with respect to resources

## Scheduling/Execution

Follows an execution path that may be interleaved with other processes

- a process has an execution state (Running, Ready, etc.) and a dispatching priority and is scheduled and dispatched by the OS

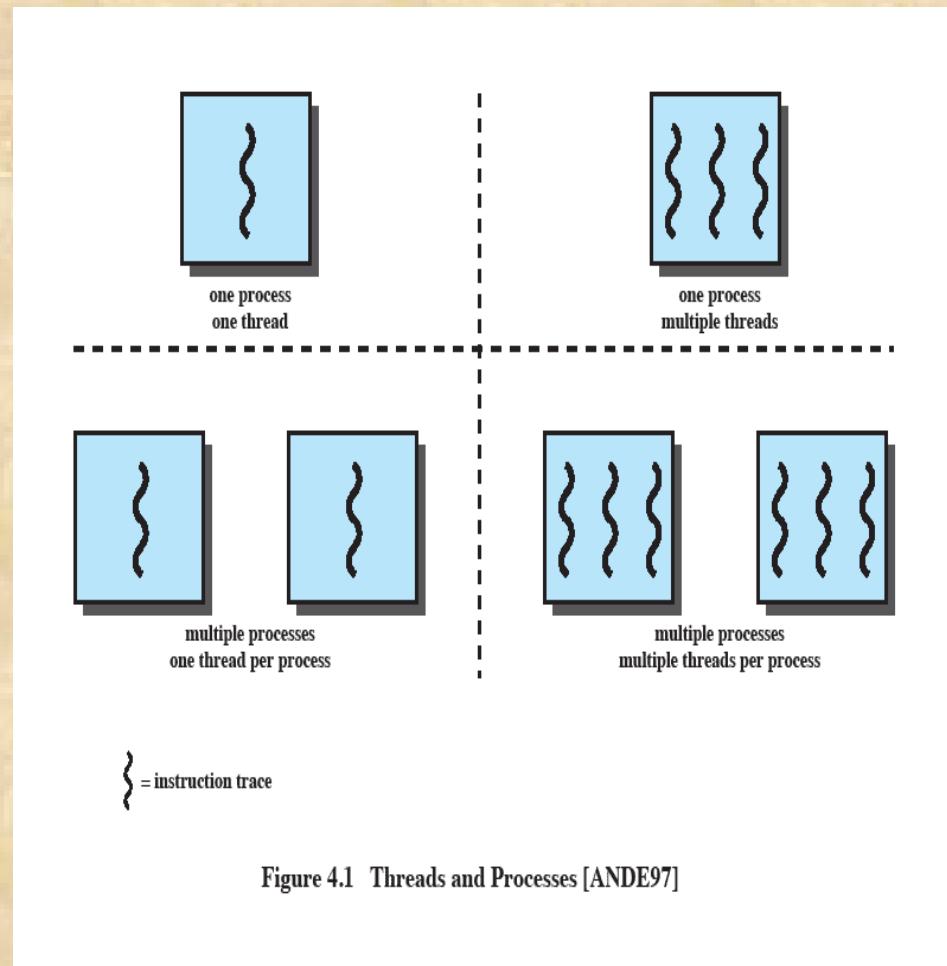


# Processes and Threads

- The unit of dispatching is referred to as a *thread* or *lightweight process*
- The unit of resource ownership is referred to as a *process* or *task*
- *Multithreading* - The ability of an OS to support multiple, concurrent paths of execution within a single process

# Single Threaded Approaches

- A single thread of execution per process, in which the concept of a thread is not recognized, is referred to as a single-threaded approach
- MS-DOS is an example



# Multithreaded Approaches

- The right half of Figure 4.1 depicts multithreaded approaches
- A Java run-time environment is an example of a system of one process with multiple threads

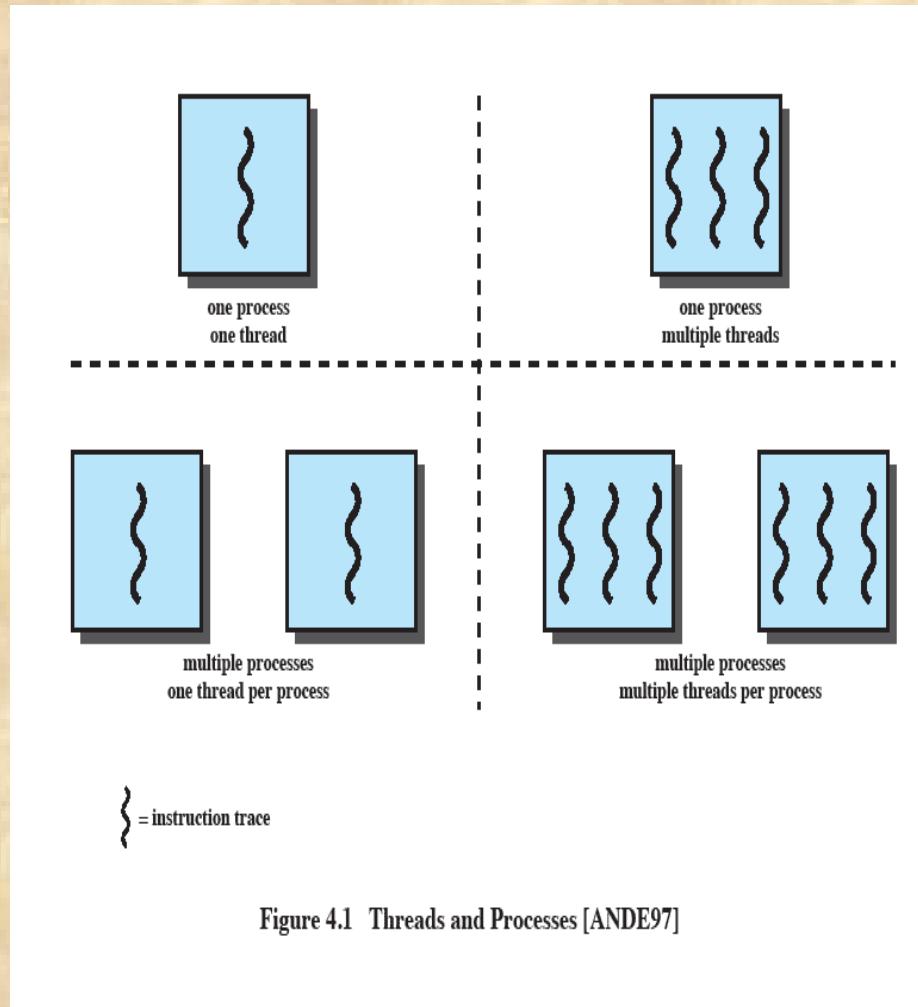
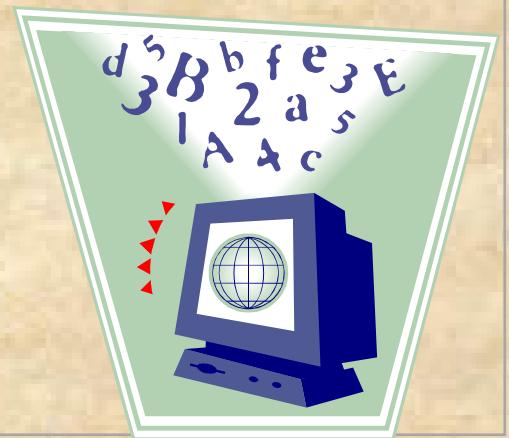


Figure 4.1 Threads and Processes [ANDE97]

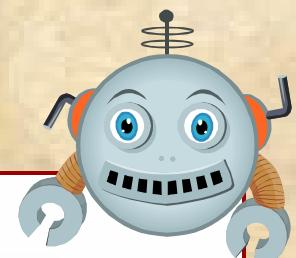
# Processes

- The unit or resource allocation and a unit of protection
- A virtual address space that holds the process image
- Protected access to:
  - processors
  - other processes
  - files
  - I/O resources



# One or More Threads in a Process

Each thread has:



- an execution state (Running, Ready, etc.)
- saved thread context when not running
- an execution stack
- some per-thread static storage for local variables
- access to the memory and resources of its process (all threads of a process share this)

# Threads vs. Processes

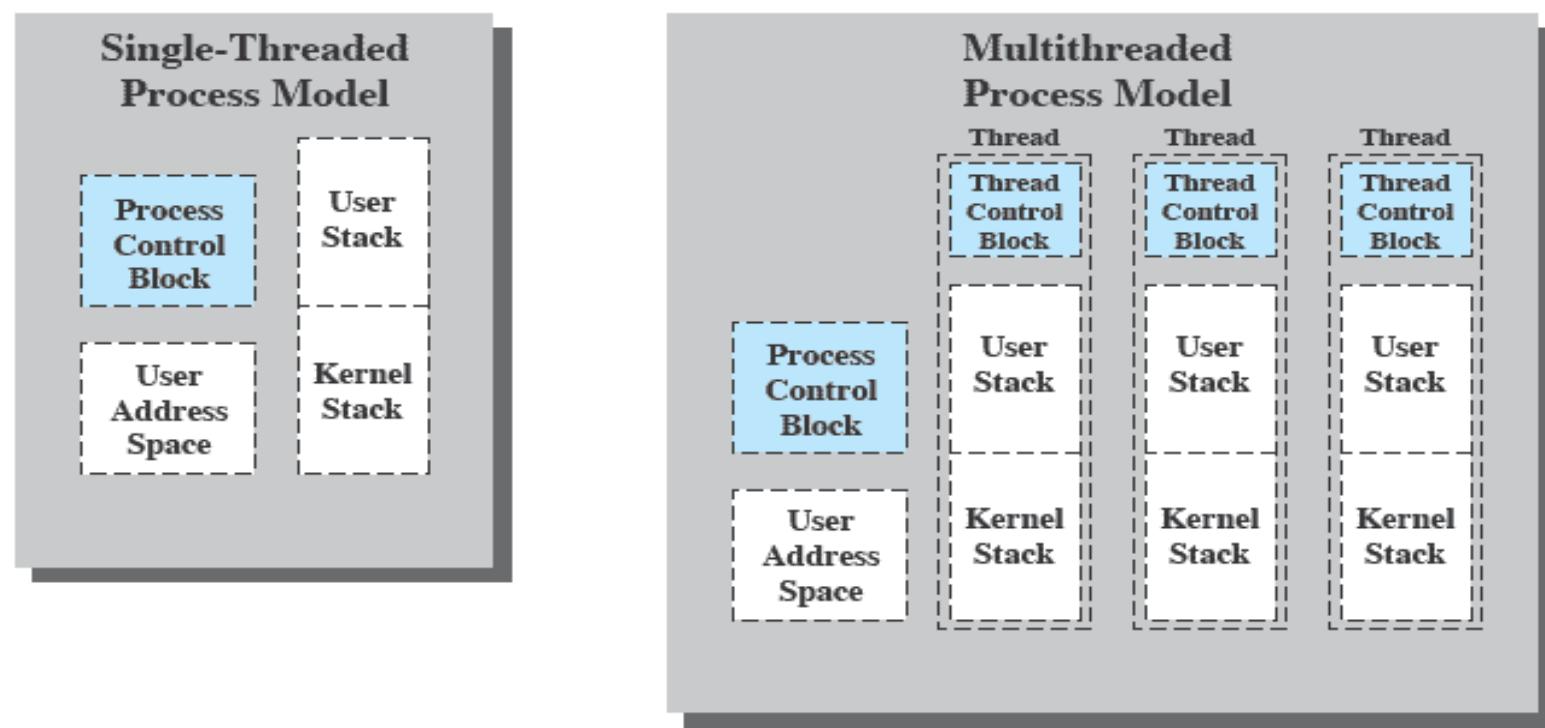
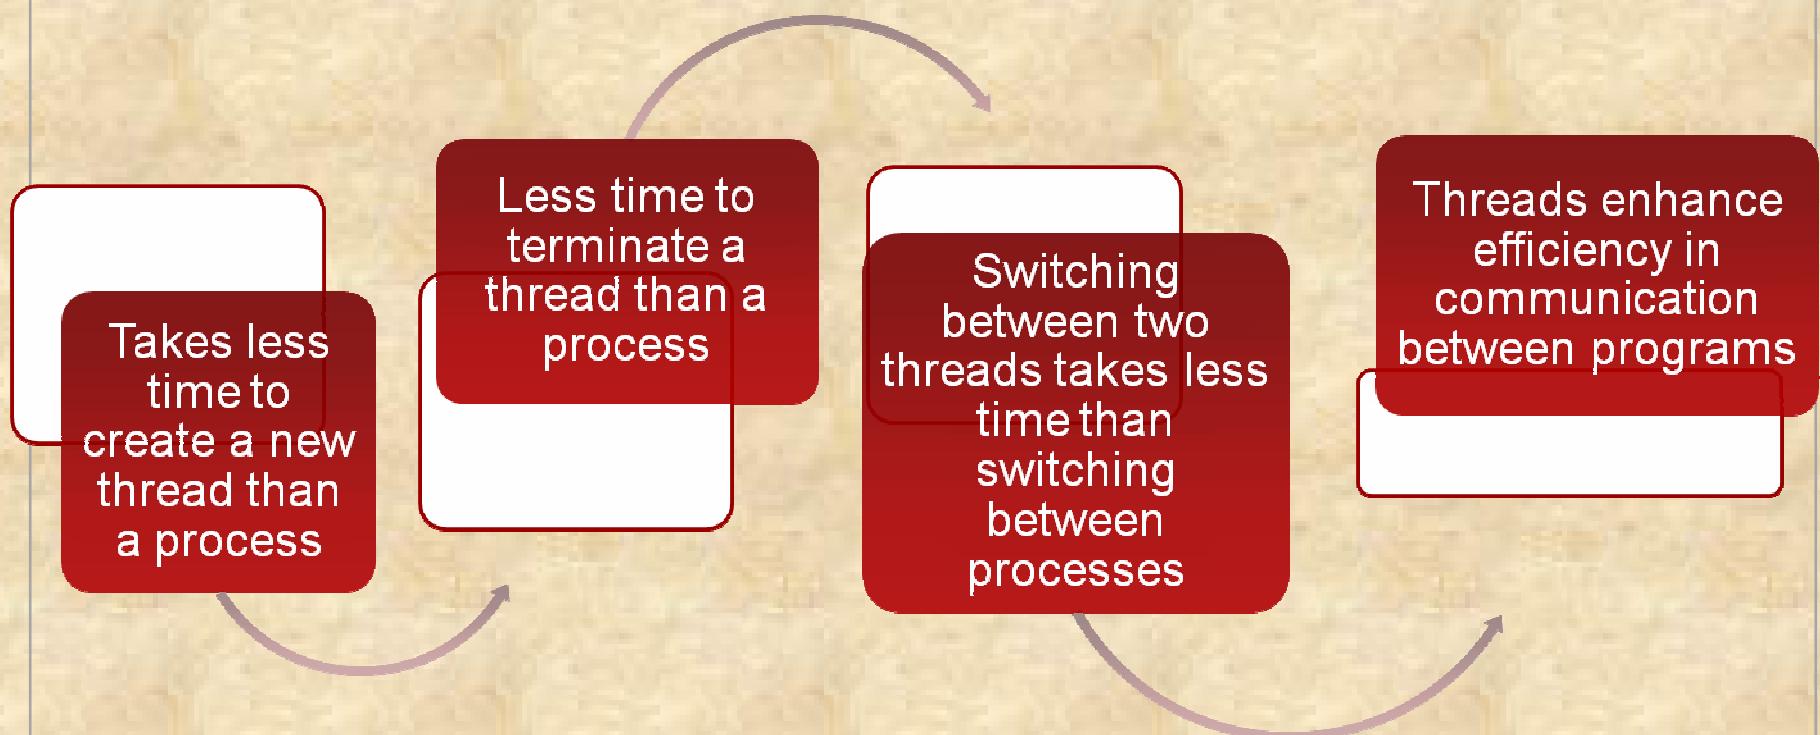


Figure 4.2 Single Threaded and Multithreaded Process Models

# Benefits of Threads



# Thread Use in a Single-User System

- Foreground and background work
- Asynchronous processing
- Speed of execution
- Modular program structure



# Threads

- In an OS that supports threads, scheduling and dispatching is done on a thread basis
- Most of the state information dealing with execution is maintained in thread-level data structures
  - ◆ suspending a process involves suspending all threads of the process
  - ◆ termination of a process terminates all threads within the process



# Thread Execution States

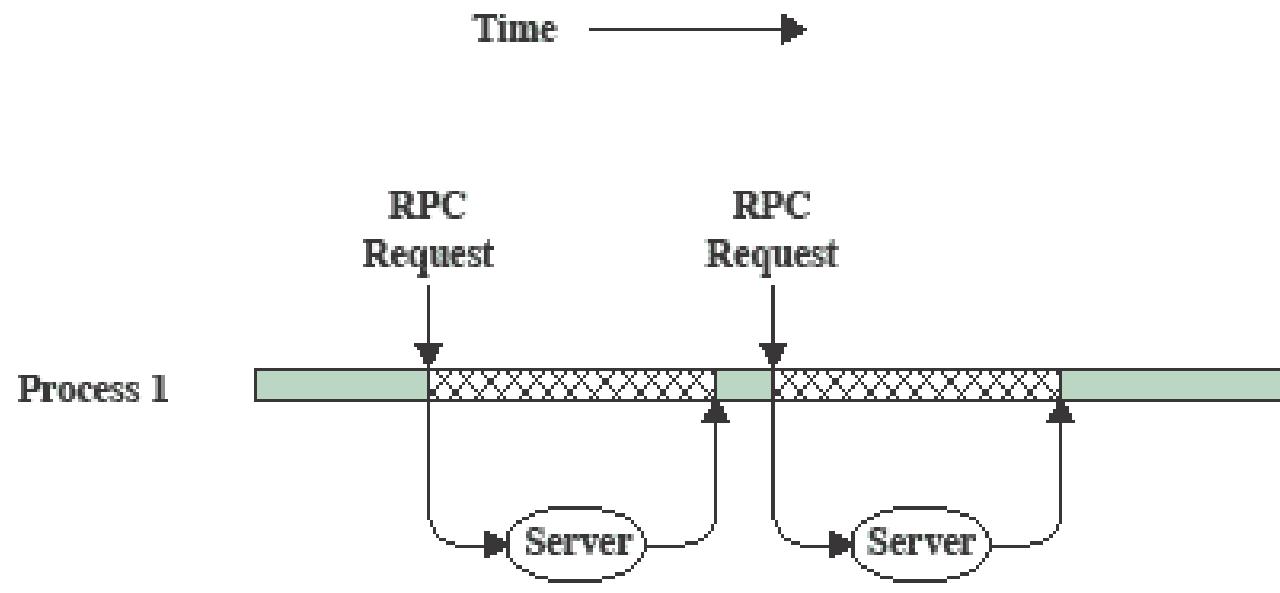
The key states for a thread are:

- Running
- Ready
- Blocked

Thread operations associated with a change in thread state are:

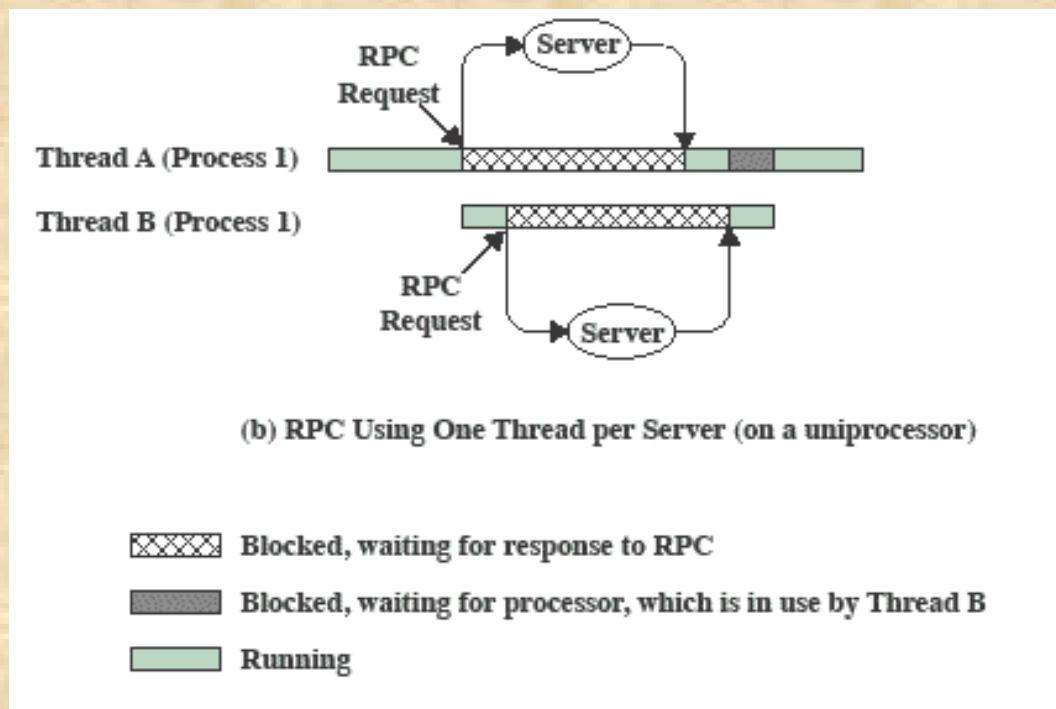
- Spawn
- Block
- Unblock
- Finish

# RPC Using Single Thread



(a) RPC Using Single Thread

# RPC Using One Thread per Server



# Multithreading on a Uniprocessor

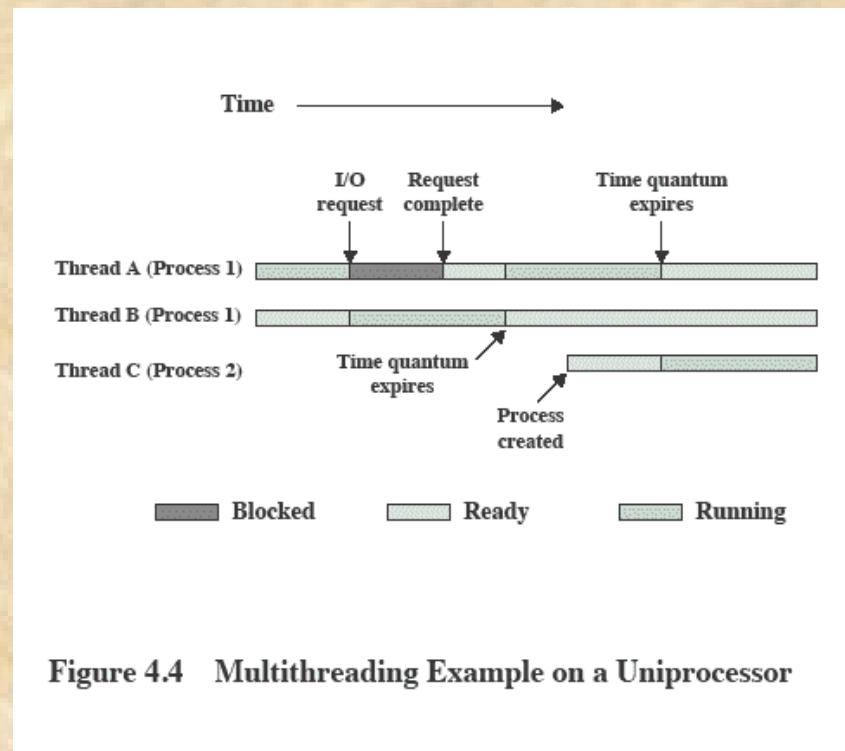


Figure 4.4 Multithreading Example on a Uniprocessor



# Thread Synchronization

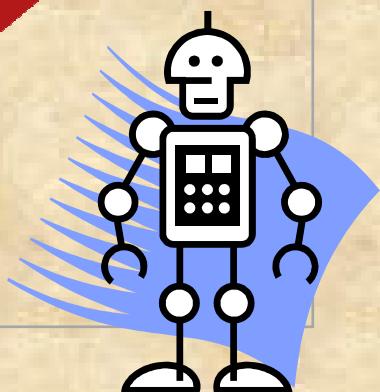
---

- It is necessary to synchronize the activities of the various threads
  - all threads of a process share the same address space and other resources
  - any alteration of a resource by one thread affects the other threads in the same process

# Types of Threads

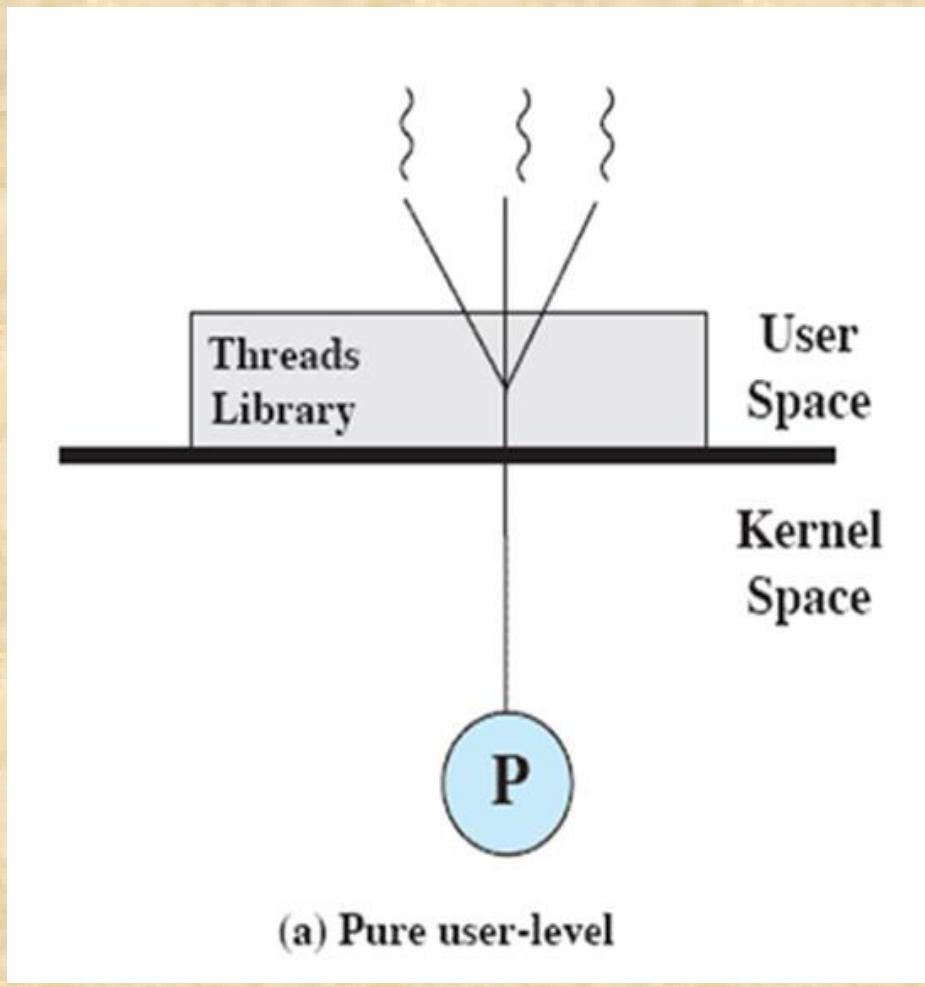
User Level  
Thread (ULT)

Kernel level  
Thread (KLT)



# User-Level Threads (ULTs)

- All thread management is done by the application
- The kernel is not aware of the existence of threads



# Relationships Between ULT States and Process States

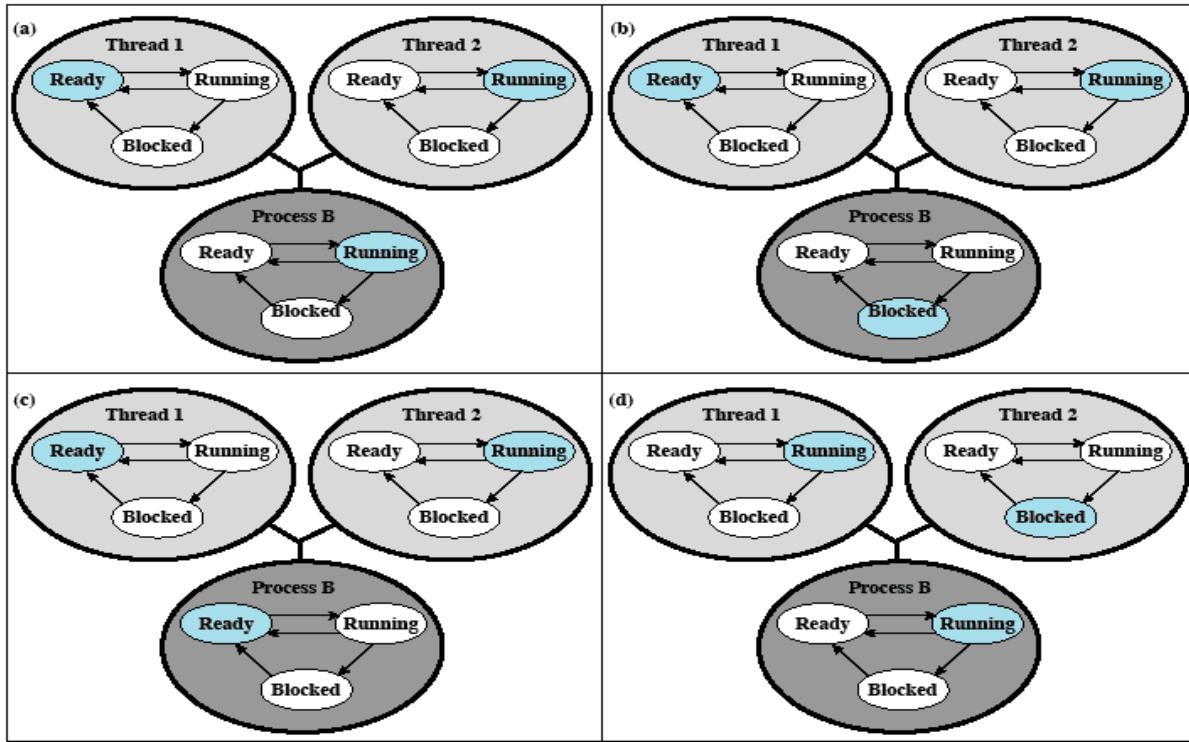


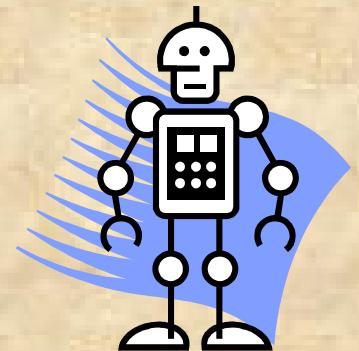
Figure 4.6 Examples of the Relationships between User-Level Thread States and Process States

# Advantages of ULTs

Thread switching does  
not require kernel mode  
privileges

Scheduling can be  
application specific

ULTs  
can run  
on any  
OS



# Disadvantages of ULTs

- In a typical OS many system calls are blocking
  - as a result, when a ULT executes a system call, not only is that thread blocked, but all of the threads within the process are blocked
- In a pure ULT strategy, a multithreaded application cannot take advantage of multiprocessing



# Overcoming ULT Disadvantages

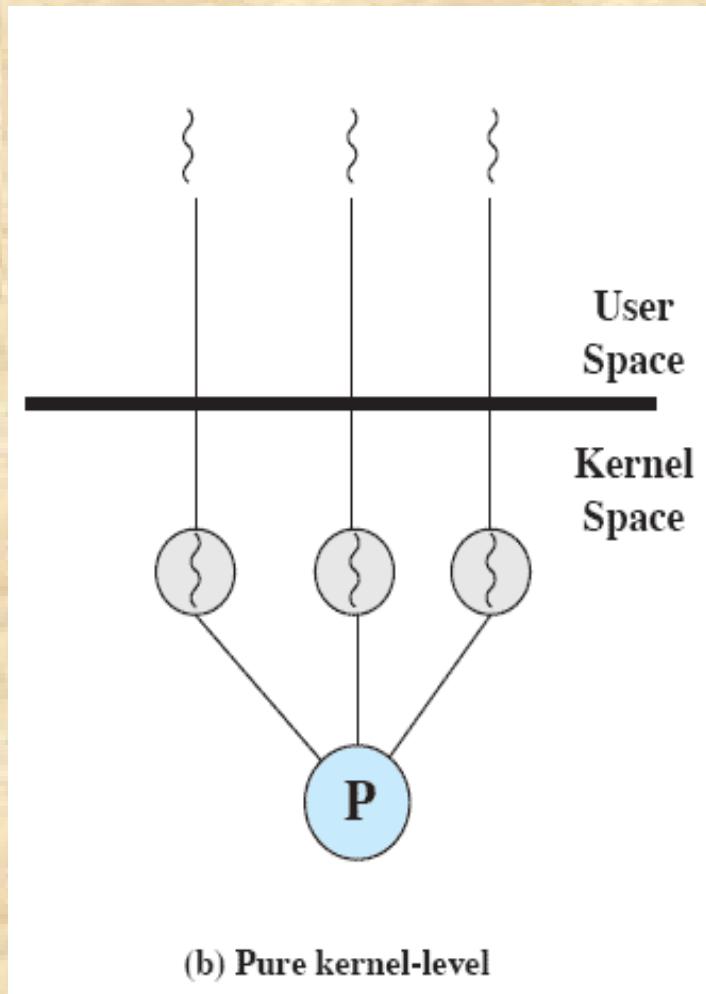
## Jacketing

- converts a blocking system call into a non-blocking system call



Writing an application as multiple processes rather than multiple threads

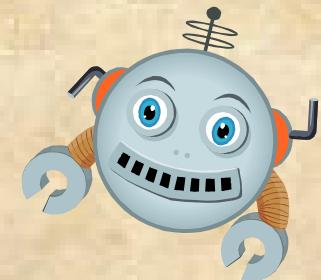
# Kernel-Level Threads (KLTs)



- ◆ Thread management is done by the kernel
  - ◆ no thread management is done by the application
- ◆ Windows is an example of this approach

# Advantages of KLTs

- The kernel can simultaneously schedule multiple threads from the same process on multiple processors
- If one thread in a process is blocked, the kernel can schedule another thread of the same process
- Kernel routines can be multithreaded

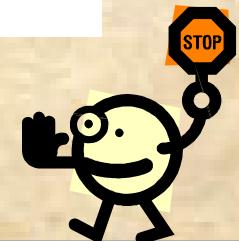


# Disadvantage of KLTs

- The transfer of control from one thread to another within the same process requires a mode switch to the kernel

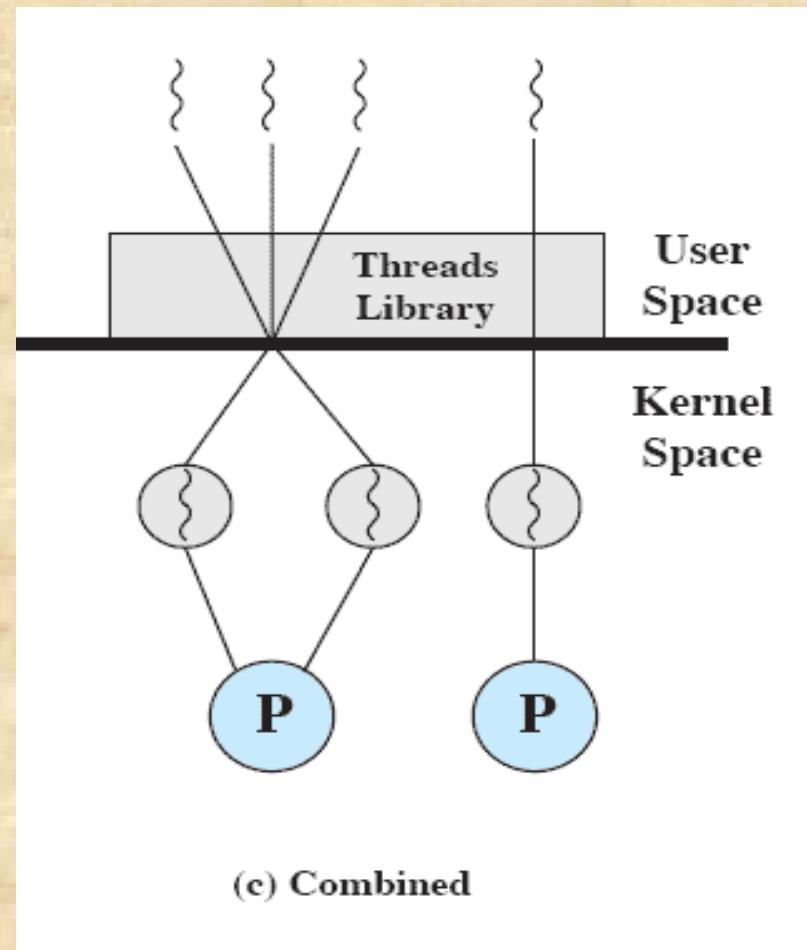
Operation	User-Level Threads	Kernel-Level Threads	Processes
Null Fork	34	948	11,300
Signal Wait	37	441	1,840

Table 4.1 Thread and Process Operation Latencies ( $\mu$ s)



# Combined Approaches

- Thread creation is done in the user space
- Bulk of scheduling and synchronization of threads is by the application
- Solaris is an example

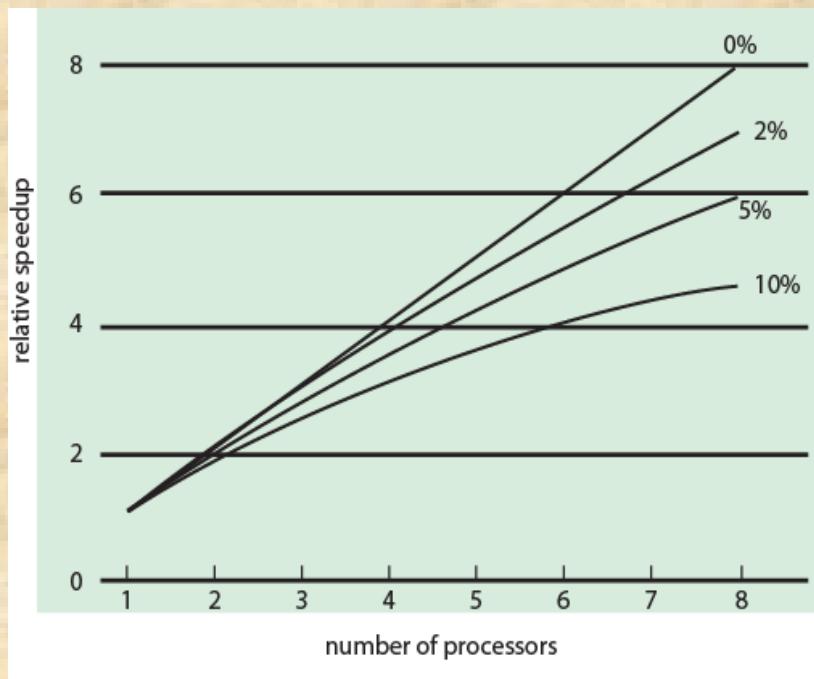


# Relationship Between Threads and Processes

Threads:Processes	Description	Example Systems
<b>1:1</b>	Each thread of execution is a unique process with its own address space and resources.	Traditional UNIX implementations
<b>M:1</b>	A process defines an address space and dynamic resource ownership. Multiple threads may be created and executed within that process.	Windows NT, Solaris, Linux, OS/2, OS/390, MACH
<b>1:M</b>	A thread may migrate from one process environment to another. This allows a thread to be easily moved among distinct systems.	Ra (Clouds), Emerald
<b>M:N</b>	Combines attributes of M:1 and 1:M cases.	TRIX

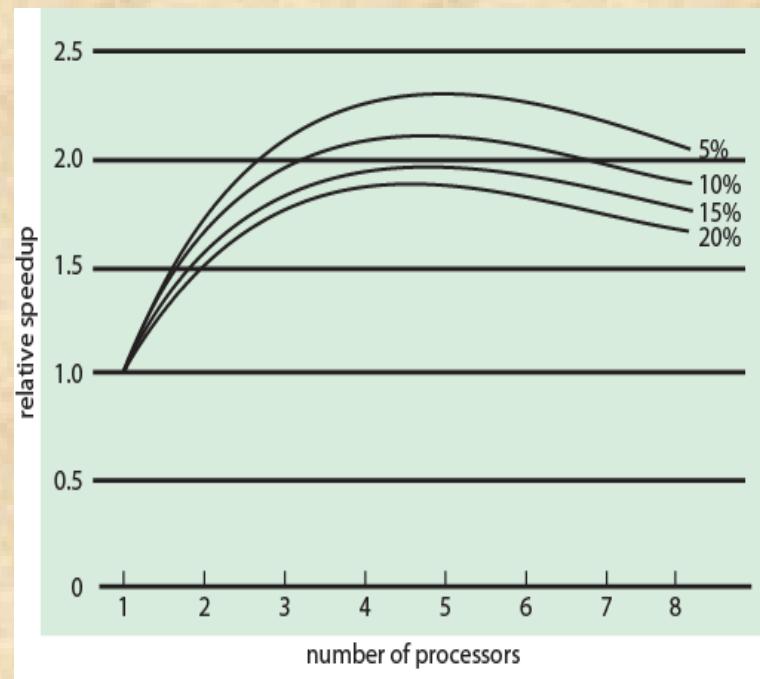
Table 4.2 Relationship between Threads and Processes

# Performance Effect of Multiple Cores



(a) Speedup with 0%, 2%, 5%, and 10% sequential portions

Figure 4.7 (a)



(b) Speedup with overheads

Figure 4.7 (b)

# Database Workloads on Multiple-Processor Hardware

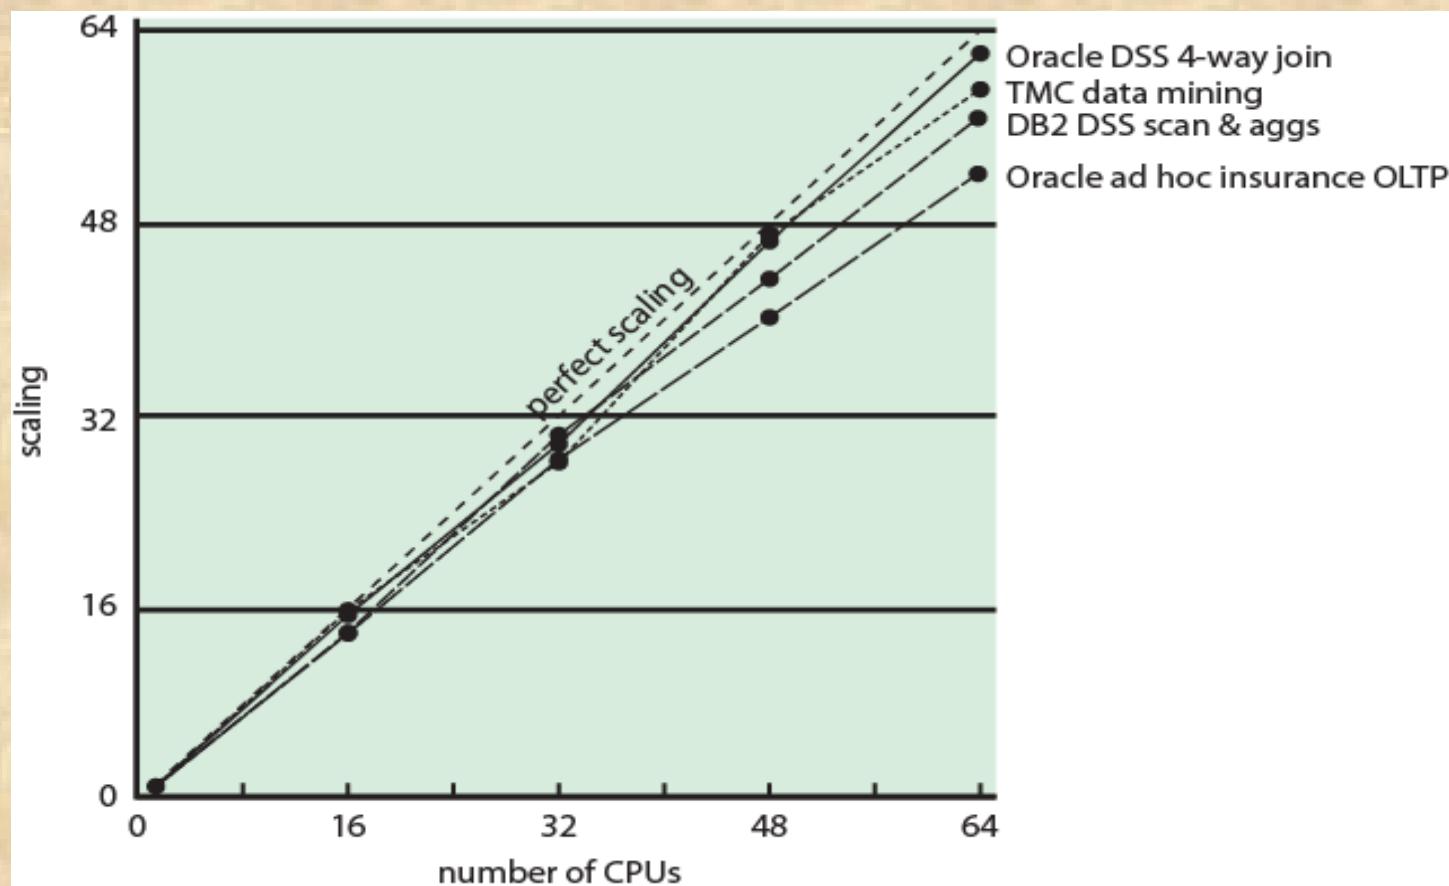


Figure 4.8 Scaling of Database Workloads on Multiple Processor Hardware

# Applications That Benefit

---

- ◆ Multithreaded native applications
  - ◆ characterized by having a small number of highly threaded processes
- ◆ Multiprocess applications
  - characterized by the presence of many single-threaded processes
- ◆ Java applications
- ◆ Multiinstance applications
  - multiple instances of the application in parallel

# Windows Processes

Processes and services provided by the Windows Kernel are relatively simple and general purpose

- implemented as objects
- created as new process or a copy of an existing
- an executable process may contain one or more threads
- both processes and thread objects have built-in synchronization capabilities



# Relationship Between Process and Resource

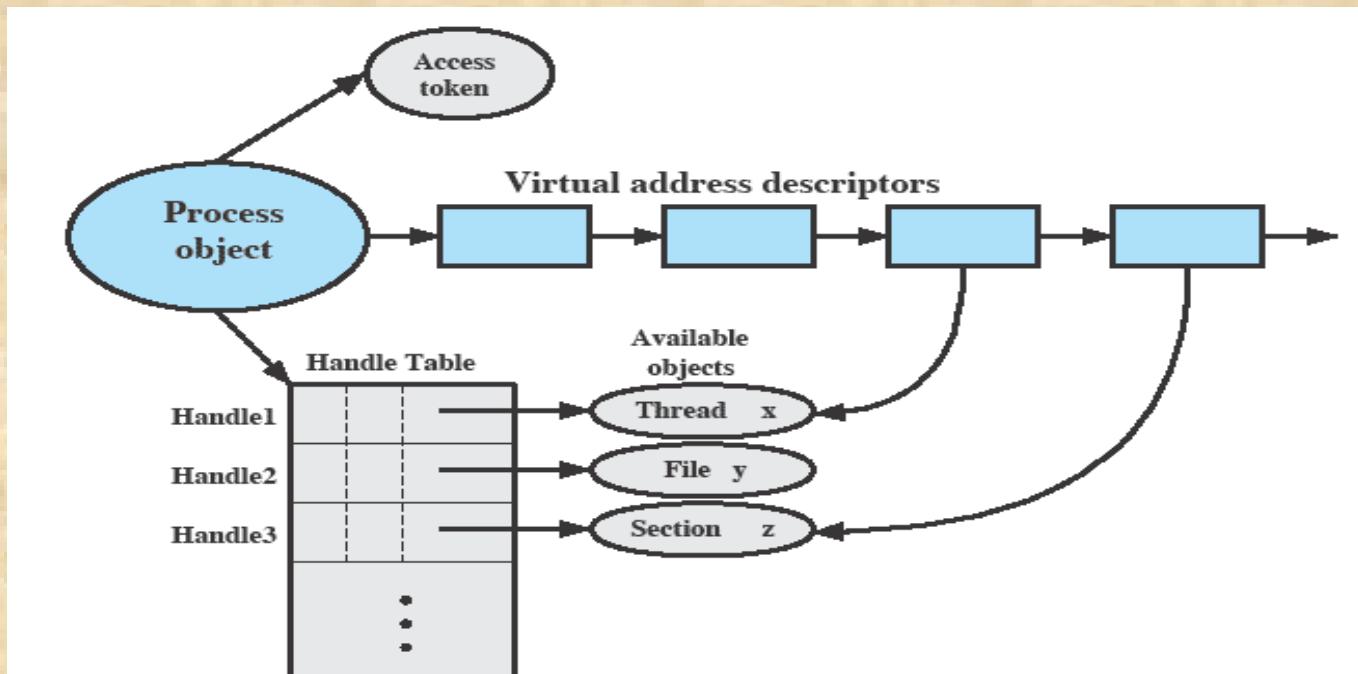


Figure 4.10 A Windows Process and Its Resources

Figure 4.12 A Windows Process and Its Resources

# Process and Thread Objects

Windows makes use of two types of process-related objects:

## Processes

- an entity corresponding to a user job or application that owns resources

## Threads

- a dispatchable unit of work that executes sequentially and is interruptible



# Windows Process and Thread Objects

## Object Type

### Process

Process ID  
Security Descriptor  
Base priority  
Default processor affinity  
Quota limits  
Execution time  
I/O counters  
VM operation counters  
Exception/debugging ports  
Exit status

## Object Body Attributes

Create process  
Open process  
Query process information  
Set process information  
Current process  
Terminate process

## Services

(a) Process object

## Object Type

### Thread

Thread ID  
Thread context  
Dynamic priority  
Base priority  
Thread processor affinity  
Thread execution time  
Alert status  
Suspension count  
Impersonation token  
Termination port  
Thread exit status

## Object Body Attributes

## Services

Create thread  
Open thread  
Query thread information  
Set thread information  
Current thread  
Terminate thread  
Get context  
Set context  
Suspend  
Resume  
Alert thread  
Test thread alert  
Register termination port

(b) Thread object



# Windows Process Object Attributes

<b>Process ID</b>	A unique value that identifies the process to the operating system.
<b>Security descriptor</b>	Describes who created an object, who can gain access to or use the object, and who is denied access to the object.
<b>Base priority</b>	A baseline execution priority for the process's threads.
<b>Default processor affinity</b>	The default set of processors on which the process's threads can run.
<b>Quota limits</b>	The maximum amount of paged and nonpaged system memory, paging file space, and processor time a user's processes can use.
<b>Execution time</b>	The total amount of time all threads in the process have executed.
<b>I/O counters</b>	Variables that record the number and type of I/O operations that the process's threads have performed.
<b>VM operation counters</b>	Variables that record the number and types of virtual memory operations that the process's threads have performed.
<b>Exception/debugging ports</b>	Interprocess communication channels to which the process manager sends a message when one of the process's threads causes an exception. Normally, these are connected to environment subsystem and debugger processes, respectively.
<b>Exit status</b>	The reason for a process's termination.

Table 4.3 Windows Process Object Attributes

# Windows Thread Object Attributes



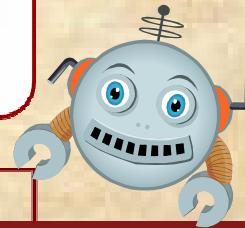
<b>Thread ID</b>	A unique value that identifies a thread when it calls a server.
<b>Thread context</b>	The set of register values and other volatile data that defines the execution state of a thread.
<b>Dynamic priority</b>	The thread's execution priority at any given moment.
<b>Base priority</b>	The lower limit of the thread's dynamic priority.
<b>Thread processor affinity</b>	The set of processors on which the thread can run, which is a subset or all of the processor affinity of the thread's process.
<b>Thread execution time</b>	The cumulative amount of time a thread has executed in user mode and in kernel mode.
<b>Alert status</b>	A flag that indicates whether a waiting thread may execute an asynchronous procedure call.
<b>Suspension count</b>	The number of times the thread's execution has been suspended without being resumed.
<b>Impersonation token</b>	A temporary access token allowing a thread to perform operations on behalf of another process (used by subsystems).
<b>Termination port</b>	An interprocess communication channel to which the process manager sends a message when the thread terminates (used by subsystems).
<b>Thread exit status</b>	The reason for a thread's termination.

Table 4.4 Windows Thread Object Attributes

# Multithreaded Process



Achieves concurrency without the overhead of using multiple processes



Threads within the same process can exchange information through their common address space and have access to the shared resources of the process

Threads in different processes can exchange information through shared memory that has been set up between the two processes

# Thread States

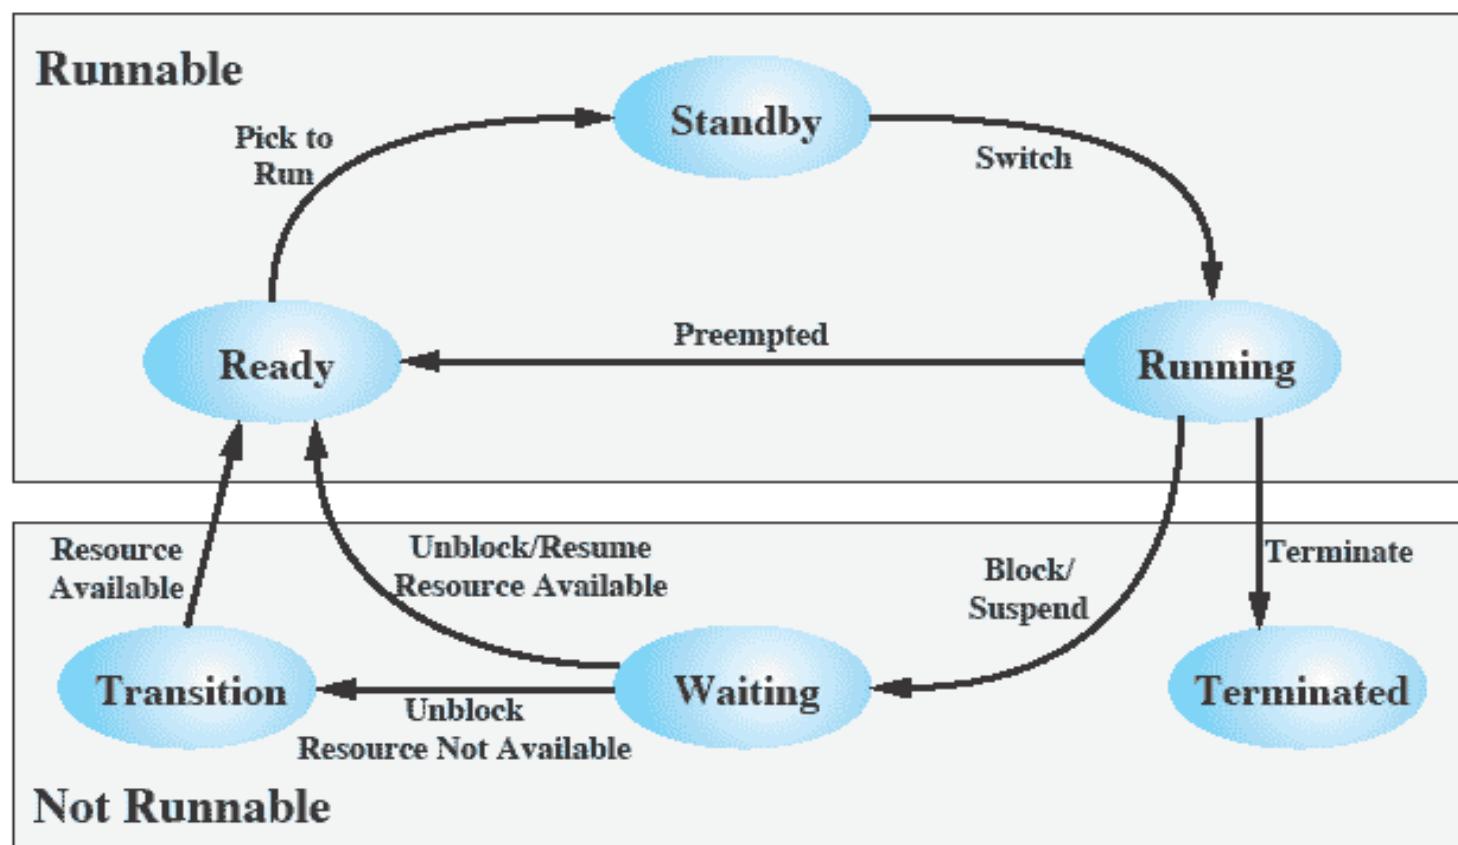


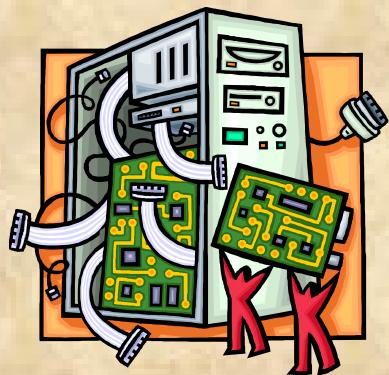
Figure 4.12 Windows Thread States

# Symmetric Multiprocessing Support (SMP)

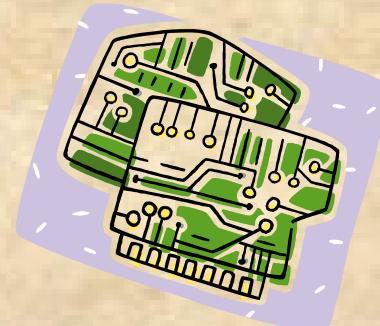
Threads of any process can run on any processor

## Soft Affinity

- the dispatcher tries to assign a ready thread to the same processor it last ran on
- helps reuse data still in that processor's memory caches from the previous execution of the thread



## Hard Affinity



# Solaris Process

:- makes use of four thread-related concepts:

## Process

- includes the user's address space, stack, and process control block

## User-level Threads

- a user-created unit of execution within a process

## Lightweight Processes (LWP)

- a mapping between ULTs and kernel threads

## Kernel Threads

- fundamental entities that can be scheduled and dispatched to run on one of the system processors

# Processes and Threads in Solaris

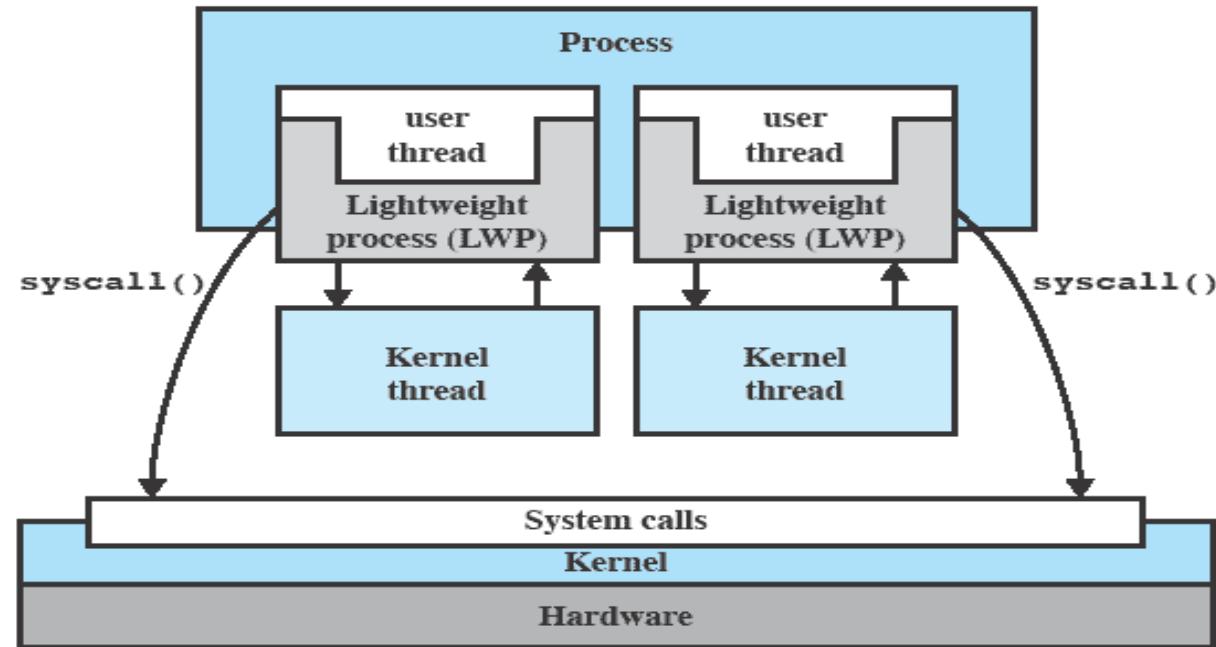
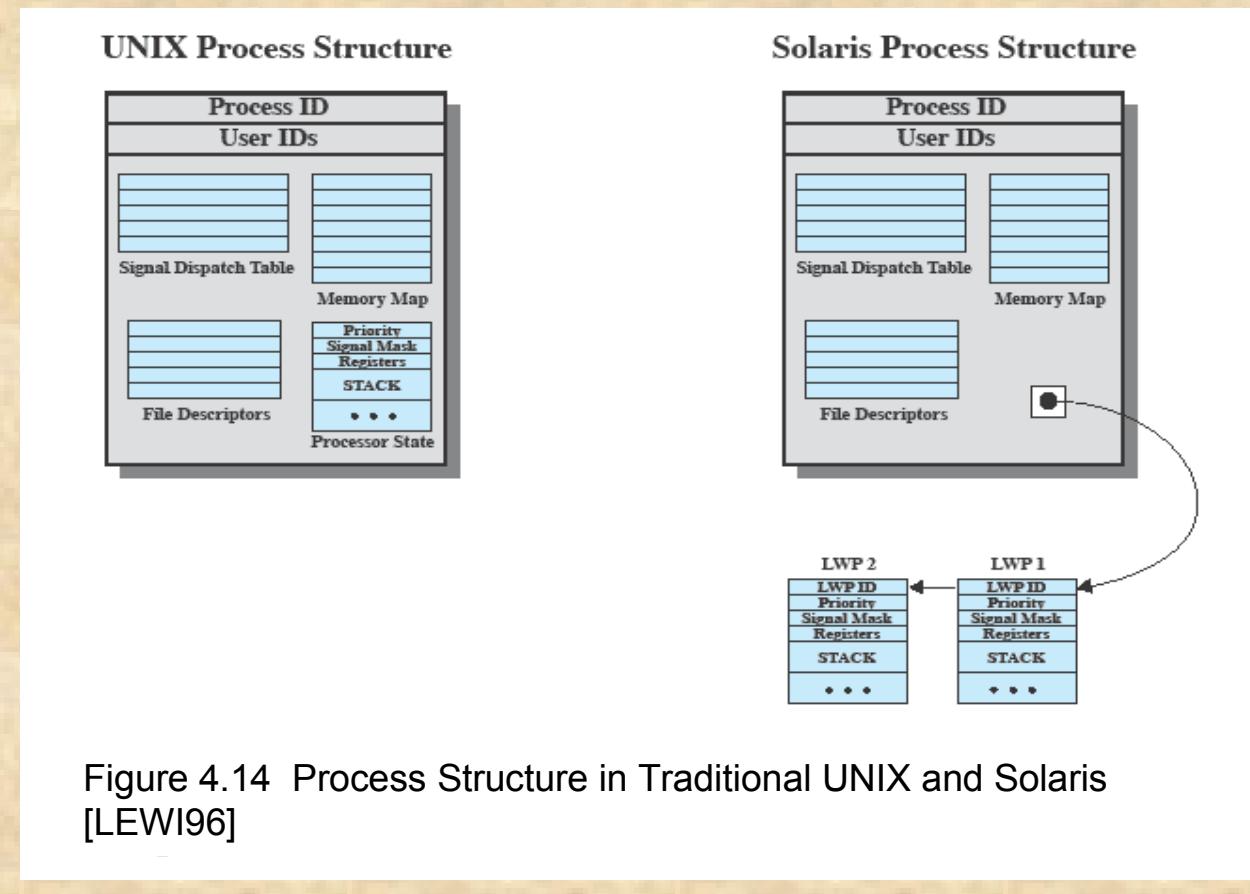


Figure 4.13 Processes and Threads in Solaris [MCDO07]

# Traditional Unix vs Solaris



# A Lightweight Process (LWP) Data Structure Includes:

- An LWP identifier
- The priority of this LWP
- A signal mask
- Saved values of user-level registers
- The kernel stack for this LWP
- Resource usage and profiling data
- Pointer to the corresponding kernel thread
- Pointer to the process structure



# Solaris Thread States

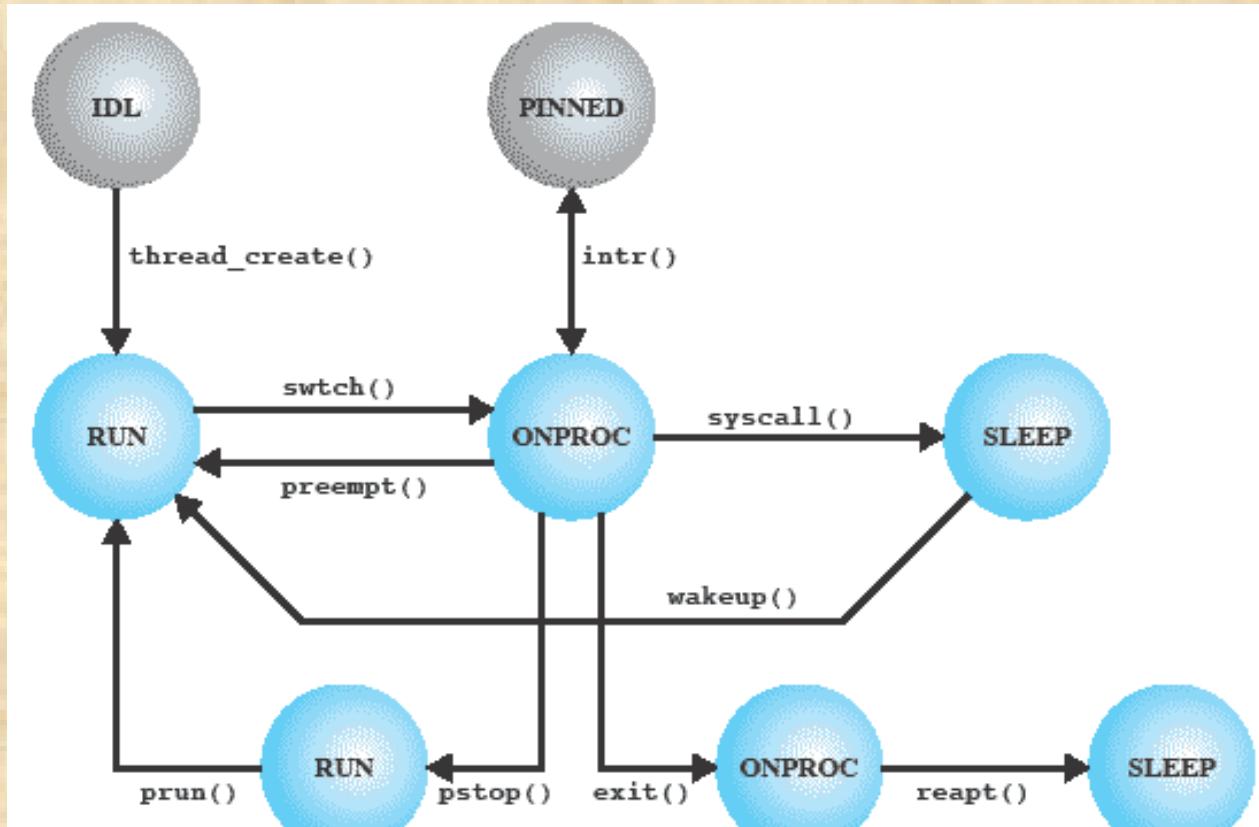


Figure 4.15 Solaris Thread States

# Interrupts as Threads

- ◆ Most operating systems contain two fundamental forms of concurrent activity:

## Processes (threads)

- cooperate with each other and manage the use of shared data structures by primitives that enforce mutual exclusion and synchronize their execution

## Interrupts

- synchronized by preventing their handling for a period of time

# Solaris Solution

---

- ◆ Solaris employs a set of kernel threads to handle interrupts
  - an interrupt thread has its own identifier, priority, context, and stack
  - the kernel controls access to data structures and synchronizes among interrupt threads using mutual exclusion primitives
  - interrupt threads are assigned higher priorities than all other types of kernel threads

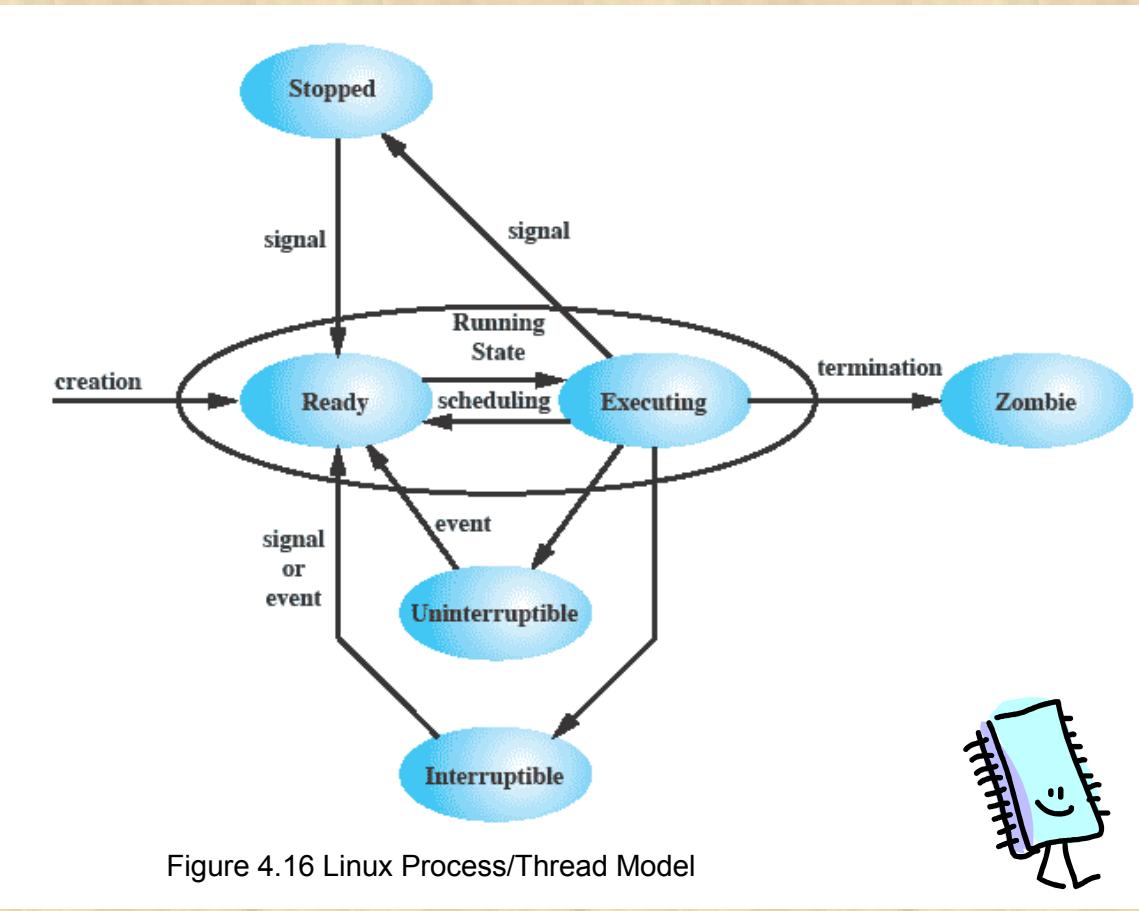
# Linux Tasks

A process, or task, in Linux is represented by a `task_struct` data structure



This structure contains information in a number of categories

# Linux Process/Thread Model

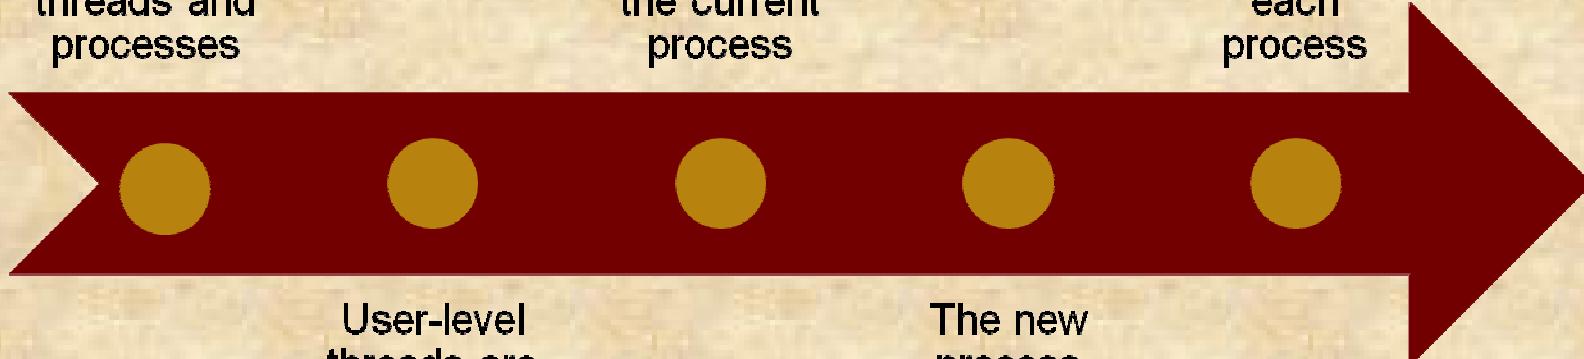


# Linux Threads

Linux does  
not  
recognize a  
distinction  
between  
threads and  
processes

A new  
process is  
created by  
copying the  
attributes of  
the current  
process

The `clone()`  
call creates  
separate  
stack  
spaces for  
each  
process



User-level  
threads are  
mapped  
into kernel-  
level  
processes

The new  
process  
can be  
cloned so  
that it  
shares  
resources

# Linux

## Clone ()

### Flags

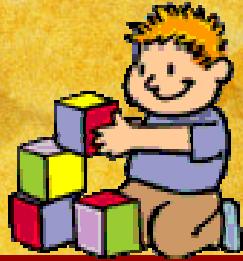


<b>CLONE_CLEARID</b>	Clear the task ID.
<b>CLONE_DETACHED</b>	The parent does not want a SIGCHILD signal sent on exit.
<b>CLONE_FILES</b>	Shares the table that identifies the open files.
<b>CLONE_FS</b>	Shares the table that identifies the root directory and the current working directory, as well as the value of the bit mask used to mask the initial file permissions of a new file.
<b>CLONE_IDLETASK</b>	Set PID to zero, which refers to an idle task. The idle task is employed when all available tasks are blocked waiting for resources.
<b>CLONE_NEWNS</b>	Create a new namespace for the child.
<b>CLONE_PARENT</b>	Caller and new task share the same parent process.
<b>CLONE_PTRACE</b>	If the parent process is being traced, the child process will also be traced.
<b>CLONE_SETTID</b>	Write the TID back to user space.
<b>CLONE_SETTLS</b>	Create a new TLS for the child.
<b>CLONE_SIGHAND</b>	Shares the table that identifies the signal handlers.
<b>CLONE_SYSVSEM</b>	Shares System V SEM_UNDO semantics.
<b>CLONE_THREAD</b>	Inserts this process into the same thread group of the parent. If this flag is true, it implicitly enforces CLONE_PARENT.
<b>CLONE_VFORK</b>	If set, the parent does not get scheduled for execution until the child invokes the <code>execve()</code> system call.
<b>CLONE_VM</b>	Shares the address space (memory descriptor and all page tables).

# Mac OS X Grand Central Dispatch (GCD)



- Provides a pool of available threads
- Designers can designate portions of applications, called *blocks*, that can be dispatched independently and run concurrently
- Concurrency is based on the number of cores available and the thread capacity of the system



# Block

- A simple extension to a language
- A block defines a self-contained unit of work
- Enables the programmer to encapsulate complex functions
- Scheduled and dispatched by queues
- Dispatched on a first-in-first-out basis
- Can be associated with an event source, such as a timer, network socket, or file descriptor



# Summary



- User-level threads
  - created and managed by a threads library that runs in the user space of a process
  - a mode switch is not required to switch from one thread to another
  - only a single user-level thread within a process can execute at a time
  - if one thread blocks, the entire process is blocked
- Kernel-level threads
  - threads within a process that are maintained by the kernel
  - a mode switch is required to switch from one thread to another
  - multiple threads within the same process can execute in parallel on a multiprocessor
  - blocking of a thread does not block the entire process
- Process/related to resource ownership
- Thread/related to program execution