Operating Systems: Internals and Design Principles

Chapter 9
Uniprocessor Scheduling

Seventh Edition
By William Stallings

Dave Bremer
Otago Polytechnic, N.Z.
©2008 Prentice Hall

Operating Systems:
Internals and Design Principles

"I take a two hour nap, from one o'clock to four."

— Yogi Berra

Processor Scheduling

- Aim is to assign processes to be executed by the processor in a way that meets system objectives, such as response time, throughput, and processor efficiency

- Broken down into three separate functions:

long term scheduling → medium term scheduling → short term scheduling

Table 9.1
Types of Scheduling

Long-term scheduling	The decision to add to the pool of processes to be executed
Medium-term scheduling	The decision to add to the number of processes that are partially or fully in main memory
Short-term scheduling	The decision as to which available process will be executed by the processor
I/O scheduling	The decision as to which process's pending I/O request shall be handled by an available I/O device

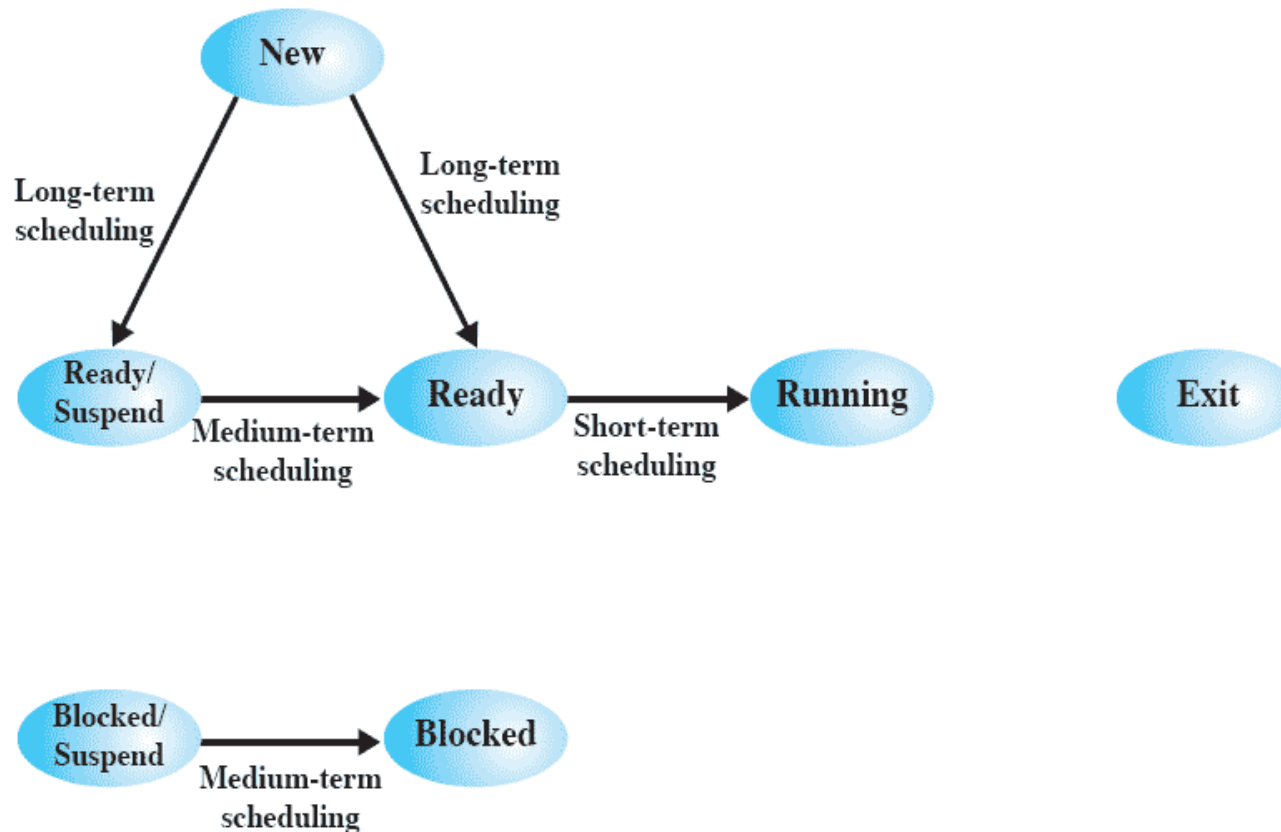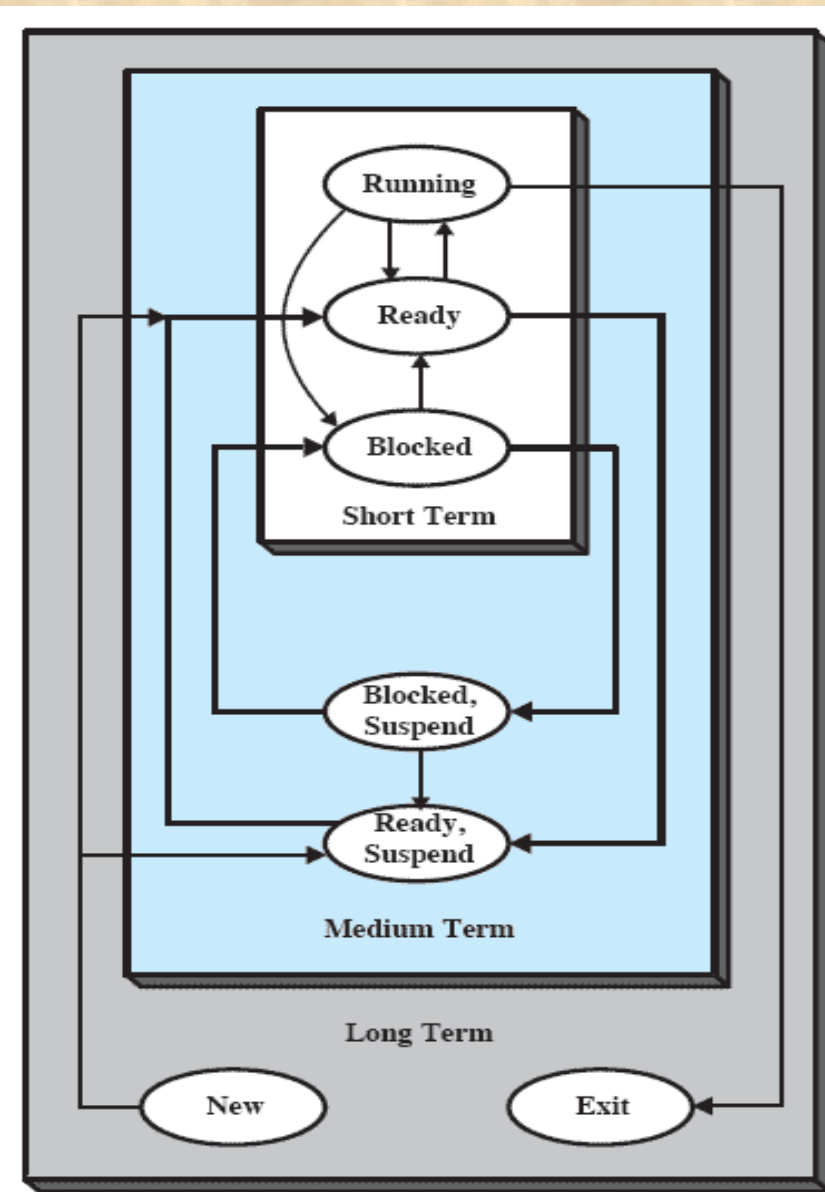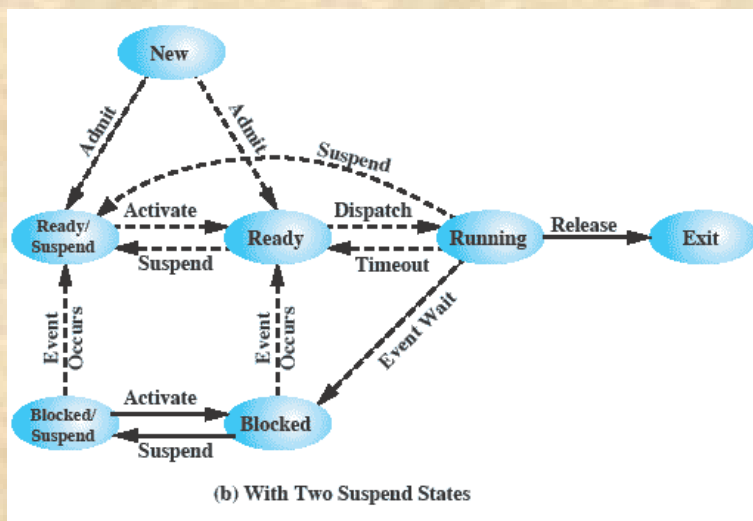# Scheduling and Process State Transitions

Figure 9.1 Scheduling and Process State Transitions

Figure 9.2 Nesting of Scheduling Functions
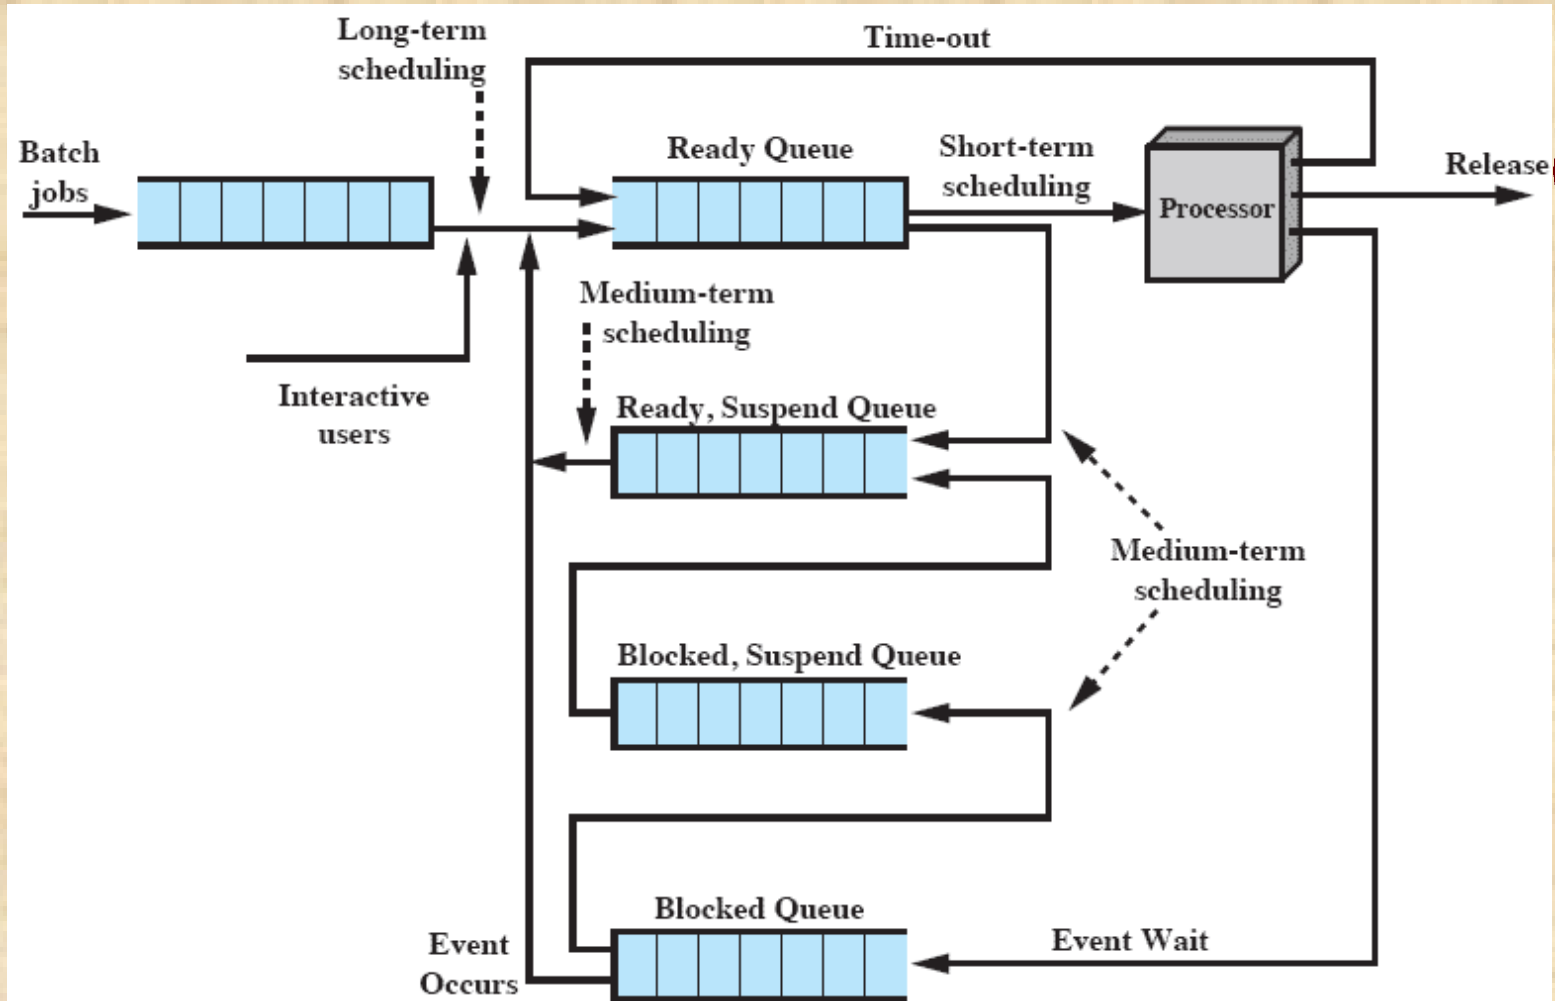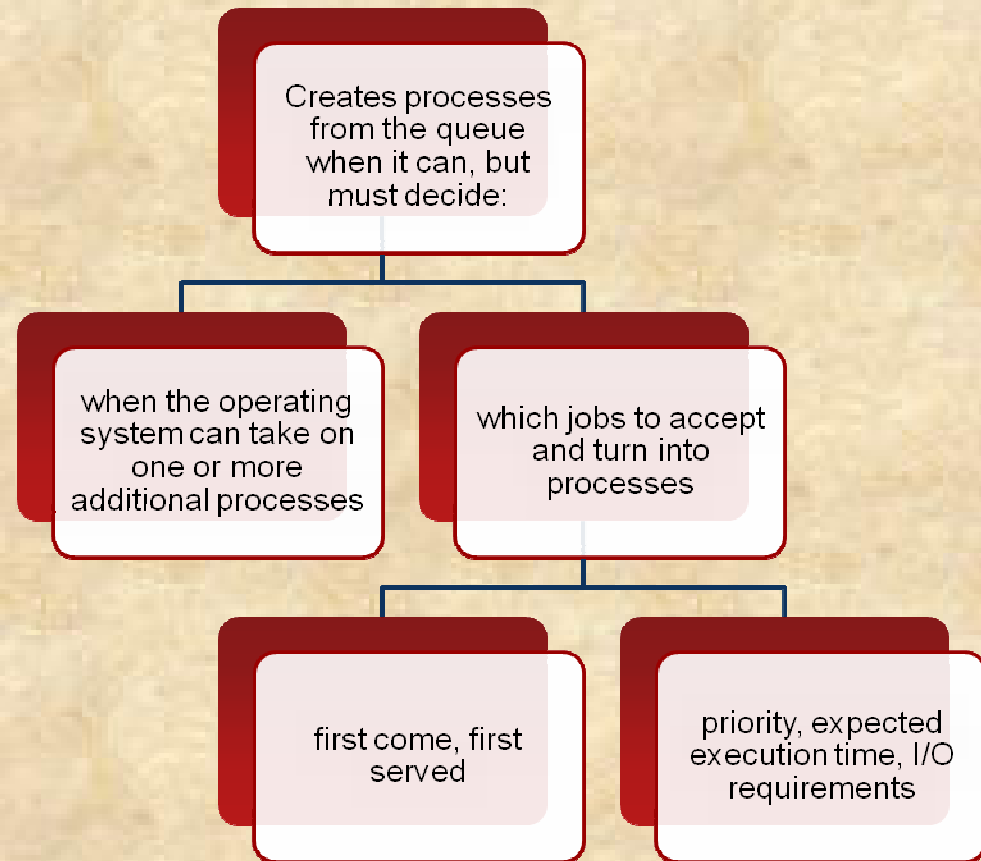
(Referencing figure 3.9b)

(b) With Two Suspend States

Figure 9.3   Queuing Diagram for Scheduling

Queuing Diagram

Long-Term Scheduler

- Determines which programs are admitted to the system for processing

- Controls the degree of multiprogramming
  - the more processes that are created, the smaller the percentage of time that each process can be executed
  - may limit to provide satisfactory service to the current set of processes

Creates processes from the queue when it can, but must decide:

when the operating system can take on one or more additional processes

which jobs to accept and turn into processes

first come, first served

priority, expected execution time, I/O requirements

Medium-Term Scheduling

- Part of the swapping function

- Swapping-in decisions are based on the need to manage the degree of multiprogramming

  - considers the memory requirements of the swapped-out processes

Short-Term Scheduling

- Known as the dispatcher

- Executes most frequently

- Makes the fine-grained decision of which process to execute next

- Invoked when an event occurs that may lead to the blocking of the current process or that may provide an opportunity to preempt a currently running process in favor of another

Examples:

- Clock interrupts
- I/O interrupts
- Operating system calls
- Signals (e.g., semaphores)

Short Term Scheduling Criteria

- Main objective is to allocate processor time to optimize certain aspects of system behavior

- A set of criteria is needed to evaluate the scheduling policy

User-oriented criteria
- relate to the behavior of the system as perceived by the individual user or process (such as response time in an interactive system)
- important on virtually all systems

System-oriented criteria
- focus in on effective and efficient utilization of the processor (rate at which processes are completed)
- generally of minor importance on single-user systems

Short-Term Scheduling Criteria: Performance

examples:
- response time
- throughput

Criteria can be classified into:

example:
- predictability

Performance-related

Non-performance related

quantitative

easily measured

qualitative

hard to measure

Table 9.2 Scheduling Criteria

User Oriented, Performance Related

Turnaround time   This is the interval of time between the submission of a process and its completion. Includes actual execution time plus time spent waiting for resources, including the processor. This is an appropriate measure for a batch job.

Response time   For an interactive process, this is the time from the submission of a request until the response begins to be received. Often a process can begin producing some output to the user while continuing to process the request. Thus, this is a better measure than turnaround time from the user's point of view. The scheduling discipline should attempt to achieve low response time and to maximize the number of interactive users receiving acceptable response time.

Deadlines   When process completion deadlines can be specified, the scheduling discipline should subordinate other goals to that of maximizing the percentage of deadlines met.

User Oriented, Other

Predictability   A given job should run in about the same amount of time and at about the same cost regardless of the load on the system. A wide variation in response time or turnaround time is distracting to users. It may signal a wide swing in system workloads or the need for system tuning to cure instabilities.

System Oriented, Performance Related

Throughput   The scheduling policy should attempt to maximize the number of processes completed per unit of time. This is a measure of how much work is being performed. This clearly depends on the average length of a process but is also influenced by the scheduling policy, which may affect utilization.

Processor utilization   This is the percentage of time that the processor is busy. For an expensive shared system, this is a significant criterion. In single-user systems and in some other systems, such as real-time systems, this criterion is less important than some of the others.

System Oriented, Other

Fairness   In the absence of guidance from the user or other system-supplied guidance, processes should be treated the same, and no process should suffer starvation.

Enforcing priorities   When processes are assigned priorities, the scheduling policy should favor higher-priority processes.
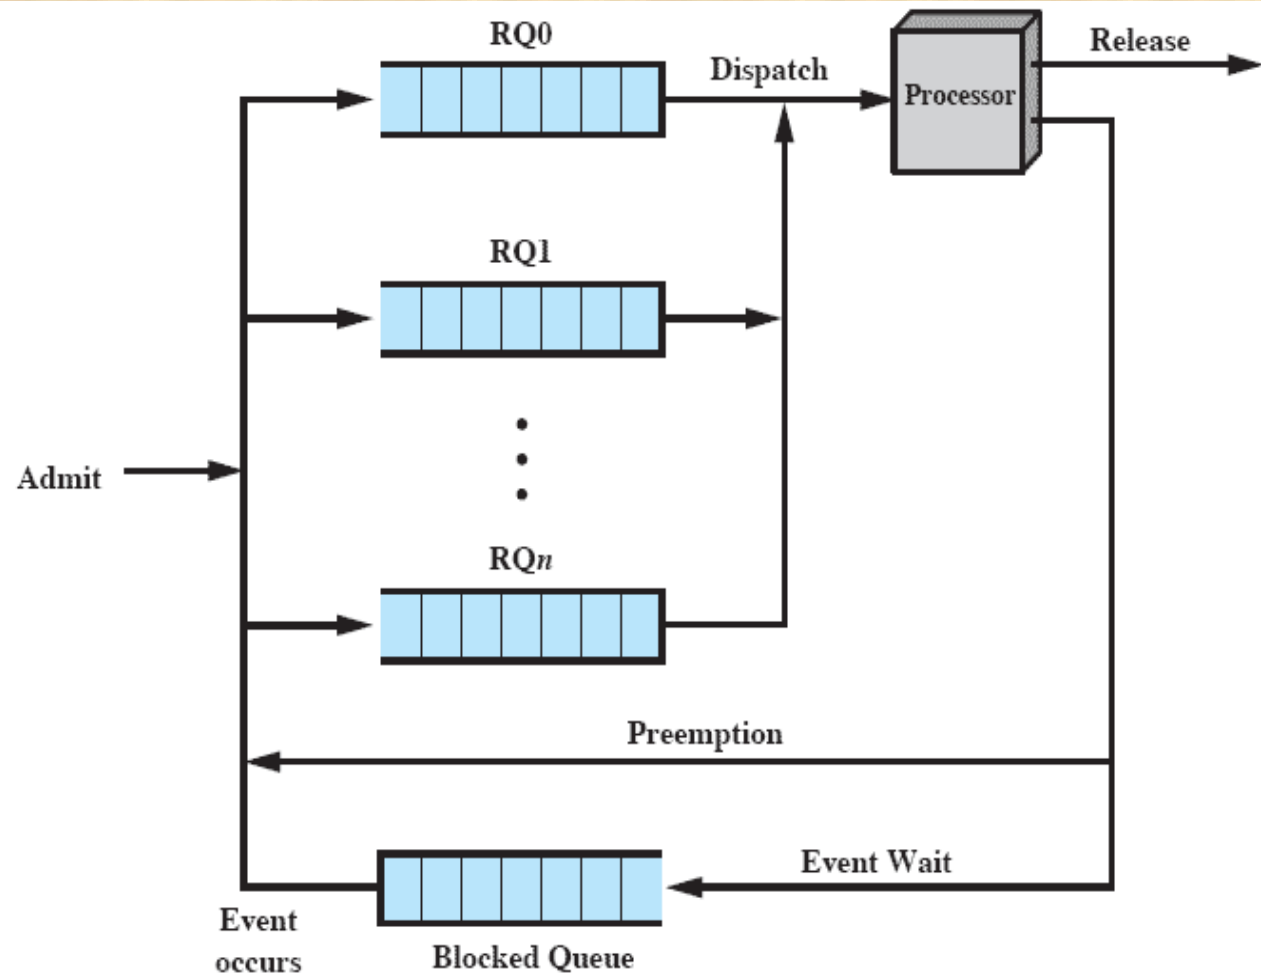
Priority Queuing

Figure 9.4   Priority Queuing

Alternative Scheduling Policies

Table 9.3 Characteristics of Various Scheduling Policies

	FCFS	Round robin	SPN	SRT	HRRN	Feedback
Selection function	$\max[w]$	constant	$\min[s]$	$\min[s-e]$	$\max\left(\dfrac{w+s}{s}\right)$	(see text)
Decision mode	Non-preemptive	Preemptive (at time quantum)	Non-preemptive	Preemptive (at arrival)	Non-preemptive	Preemptive (at time quantum)
Throughput	Not emphasized	May be low if quantum is too small	High	High	High	Not emphasized
Response time	May be high, especially if there is a large variance in process execution times	Provides good response time for short processes	Provides good response time for short processes	Provides good response time	Provides good response time	Not emphasized
Overhead	Minimum	Minimum	Can be high	Can be high	Can be high	Can be high
Effect on processes	Penalizes short processes; penalizes I/O bound processes	Fair treatment	Penalizes long processes	Penalizes long processes	Good balance	May favor I/O bound processes
Starvation	No	No	Possible	Possible	No	Possible

Selection Function

- Determines which process, among ready processes, is selected next for execution

- May be based on priority, resource requirements, or the execution characteristics of the process

- If based on execution characteristics then important quantities are:
  - w = time spent in system so far, waiting
  - e = time spent in execution so far
  - s = total service time required by the process, including e; generally, this quantity must be estimated or supplied by the user

Decision Mode

- Specifies the instants in time at which the selection function is exercised

- Two categories:
  - Nonpreemptive
  - Preemptive

Nonpreemptive vs Preemptive

Nonpreemptive

- once a process is in the running state, it will continue until it terminates or blocks itself for I/O

Preemptive

- currently running process may be interrupted and moved to ready state by the OS

- preemption may occur when new process arrives, on an interrupt, or periodically

Table 9.4

Process Scheduling Example

Process	Arrival Time	Service Time
A	0	3
B	2	6
C	4	4
D	6	5
E	8	2

Table 9.5

Comparison of Scheduling Policies

Process	A	B	C	D	E	
Arrival Time	0	2	4	6	8	
Service Time (T_s)	3	6	4	5	2	Mean
FCFS						
Finish Time	3	9	13	18	20	
Turnaround Time (T_r)	3	7	9	12	12	8.60
T_r/T_s	1.00	1.17	2.25	2.40	6.00	2.56
RR $q = 1$						
Finish Time	4	18	17	20	15	
Turnaround Time (T_r)	4	16	13	14	7	10.80
T_r/T_s	1.33	2.67	3.25	2.80	3.50	2.71
RR $q = 4$						
Finish Time	3	17	11	20	19	
Turnaround Time (T_r)	3	15	7	14	11	10.00
T_r/T_s	1.00	2.5	1.75	2.80	5.50	2.71
SPN						
Finish Time	3	9	15	20	11	
Turnaround Time (T_r)	3	7	11	14	3	7.60
T_r/T_s	1.00	1.17	2.75	2.80	1.50	1.84
SRT						
Finish Time	3	15	8	20	10	
Turnaround Time (T_r)	3	13	4	14	2	7.20
T_r/T_s	1.00	2.17	1.00	2.80	1.00	1.59
HRRN						
Finish Time	3	9	13	20	15	
Turnaround Time (T_r)	3	7	9	14	7	8.00
T_r/T_s	1.00	1.17	2.25	2.80	3.5	2.14
FB $q = 1$						
Finish Time	4	20	16	19	11	
Turnaround Time (T_r)	4	18	12	13	3	10.00
T_r/T_s	1.33	3.00	3.00	2.60	1.5	2.29
FB $q = 2i$						
Finish Time	4	17	18	20	14	

First-Come-First-Served (FCFS)

- Simplest scheduling policy

- Also known as first-in-first-out (FIFO) or a strict queuing scheme

- When the current process ceases to execute, the longest process in the Ready queue is selected

- Performs much better for long processes than short ones

- Tends to favor processor-bound processes over I/O-bound processes
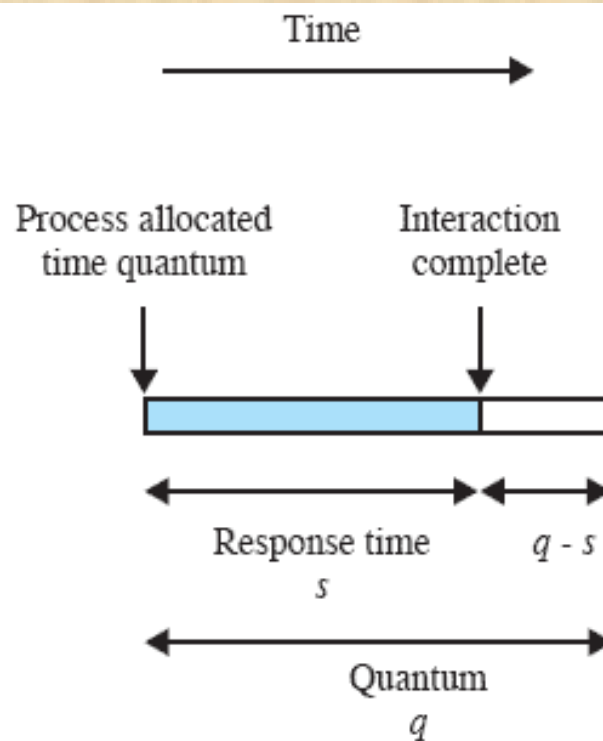
First-Come-First Served (FCFS)

Round Robin

- Uses preemption based on a clock

- Also known as **time slicing** because each process is given a slice of time before being preempted

- Principal design issue is the length of the time quantum, or slice, to be used

- Particularly effective in a general-purpose time-sharing system or transaction processing system

- One drawback is its relative treatment of processor-bound and I/O-bound processes

Round-Robin (RR), $q = 1$

Figure 9.6a

Effect of
Size of
Preemption
Time
Quantum

Time

Process allocated
time quantum

Interaction
complete

Response time
s

q - s

Quantum
q

(a) Time quantum greater than typical interaction

Figure 9.6b

Effect of Size of Preemption Time Quantum
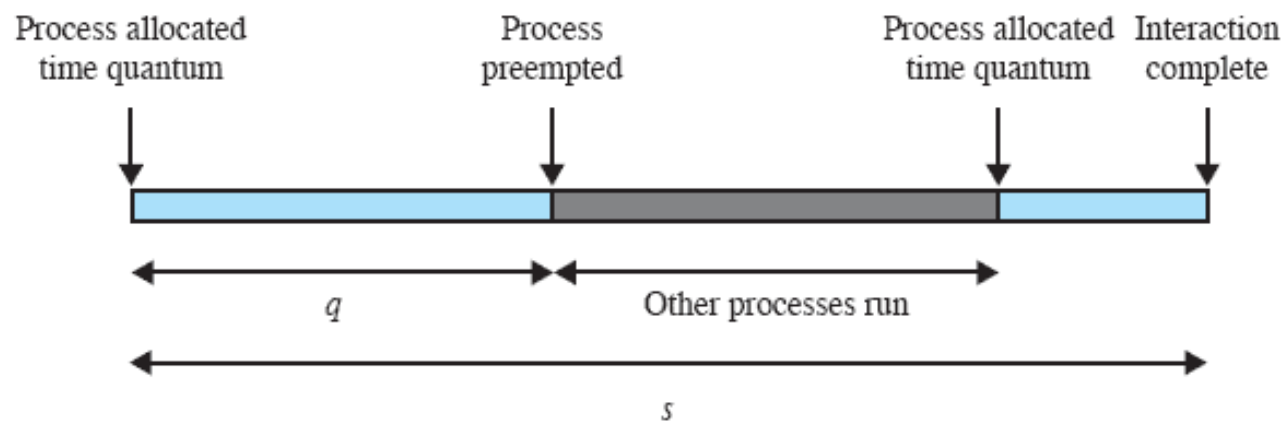
(b) Time quantum less than typical interaction

Figure 9.6 Effect of Size of Preemption Time Quantum

Virtual Round Robin (VRR)

Figure 9.7   Queuing Diagram for Virtual Round-Robin Scheduler

Shortest Process Next (SPN)

- Nonpreemptive policy in which the process with the shortest expected processing time is selected next

- A short process will jump to the head of the queue

- Possibility of starvation for longer processes

- One difficulty is the need to know, or at least estimate, the required processing time of each process

- If the programmer's estimate is substantially under the actual running time, the system may abort the job

Shortest Process Next (SPN)

A
B
C
D
E

Exponential Smoothing Coefficients

Figure 9.8   Exponential Smoothing Coefficients

Use Of Exponential Averaging

(a) Increasing function

Use Of Exponential Averaging

(b) Decreasing function

Shortest Remaining Time (SRT)

- Preemptive version of SPN

- Scheduler always chooses the process that has the shortest expected remaining processing time

- Risk of starvation of longer processes

- Should give superior turnaround time performance to SPN because a short job is given immediate preference to a running longer job

Shortest Remaining Time (SRT)

Highest Response Ratio Next (HRRN)

- Chooses next process with the greatest ratio

- Attractive because it accounts for the age of the process

- While shorter jobs are favored, aging without service increases the ratio so that a longer process will eventually get past competing shorter jobs

$$Ratio = \frac{time\ spent\ waiting + expected\ service\ time}{expected\ service\ time}$$

**Highest Response Ratio Next (HRRN)**

Feedback Scheduling
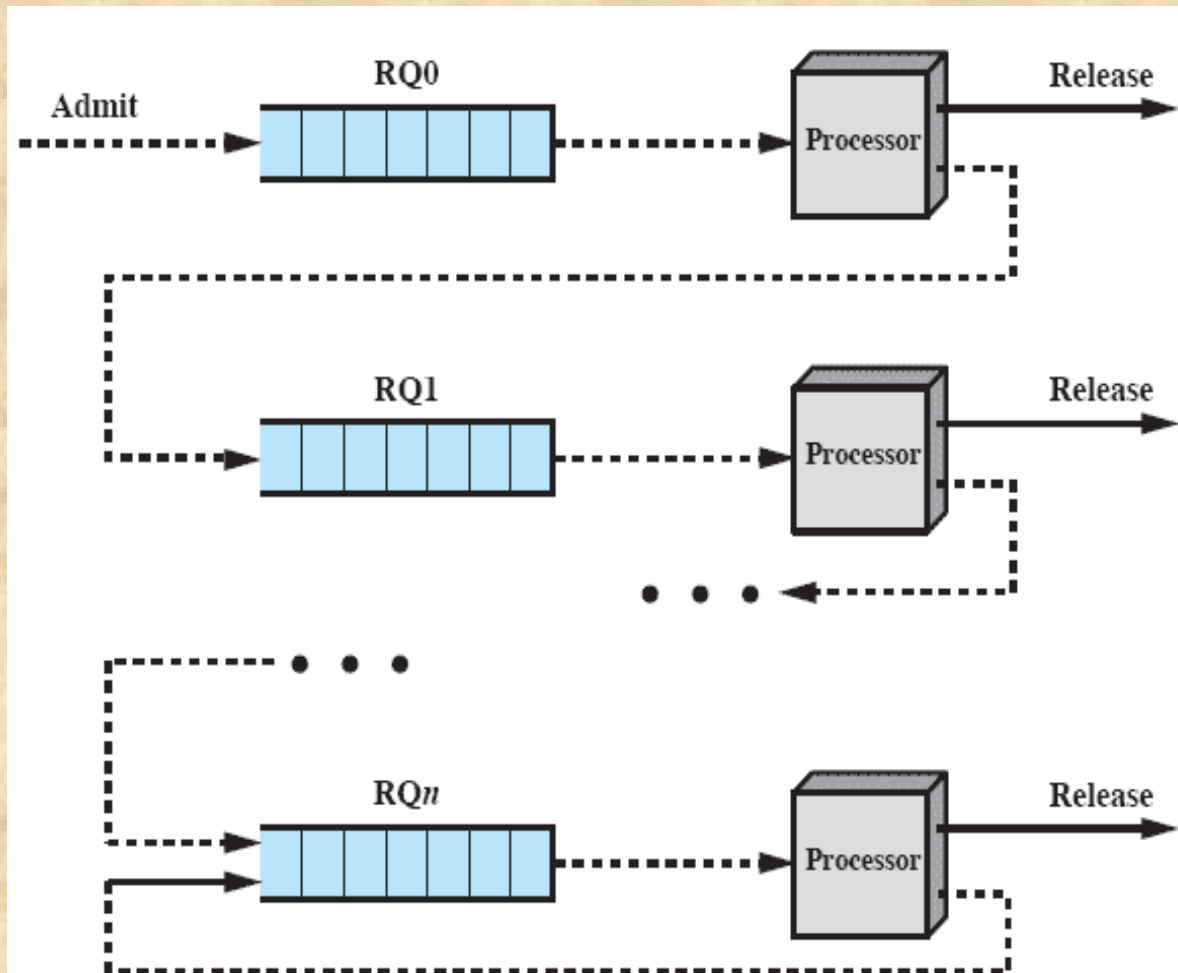
Figure 9.10   Feedback Scheduling

Feedback Performance

Feedback
$q = 1$

Feedback
$q = 2^i$

Performance Comparison

- Any scheduling discipline that chooses the next item to be served independent of service time obeys the relationship:

$$\frac{T_r}{T_s} = \frac{1}{1 - \rho}$$

where

T_r = turnaround time or residence time; total time in system, waiting plus execution

T_s = average service time; average time spent in Running state

ρ = processor utilization

Table 9.6

Formulas for Single-Server Queues with Two Priority Categories

Assumptions: **1.** Poisson arrival rate.

2. Priority 1 items are serviced before priority 2 items.

3. First-come-first-served dispatching for items of equal priority.

4. No item is interrupted while being served.

5. No items leave the queue (lost calls delayed).

(a) General formulas

$$\lambda = \lambda_1 + \lambda_2$$

$$\rho_1 = \lambda_1 T_{s1}; \quad \rho_2 = \lambda_2 T_{s2}$$

$$\rho = \rho_1 + \rho_2$$

$$T_s = \frac{\lambda_1}{\lambda} T_{s1} + \frac{\lambda_2}{\lambda} T_{s2}$$

$$T_r = \frac{\lambda_1}{\lambda} T_{r1} + \frac{\lambda_2}{\lambda} T_{r2}$$

$x = 2\sqrt{\mu_{23}}$

$\frac{3}{3}$

| **(b) No interrupts; exponential service times**	**(c) Preemptive-resume queuing discipline; exponential service times**
$$T_{r1} = T_{s1} + \frac{\rho_1 T_{s1} + \rho_2 T_{s2}}{1 - \rho_1}$$ $$T_{r2} = T_{s2} + \frac{T_{r1} - T_{s1}}{1 - \rho}$$	$$T_{r1} = T_{s1} + \frac{\rho_1 T_{s1}}{1 - \rho_1}$$ $$T_{r2} = T_{s2} + \frac{1}{1 - \rho_1}\left(\rho_1 T_{s2} + \frac{\rho T_s}{1 - \rho}\right)$$

Overall Normalized Response Time

Figure 9.11 Overall Normalized Response Time

Normalized Response Time for Shorter Processes


Figure 9.12 Normalized Response Time for Shorter Processes

Normalized Response Time for Longer Processes

**Figure 9.13 Normalized Response Time for Longer Processes**

Results

Simulation

Figure 9.14 Simulation Results for Normalized Turnaround Time

Figure 9.15 Simulation Results for Waiting Time

Fair-Share Scheduling

- Scheduling decisions based on the process sets

- Each user is assigned a share of the processor

- Objective is to monitor usage to give fewer resources to users who have had more than their fair share and more to those who have had less than their fair share

Fair-Share Scheduler

Time	Process A Priority	Process CPU count	Group CPU count	Process B Priority	Process CPU count	Group CPU count	Process C Priority	Process CPU count	Group CPU count
0	60	0 1 2 · · 60	0 1 2 · · 60	60	0	0	60	0	0
1	90	30	30	60	0 1 2 · · 60	0 1 2 · · 60	60	0	0 1 2 · · 60
2	74	15 16 17 · · 75	15 16 17 · · 75	90	30	30	75	0	30
3	96	37	37	74	15 16 17 · · 75	15 16 17 · · 75	67	0 1 2 · · 60	15 16 17 · · 75
4	78	18 19 20 · · 78	18 19 20 · · 78	81	7	37	93	30	37
5	98	39	39	70	3	18	76	15	18

Group 1 (Process A)   Group 2 (Process B, Process C)

Colored rectangle represents executing process

Figure 9.16 Example of Fair Share Scheduler—Three Processes, Two Groups

Traditional UNIX Scheduling

- Used in both SVR3 and 4.3 BSD UNIX
    - these systems are primarily targeted at the time-sharing interactive environment

    - Designed to provide good response time for interactive users while ensuring that low-priority background jobs do not starve

    - Employs multilevel feedback using round robin within each of the priority queues

    - Makes use of one-second preemption

    - Priority is based on process type and execution history

Scheduling Formula

$$CPU_j(i) = \frac{CPU_j(i-1)}{2}$$

$$P_j(i) = Base_j + \frac{CPU_j(i)}{2} + nice_j$$

where

$CPU_j(i)$ = measure of processor utilization by process j through interval i

$P_j(i)$ = priority of process j at beginning of interval i; lower values equal higher priorities

$Base_j$ = base priority of process j

$nice_j$ = user-controllable adjustment factor

Bands

- Used to optimize access to block devices and to allow the operating system to respond quickly to system calls

- In decreasing order of priority, the bands are:

Swapper

↓

Block I/O device control

↓

File manipulation

↓

Character I/O device control

↓

User processes

Example of Traditional UNIX Process Scheduling

Figure 9.17 Example of Traditional UNIX Process Scheduling

Summary

- The operating system must make three types of scheduling decisions with respect to the execution of processes:

  - Long-term – determines when new processes are admitted to the system

  - Medium-term – part of the swapping function and determines when a program is brought into main memory so that it may be executed

  - Short-term – determines which ready process will be executed next by the processor

  - From a user's point of view, response time is generally the most important characteristic of a system; from a system point of view, throughput or processor utilization is important

  - Algorithms:

  - FCFS, Round Robin, SPN, SRT, HRRN, Feedback